

Multilayer Perceptrons (MLPs) and Activation functions

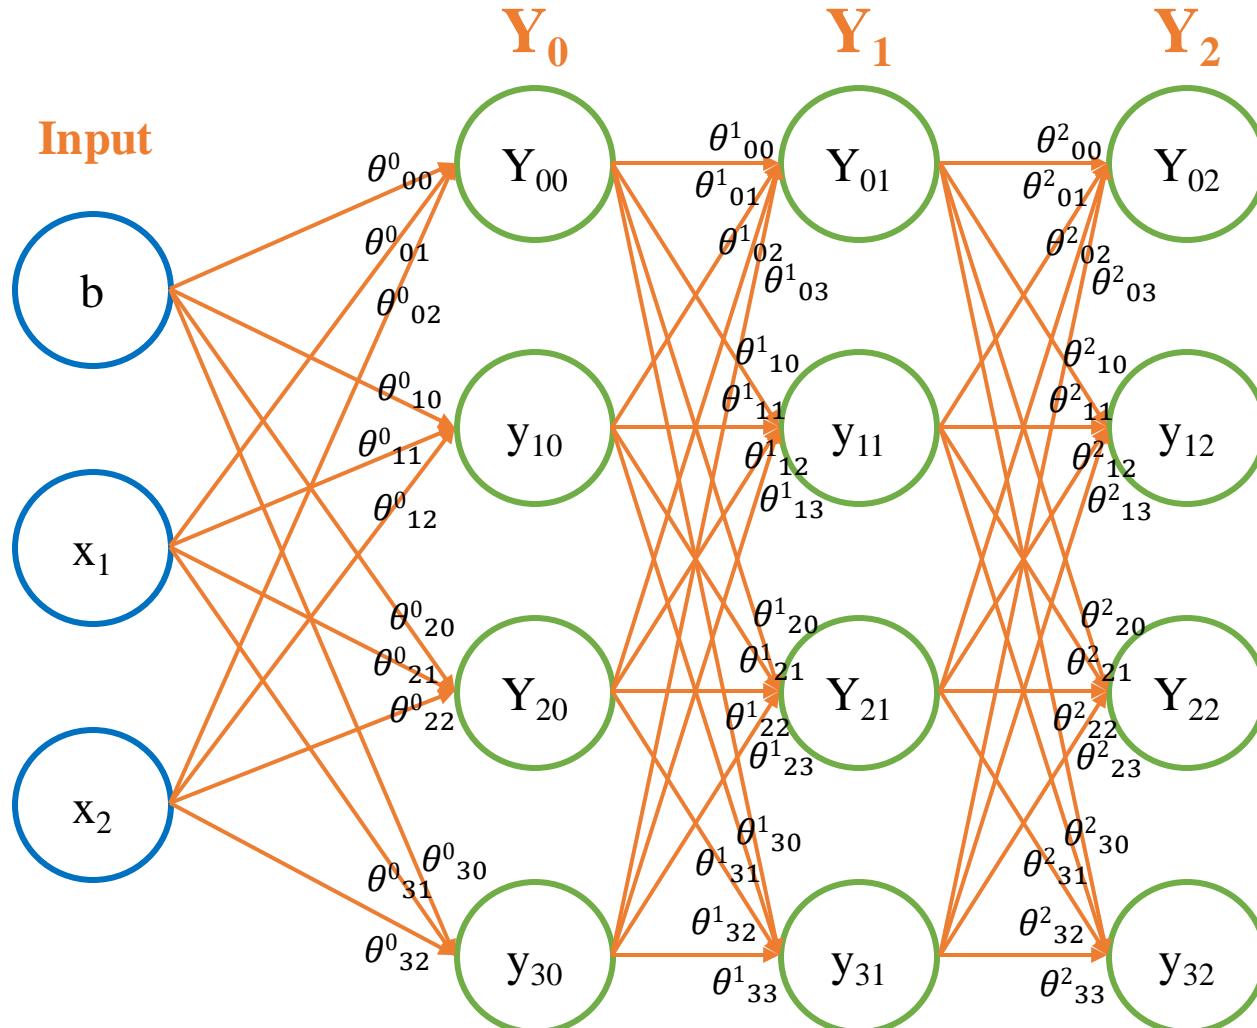
Exercise

Dinh-Thang Duong – TA

Yen-Linh Vu – STA

Getting Started

❖ Objectives



Our objectives:

- Review the Multilayer Perceptrons and Activation functions concepts.
- Practice how to implement Multilayer Perceptrons with PyTorch on 3 problems: Car Performance Prediction, Non-linear Data Classification and Sentiment Image Analysis.
- Examine the performance Multilayer Perceptrons on multiple configurations.

Outline

- Review
- Car Performance Prediction
- Non-linear Data Classification
- Sentiment Image Analysis
- Question

Review

Review

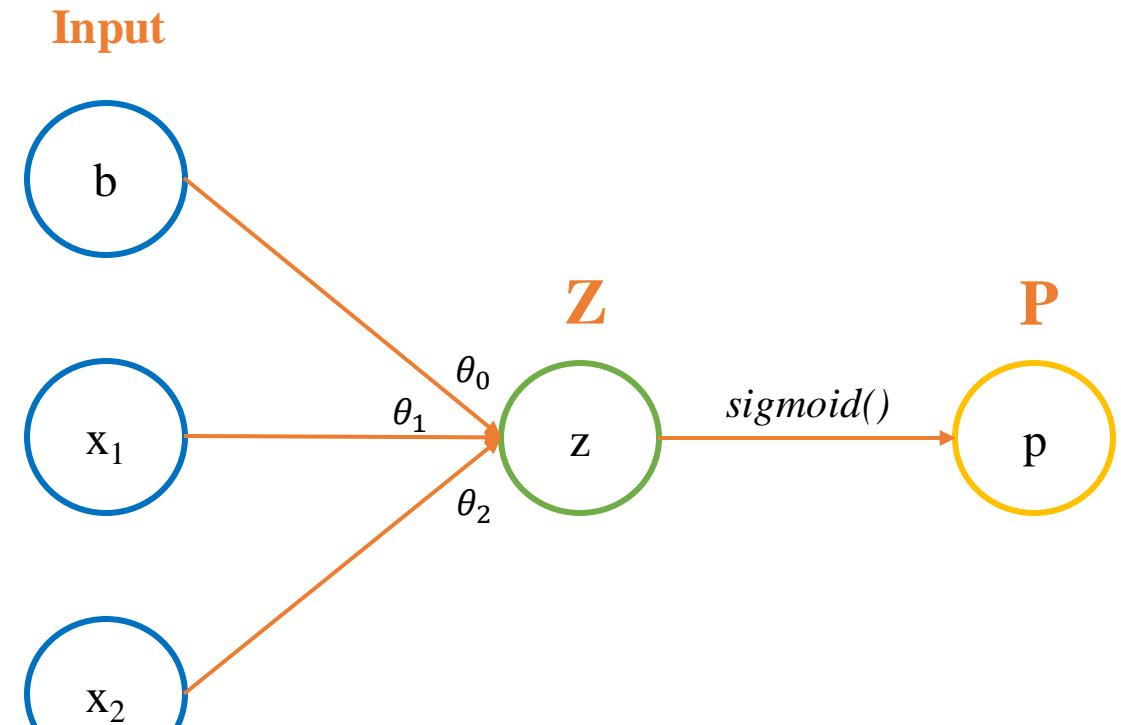
❖ Getting Started

x : Input data x_i

z : Linear value (z)

p : Activation value (p)

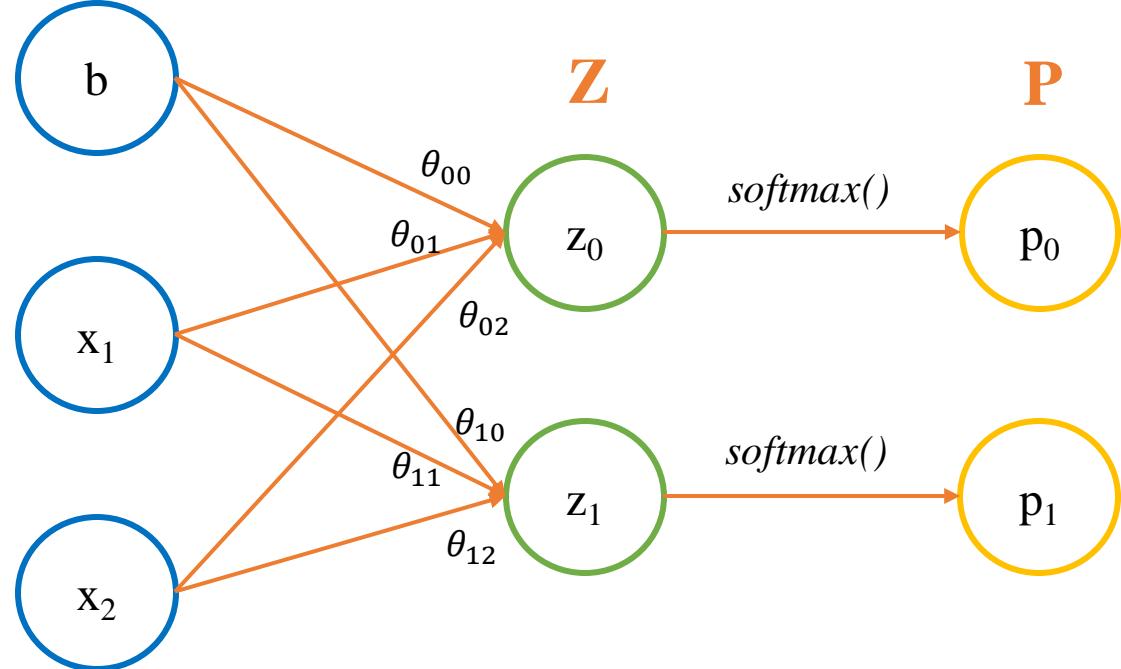
→ : Direction



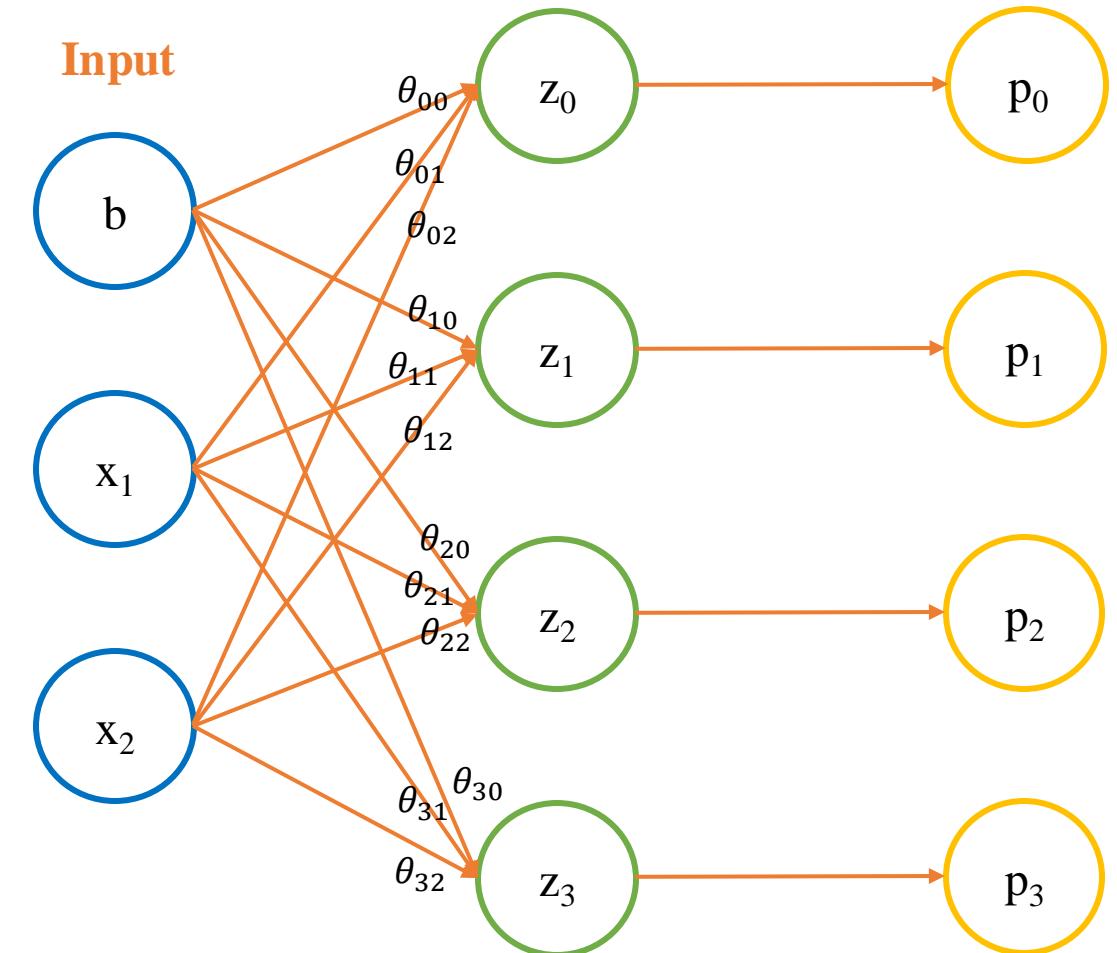
Logistic Regression Computation

Review

❖ Softmax Regression Illustration

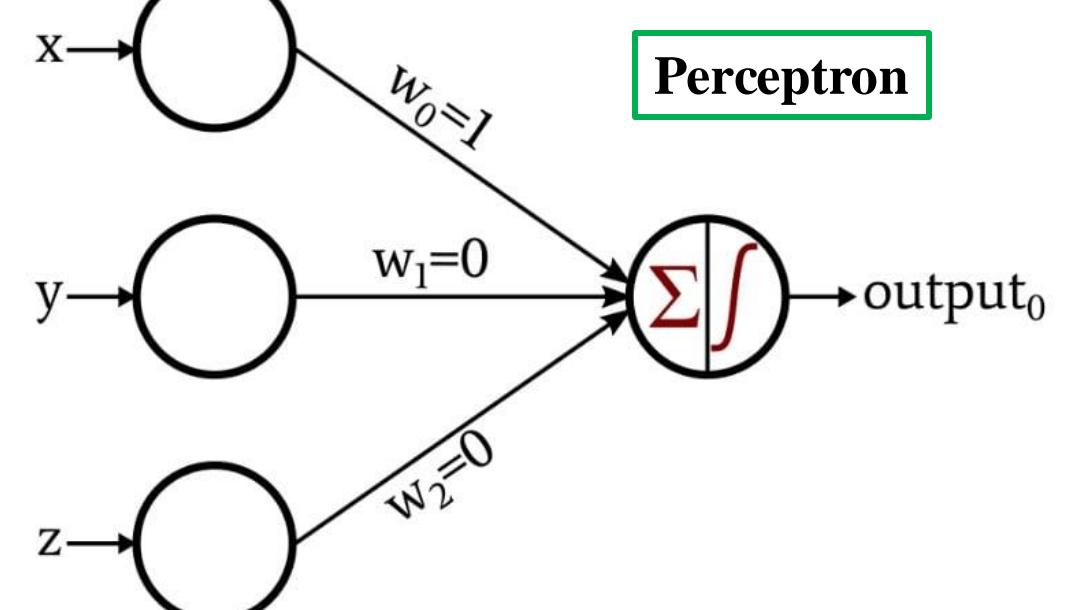
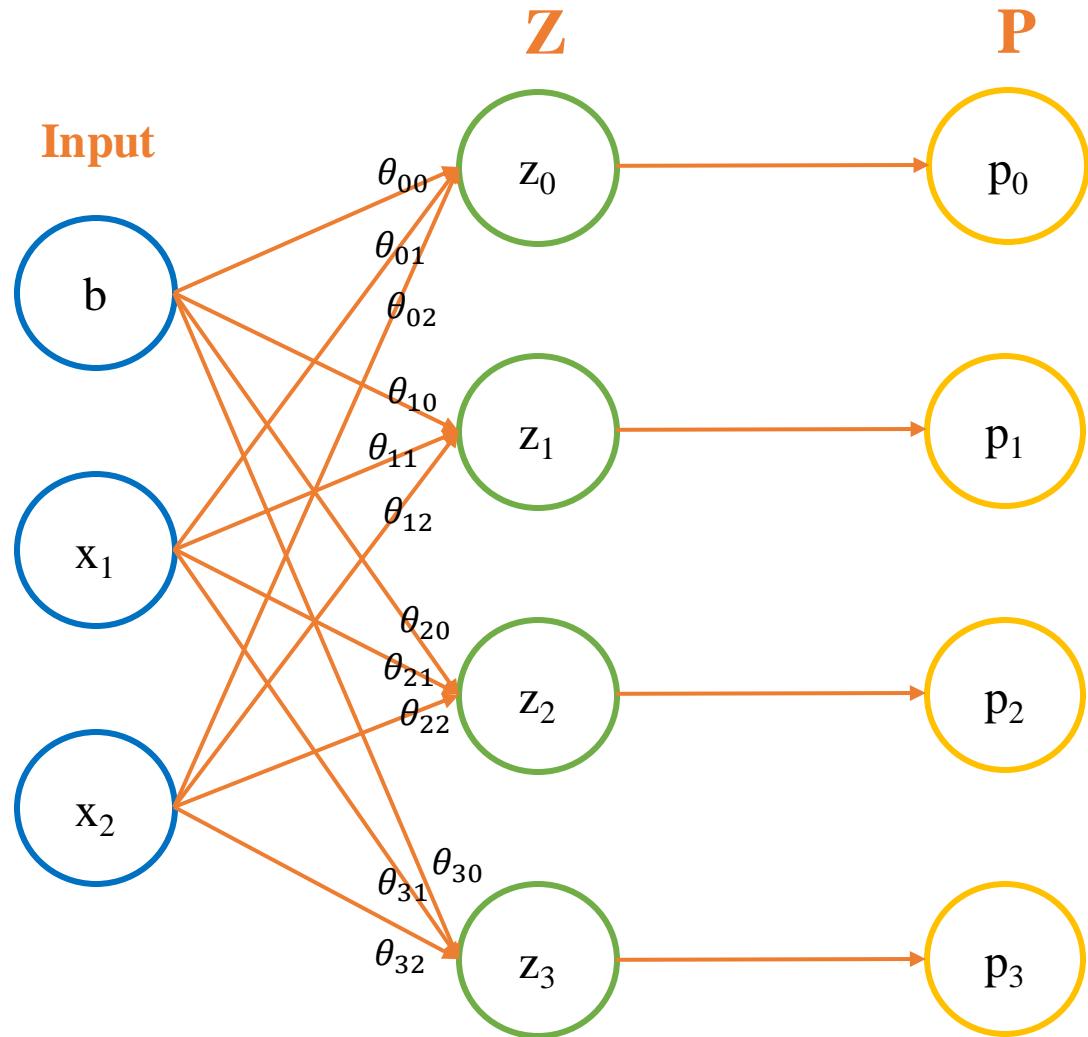


Softmax Regression (2 class)



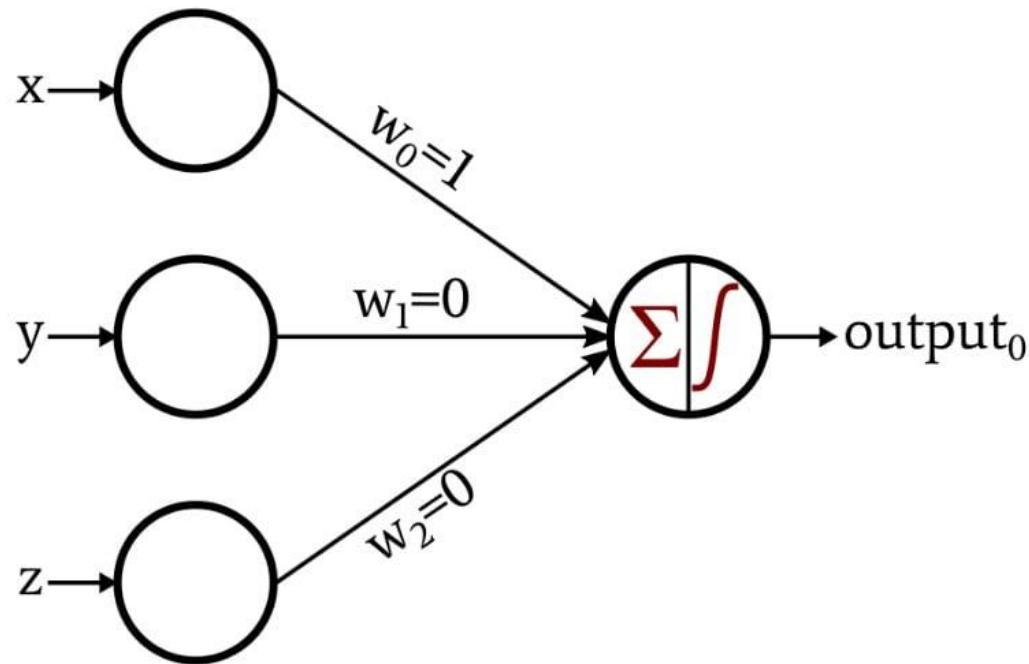
Softmax Regression (4 class)

❖ Perceptron



Review

❖ Perceptron

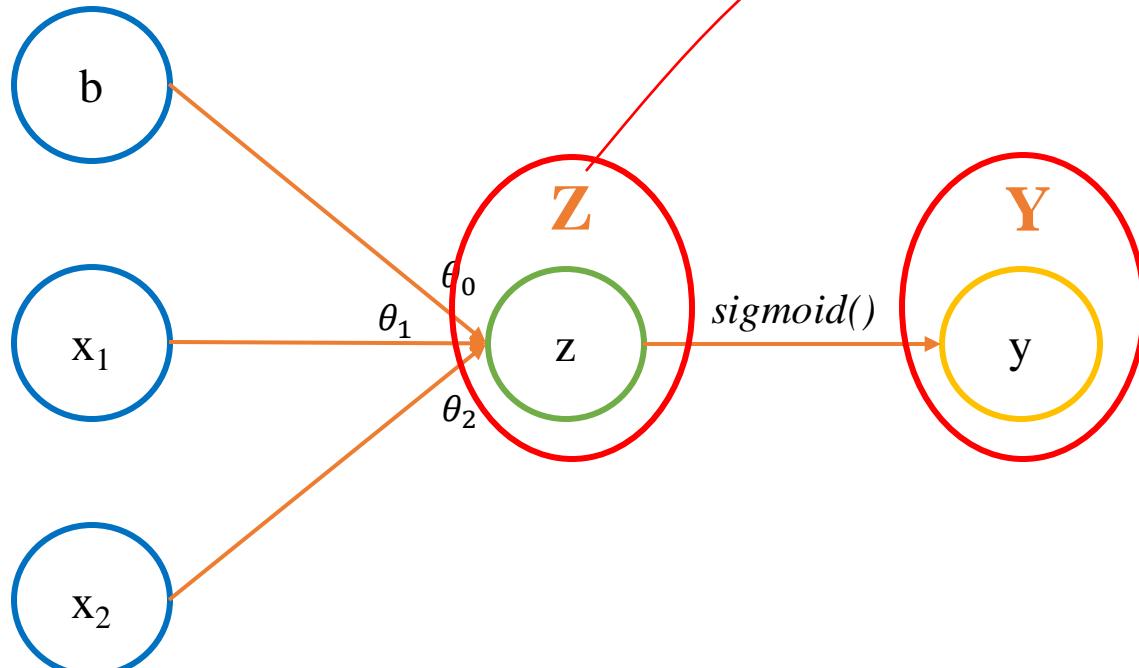


Perceptron: A simple neural network with a single layer that processes inputs using weighted sums, a bias, and an activation function for binary classification of linearly separable data.

Review

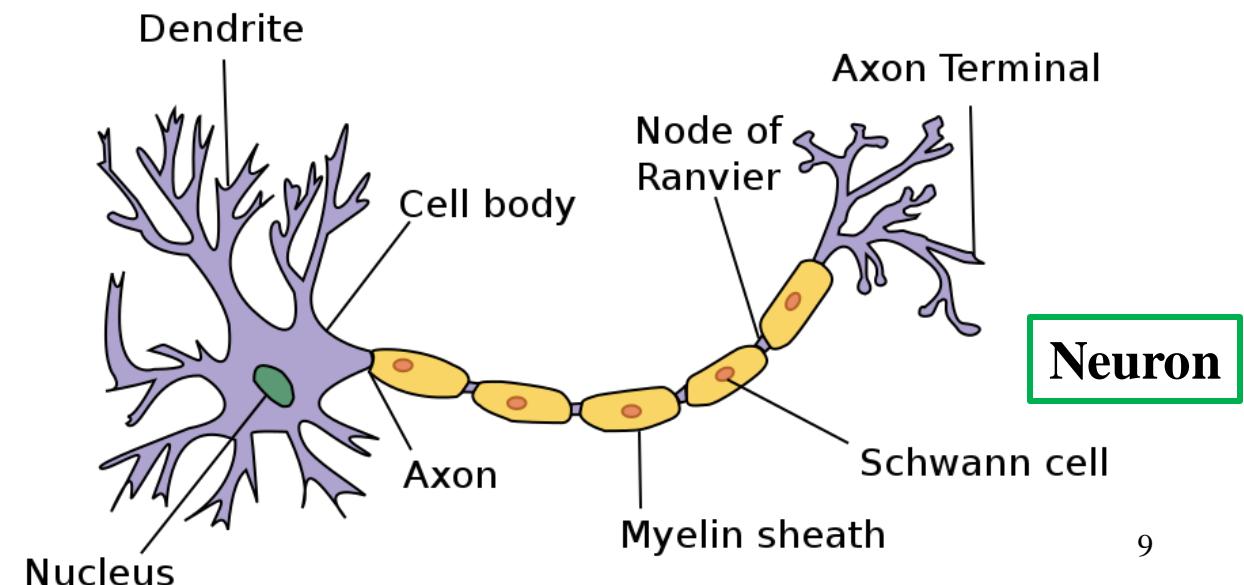
❖ Perceptron

Input



1. Dot product $Z = X \cdot \theta$

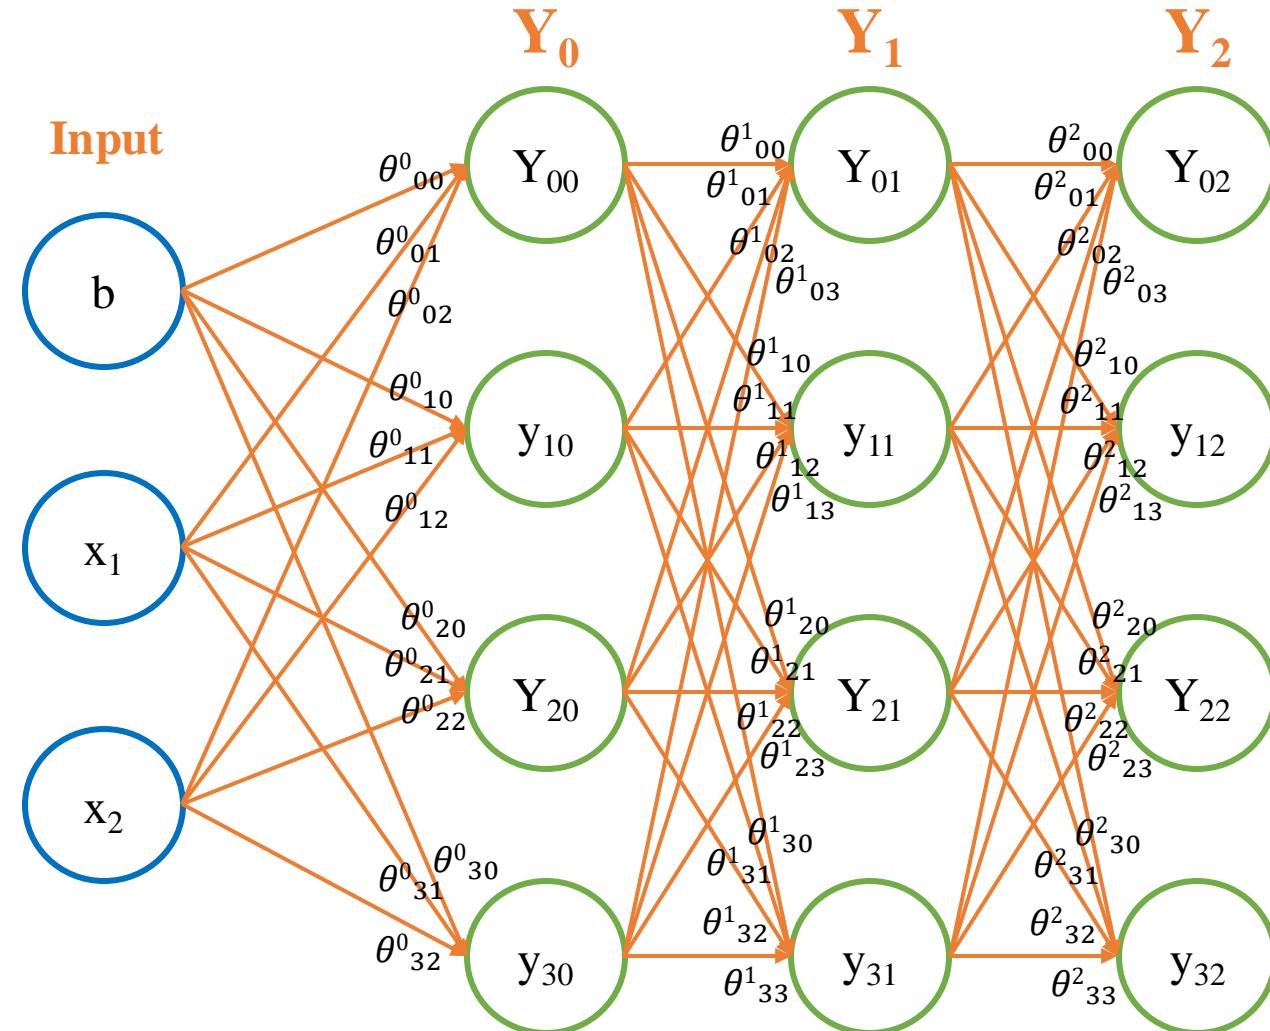
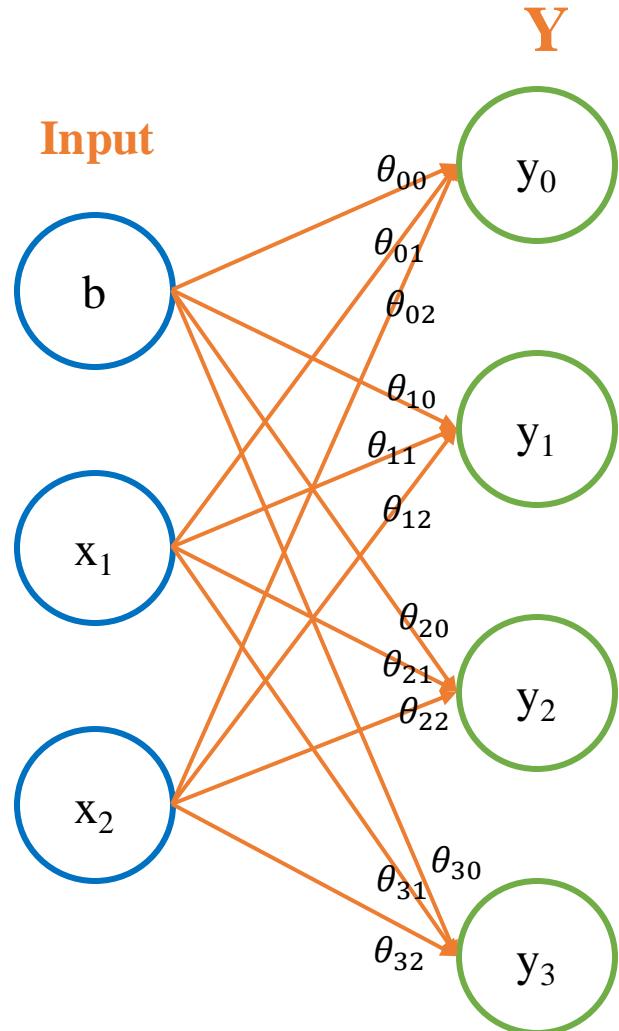
2. Output of some function (activation function):
- e.g: $y = \text{sigmoid}(Z)$



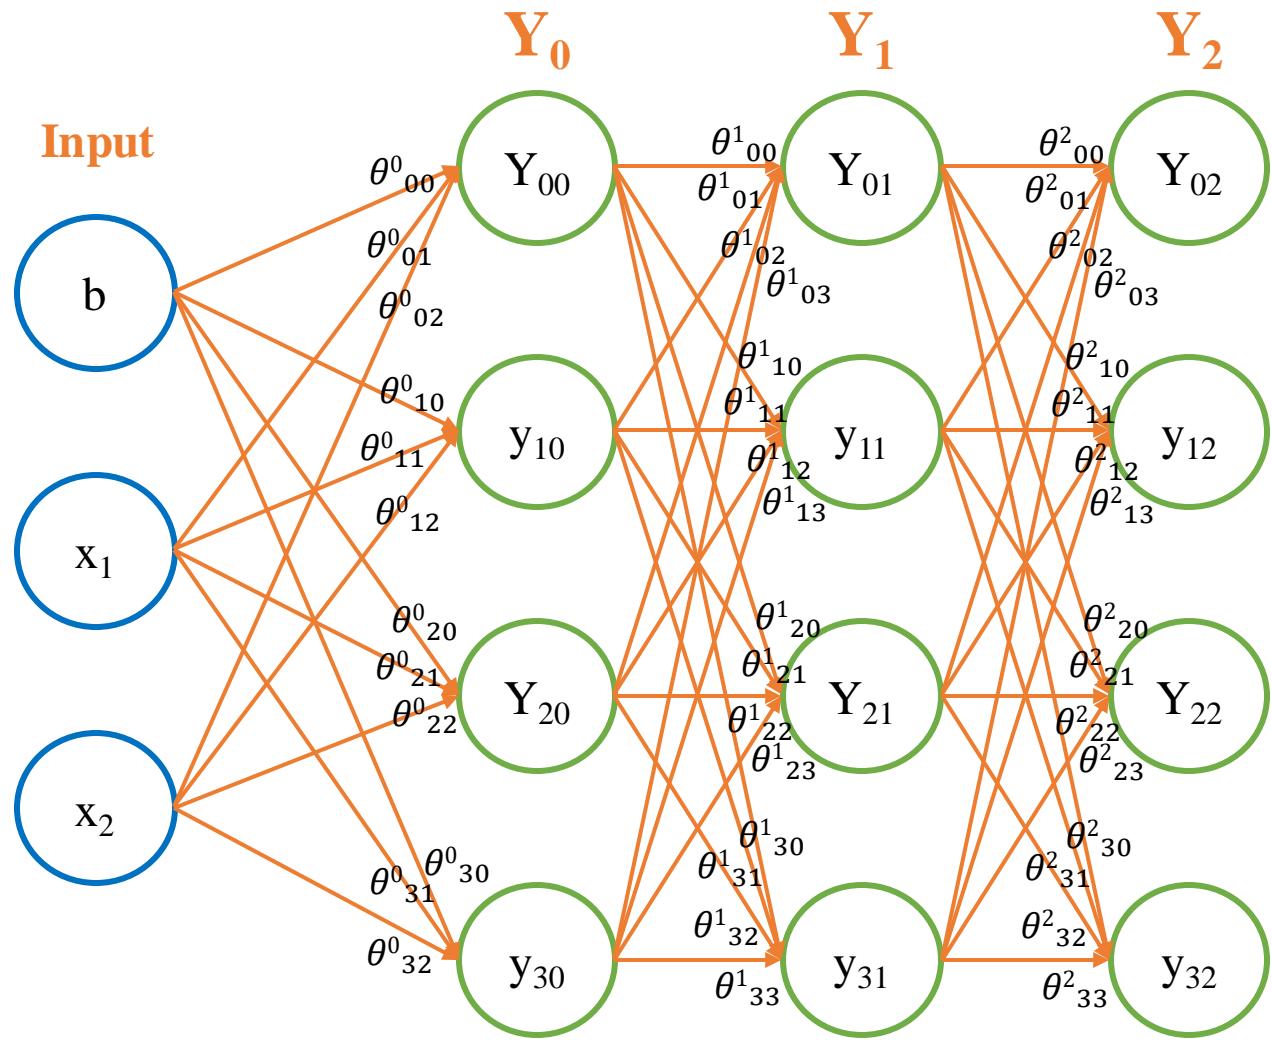
Neuron

Review

❖ Multilayer Perceptron



❖ Multilayer Perceptron



Multilayer Perceptron (MLP): A type of neural network with an input layer, one or more hidden layers, and an output layer, using non-linear activation functions to handle both linear and non-linear problems.

Review

❖ Multilayer Perceptron



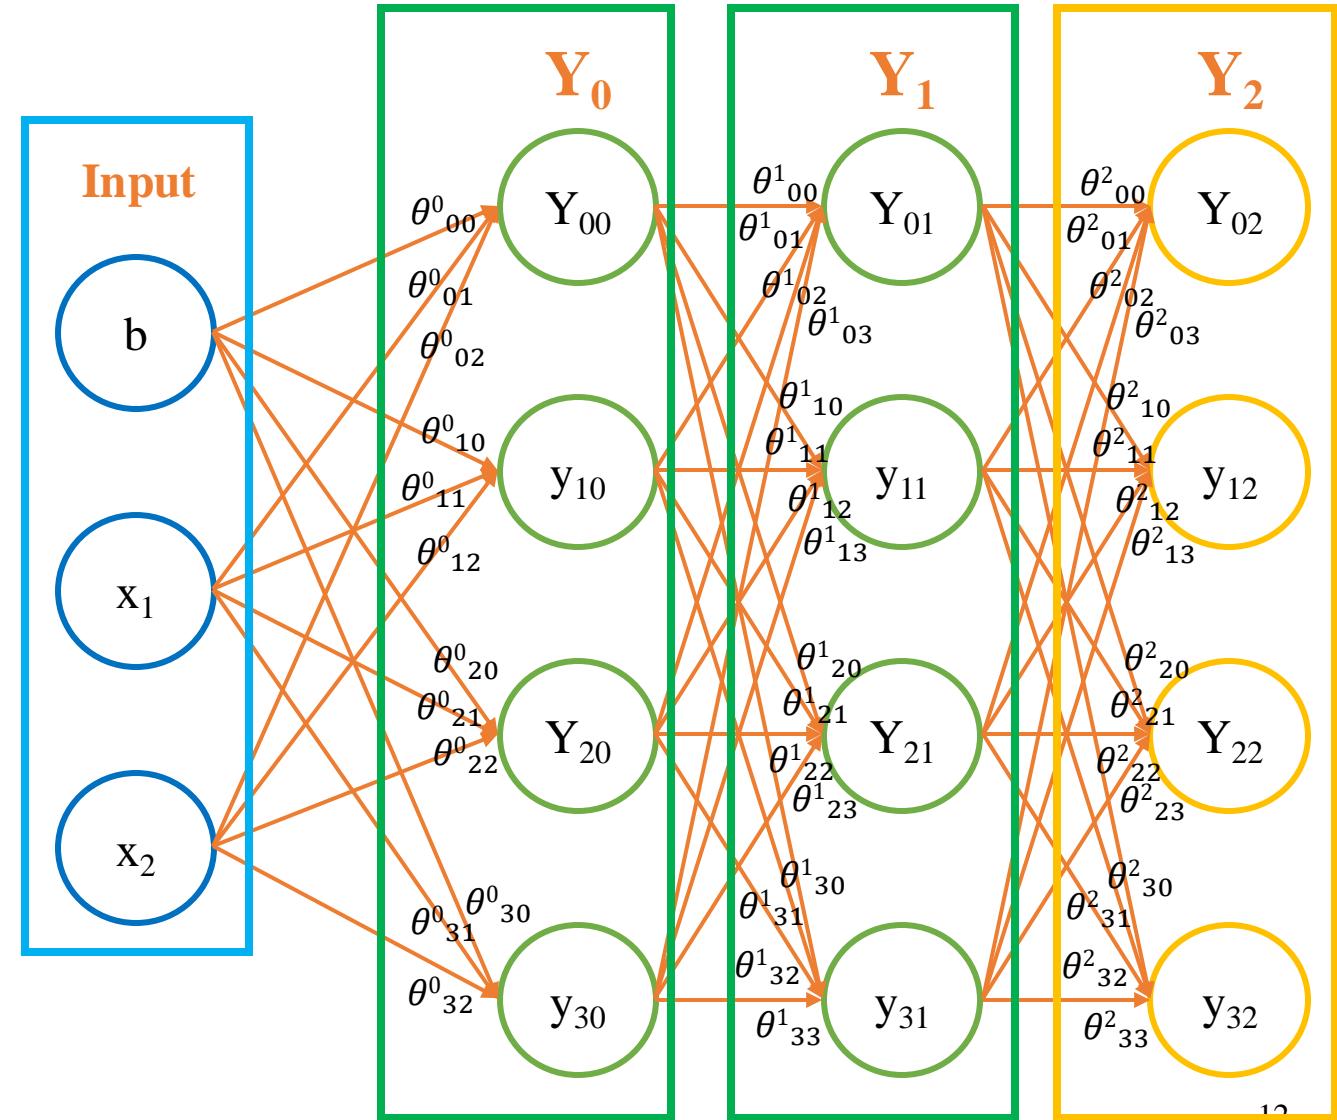
Input layer



Hidden layer

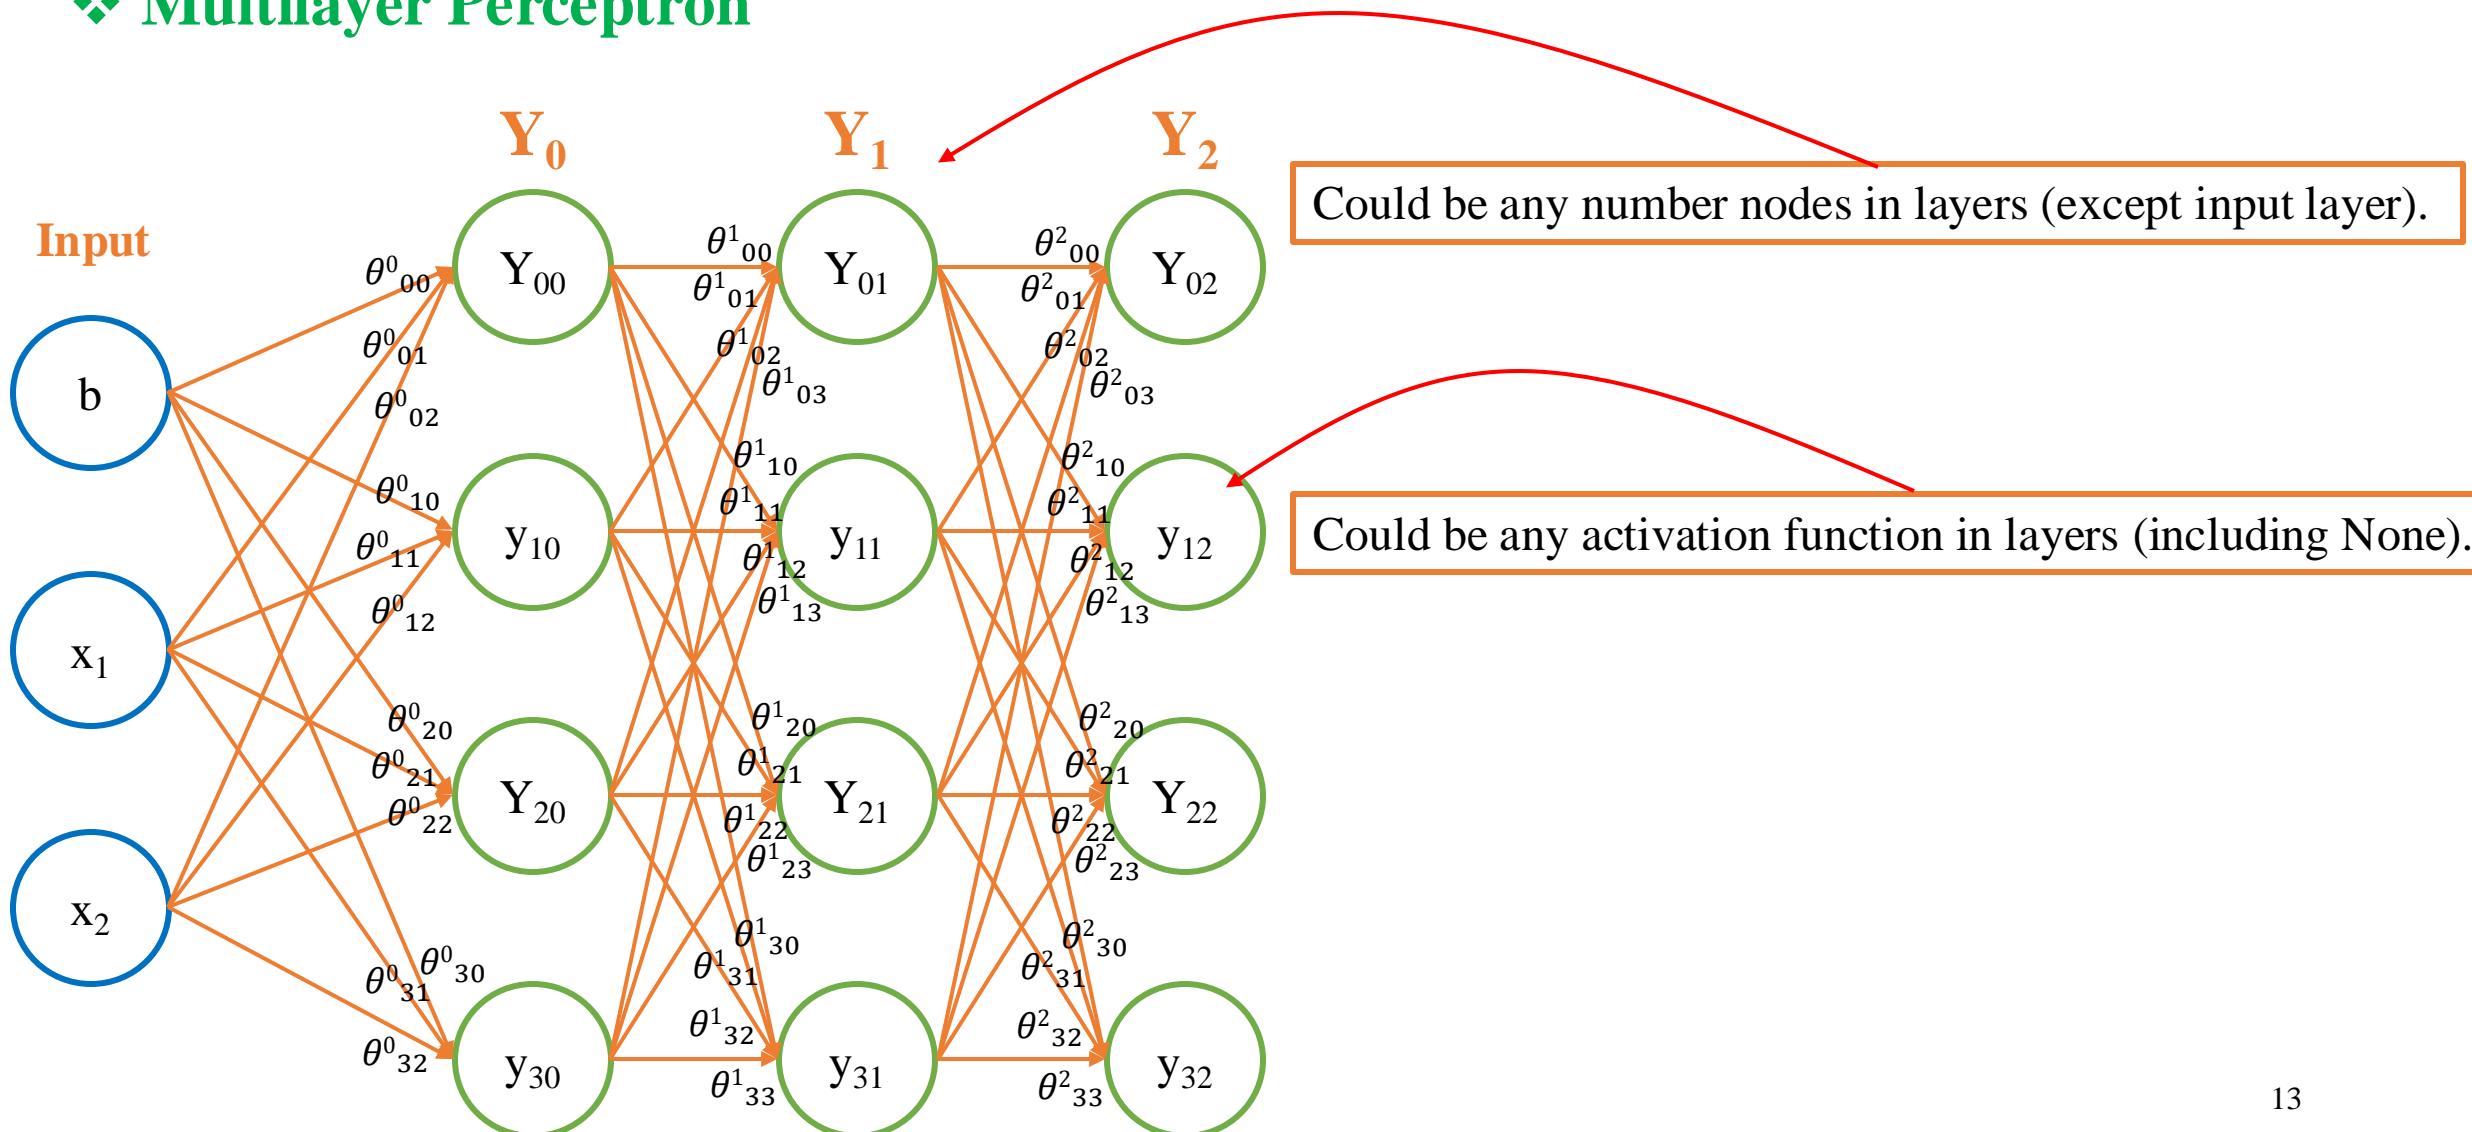


Output layer



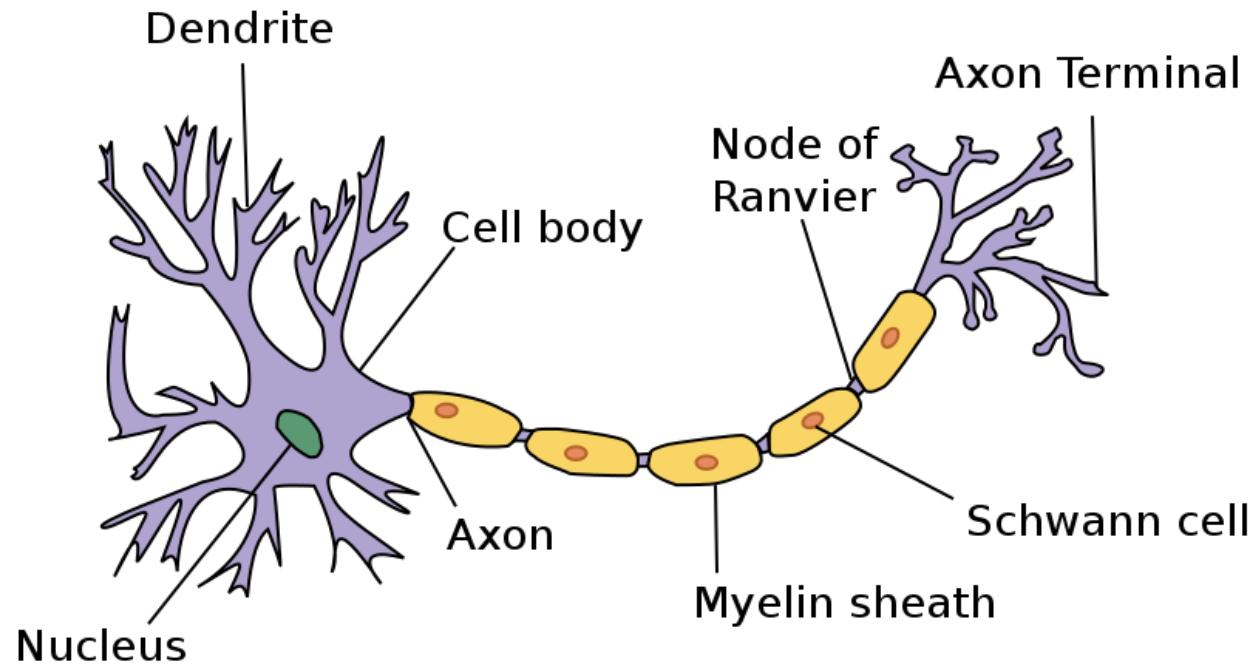
Review

❖ Multilayer Perceptron



Review

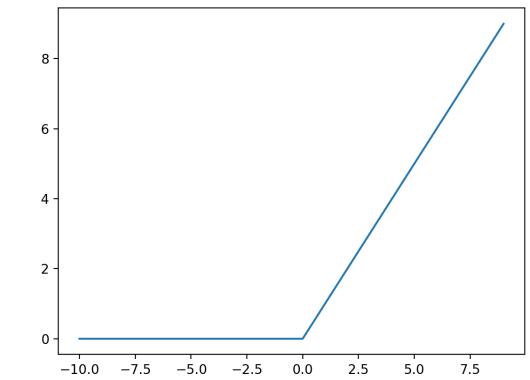
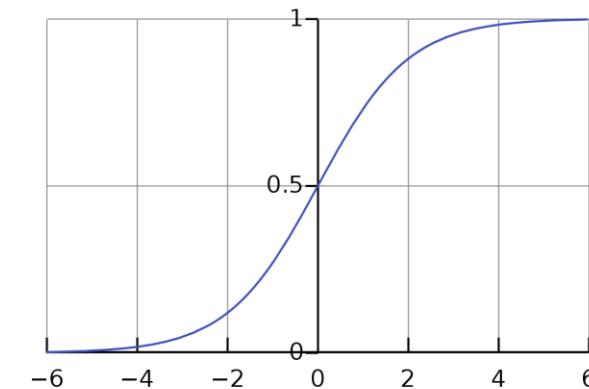
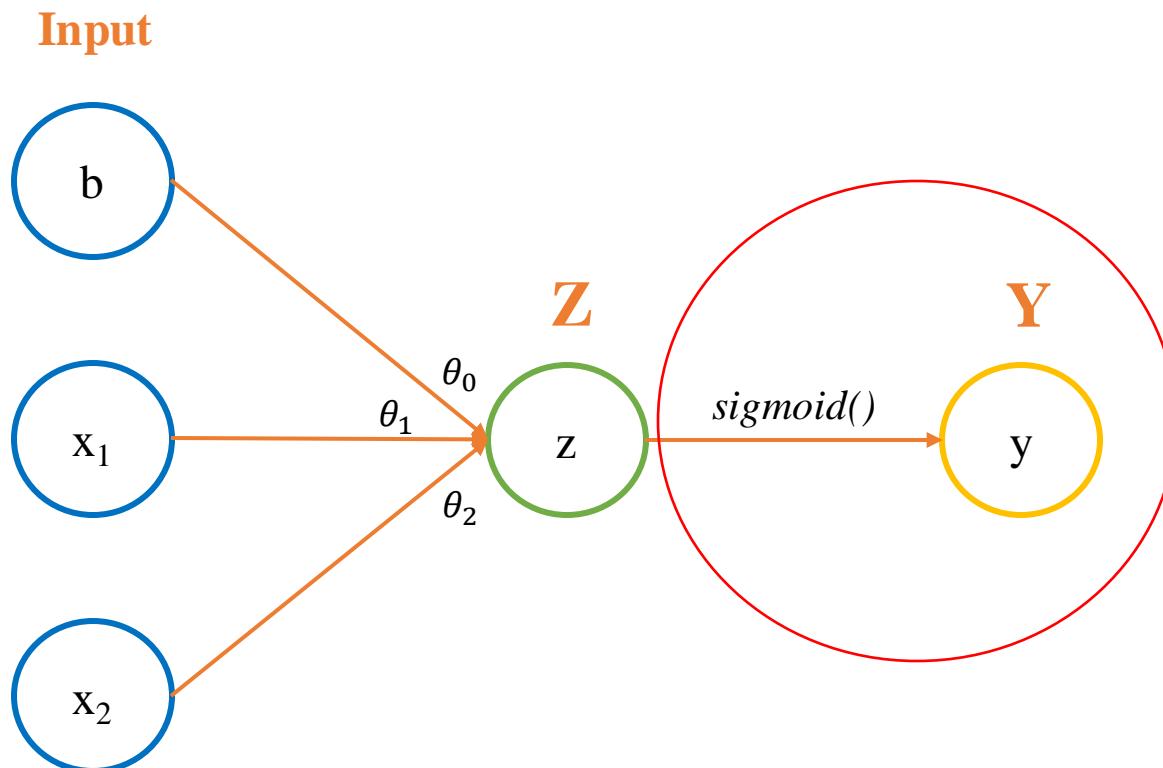
❖ Activation function



Activation Function: A function in neural networks that transforms the weighted sum of inputs into an output, introducing non-linearity and enabling the network to learn complex patterns in data.

Review

❖ Activation function



Activation functions:

- ❖ Final output of a node.
- ❖ Bring non-linearity to the network.

Review

❖ Activation function

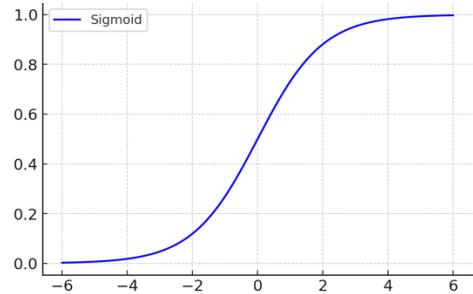
$\text{sigmoid}(x) = \frac{1}{1 + e^{-x}}$	2010 $\text{ReLU}(x) = \begin{cases} 0 & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$	2017 $\text{SELU}(x) = \begin{cases} \lambda x & \text{if } x \geq 0 \\ \lambda\alpha(e^x - 1) & \text{if } x < 0 \end{cases}$ $\lambda \approx 1.0507$ $\alpha \approx 1.6733$
$\tanh(x) = \frac{2}{1 + e^{-2x}} - 1$	2015 $\text{ELU}(x) = \begin{cases} \alpha(e^x - 1) & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$	
2001 $\text{softplus}(x) = \log(1 + e^x)$	2015 $\text{PReLU}(x) = \begin{cases} \alpha x & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$	2017 $\text{swish}(x) = x * \frac{1}{1 + e^{-x}}$

Some famous activation functions.

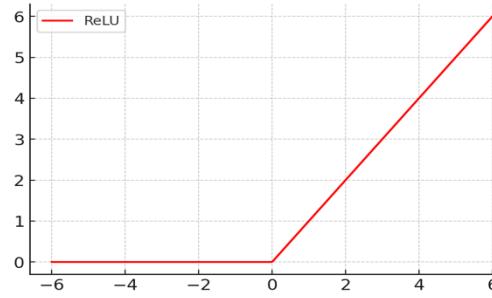
Review

❖ Activation function: Graphs

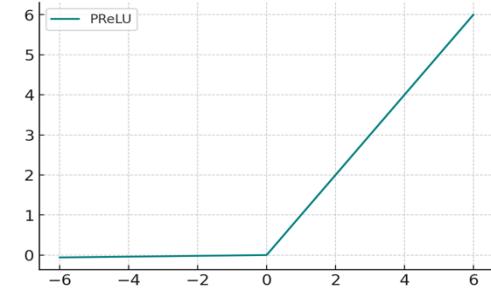
Sigmoid Function



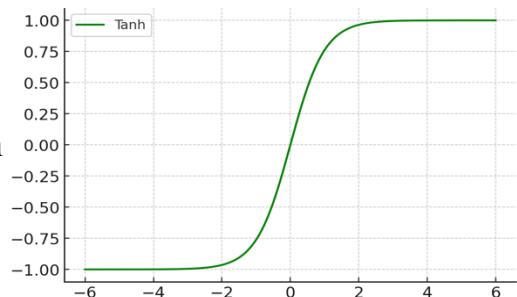
ReLU Function



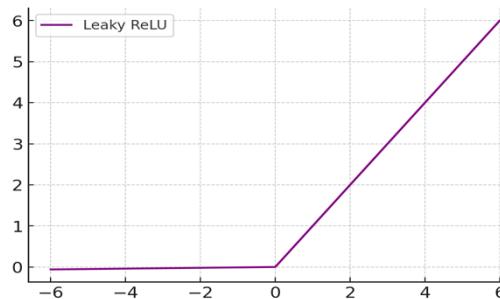
PReLU Function



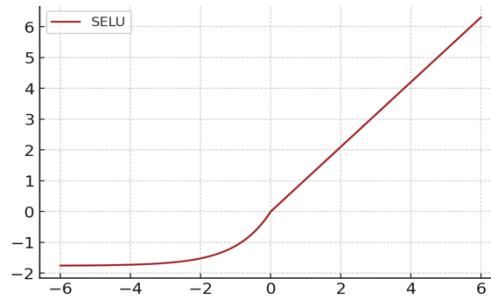
Tanh Function



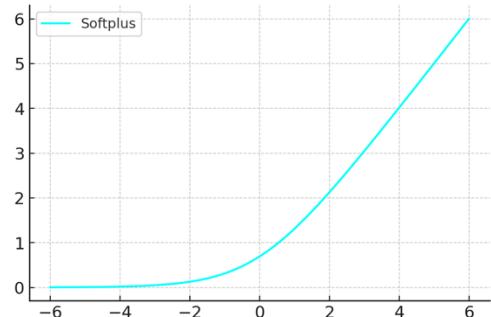
Leaky ReLU Function



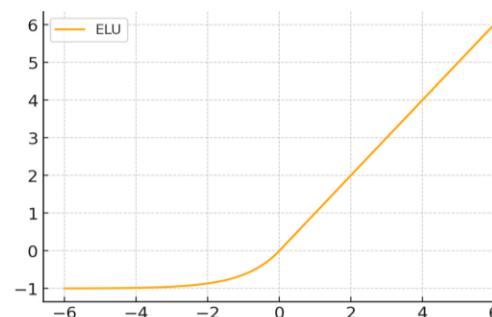
SELU Function



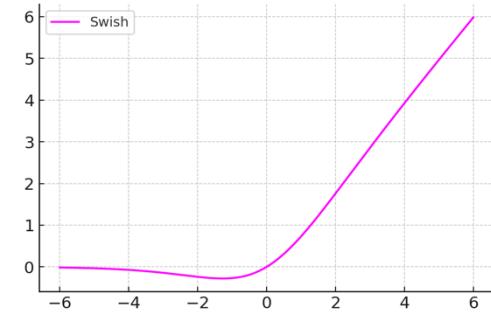
Softplus Function



ELU Function

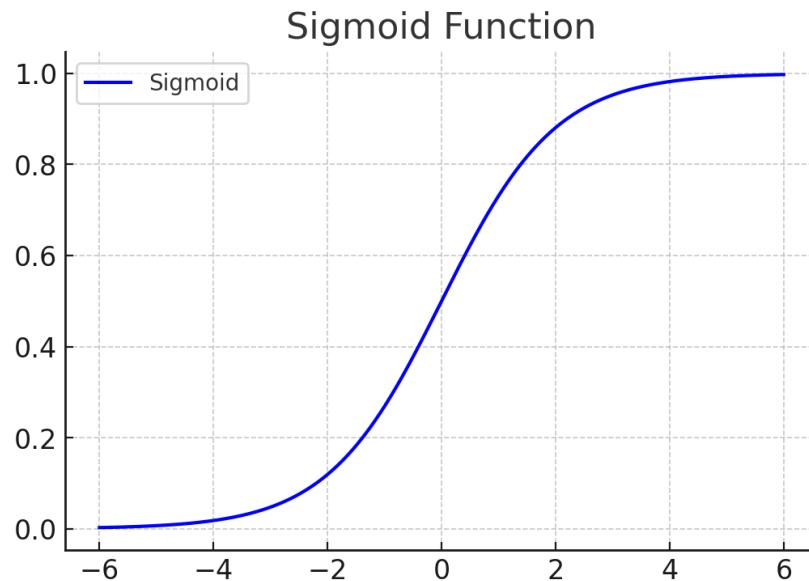


Swish Function



Review

❖ Activation function: Sigmoid

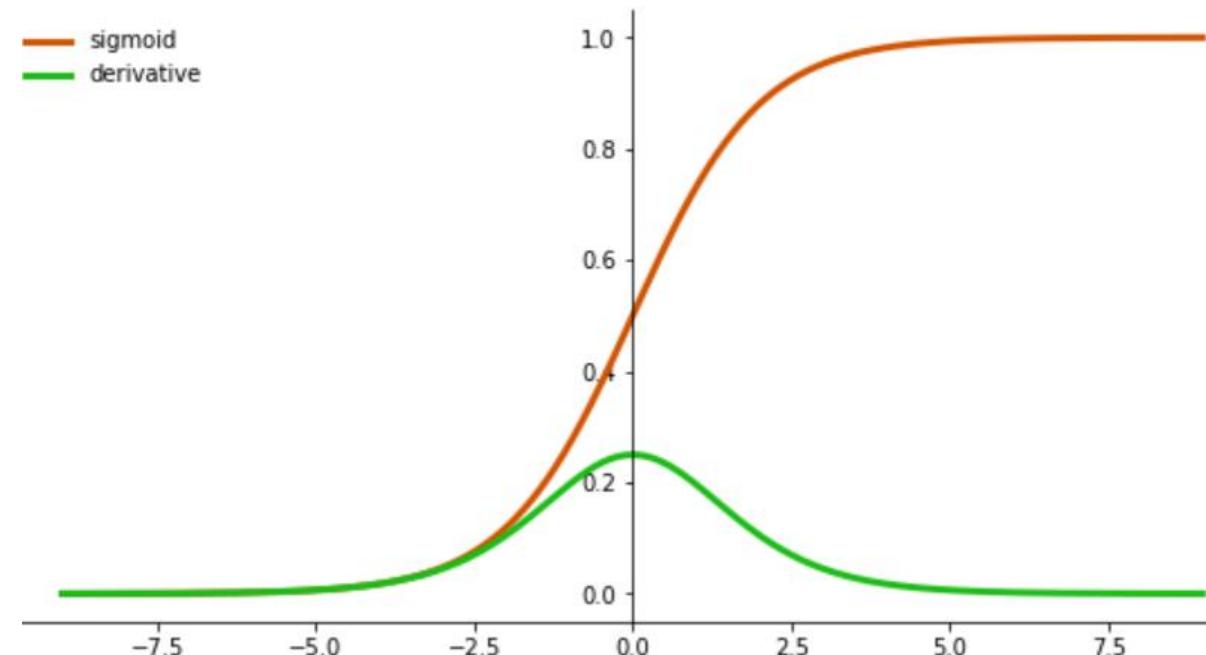


Domain: $(-\infty, +\infty)$,
Range: $(0, 1)$

$$\text{sigmoid}(x) = \frac{1}{1 + e^{-x}}$$

$$\text{sigmoid}'(x) = \text{sigmoid}(x)(1 - \text{sigmoid}(x))$$

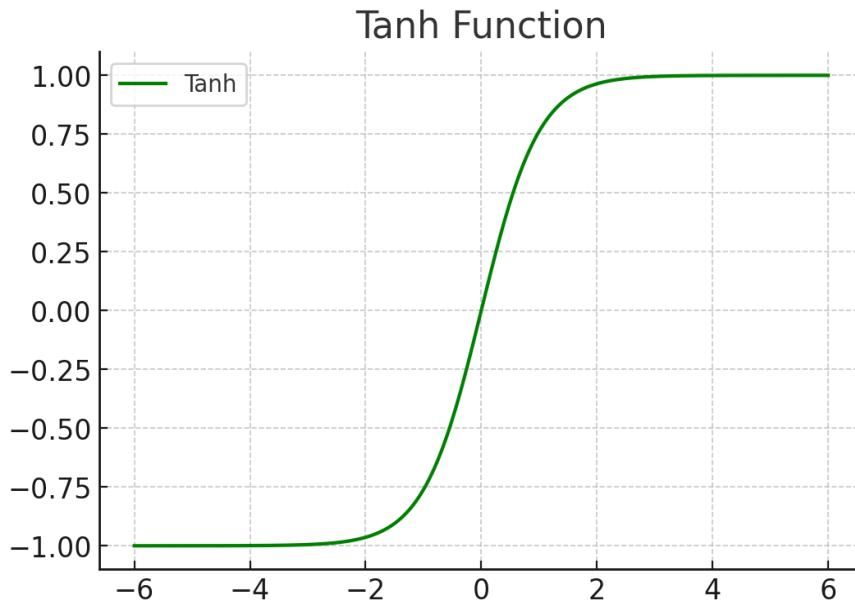
— sigmoid
— derivative



Review

❖ Activation function: Tanh

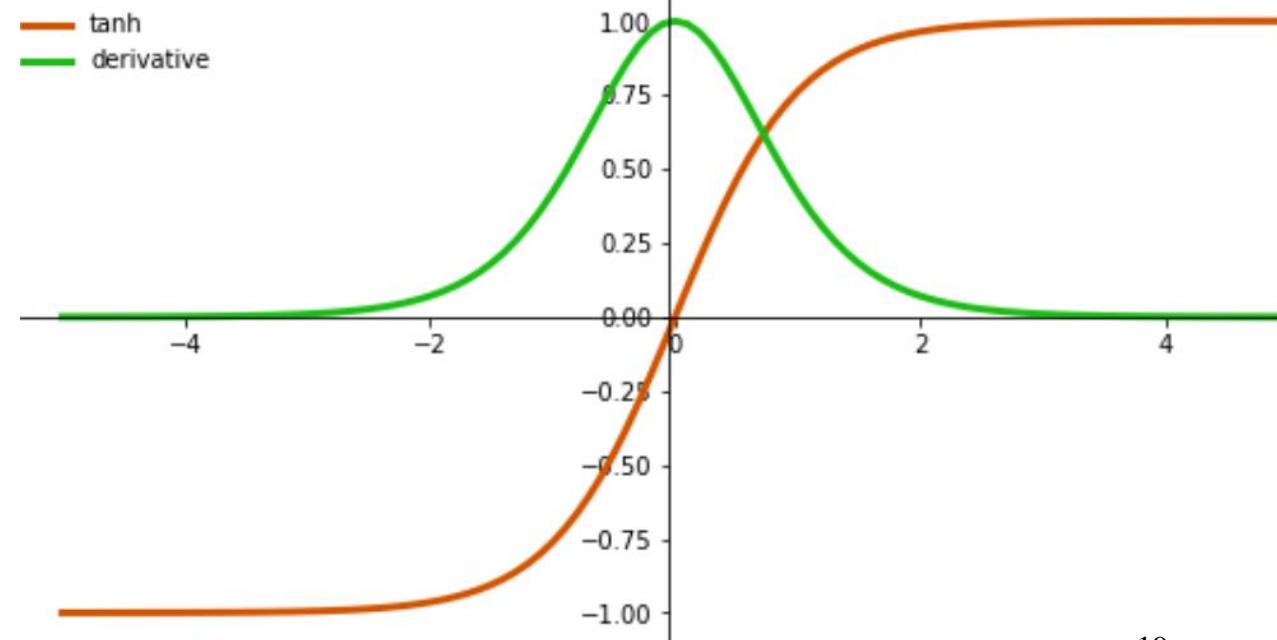
$$\tanh(x) = \frac{2}{1 + e^{-2x}} - 1$$



Domain: $(-\infty, +\infty)$,
Range: $(-1, 1)$

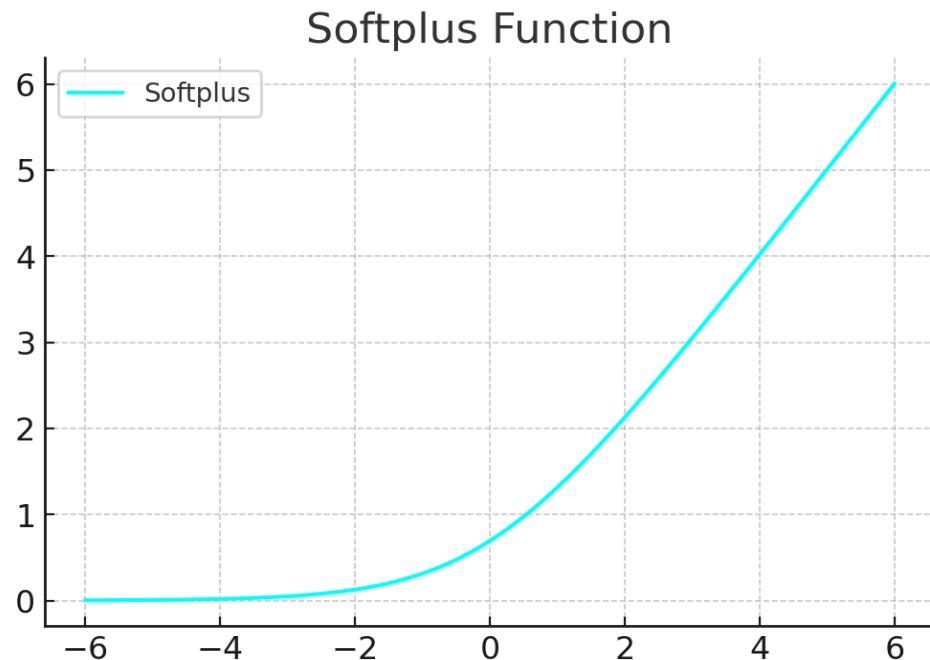
$$\begin{aligned}\tanh(x) &= \frac{e^x - e^{-x}}{e^x + e^{-x}} \\ &= \frac{2}{1 + e^{-2x}} - 1 \\ &= 1 - \frac{2}{e^{2x} + 1}\end{aligned}$$

$$\tanh'(x) = 1 - \tanh^2(x)$$



Review

❖ Activation function: Softplus



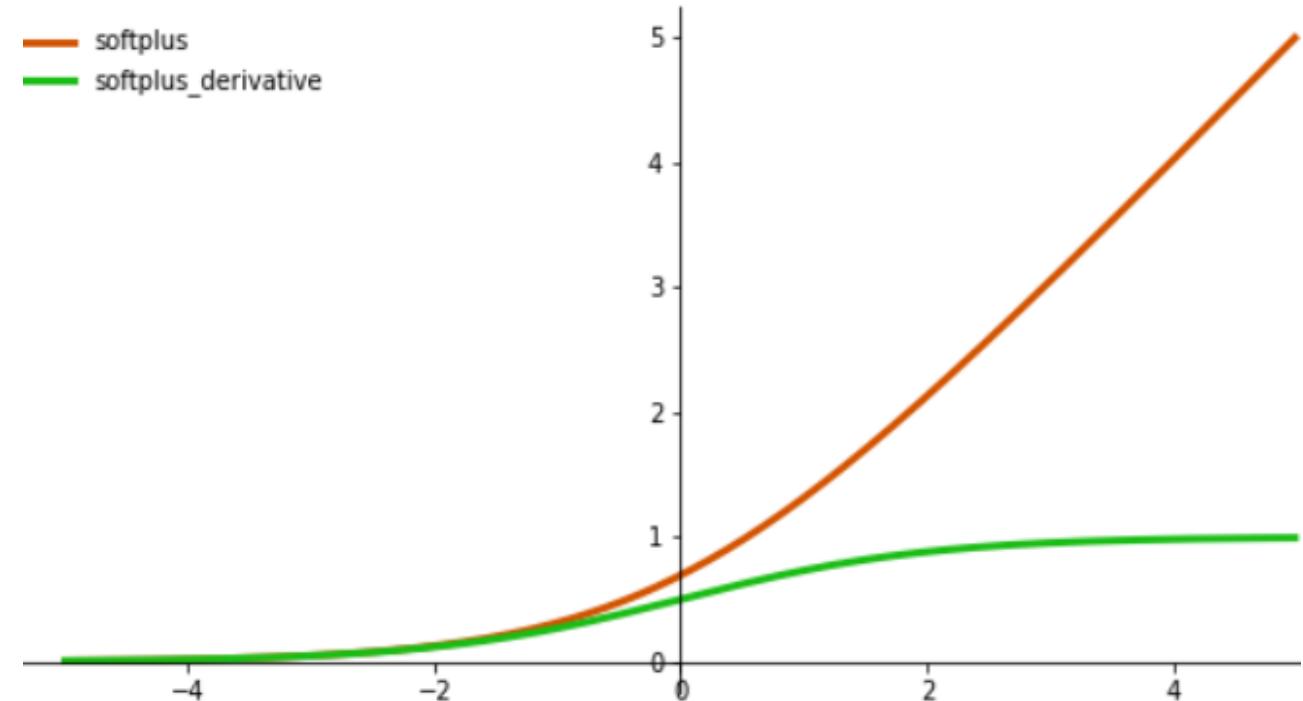
Domain: $(-\infty, +\infty)$,
Range: $(0, +\infty)$

2001

$$\text{softplus}(x) = \log(1 + e^x)$$

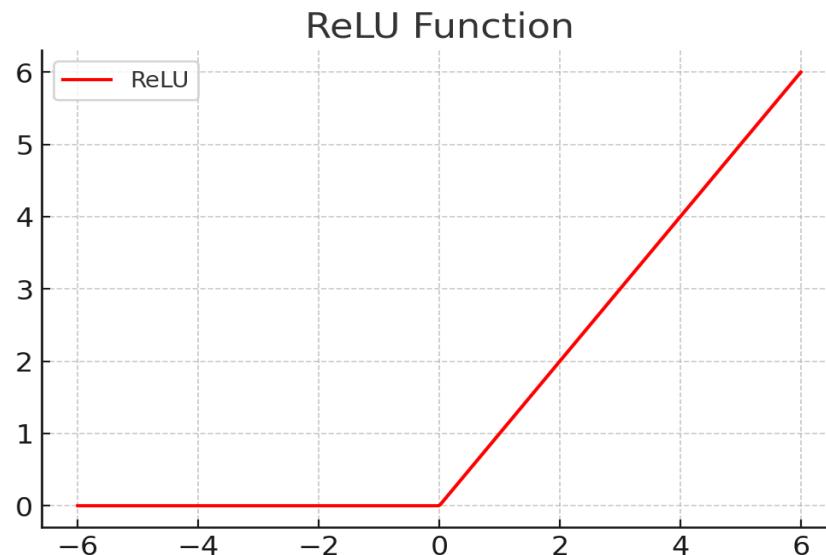
$$\text{softplus}'(x) = \frac{1}{1 + e^{-x}}$$

— softplus
— softplus_derivative



Review

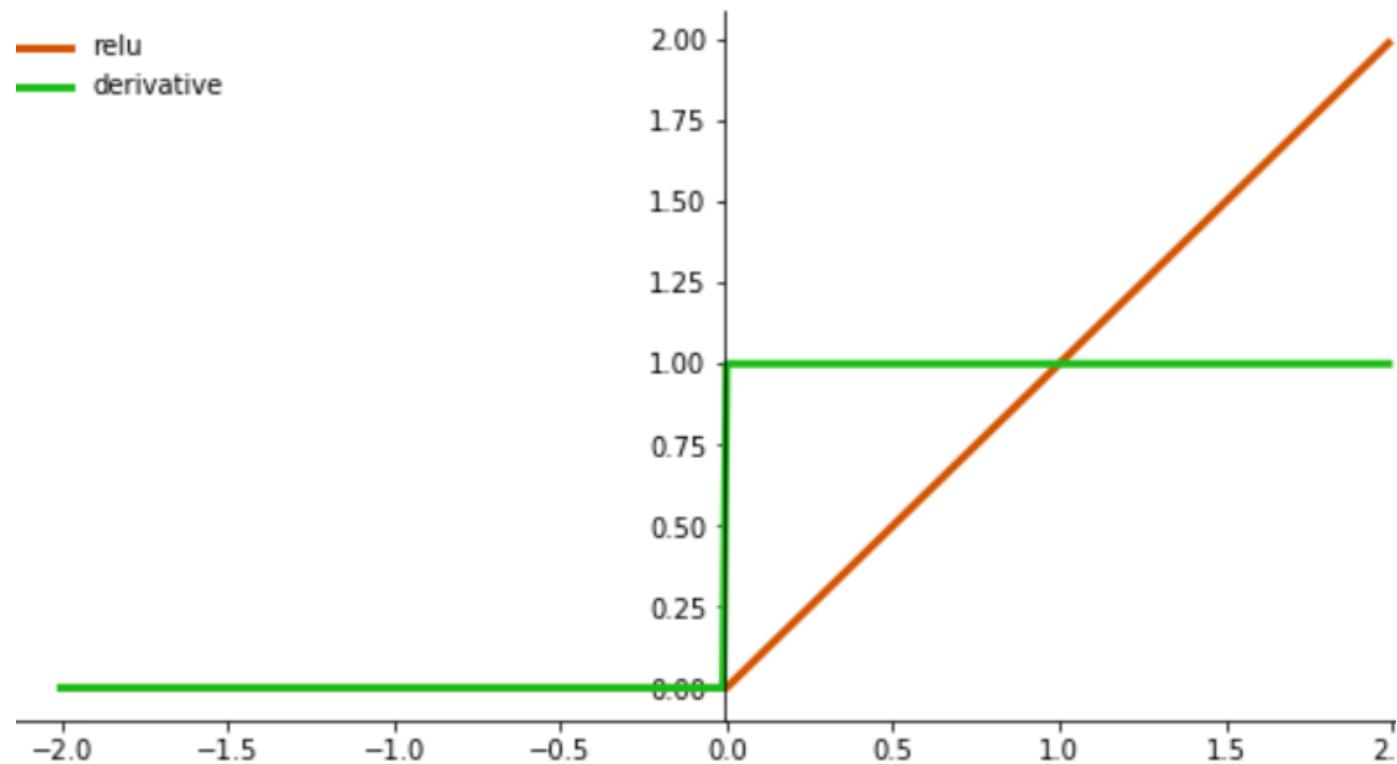
❖ Activation function: ReLU



2010

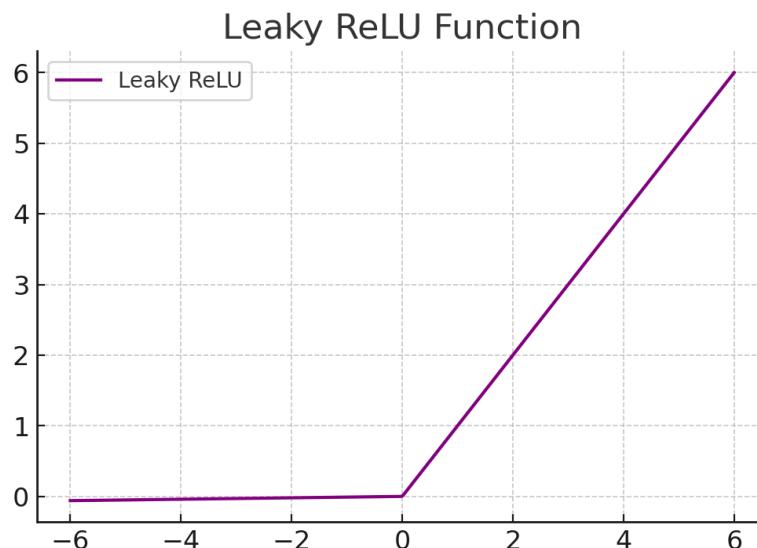
$$\text{ReLU}(x) = \begin{cases} 0 & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$$

$$\text{ReLU}'(x) = \begin{cases} 0 & \text{if } x \leq 0 \\ 1 & \text{if } x > 0 \end{cases}$$

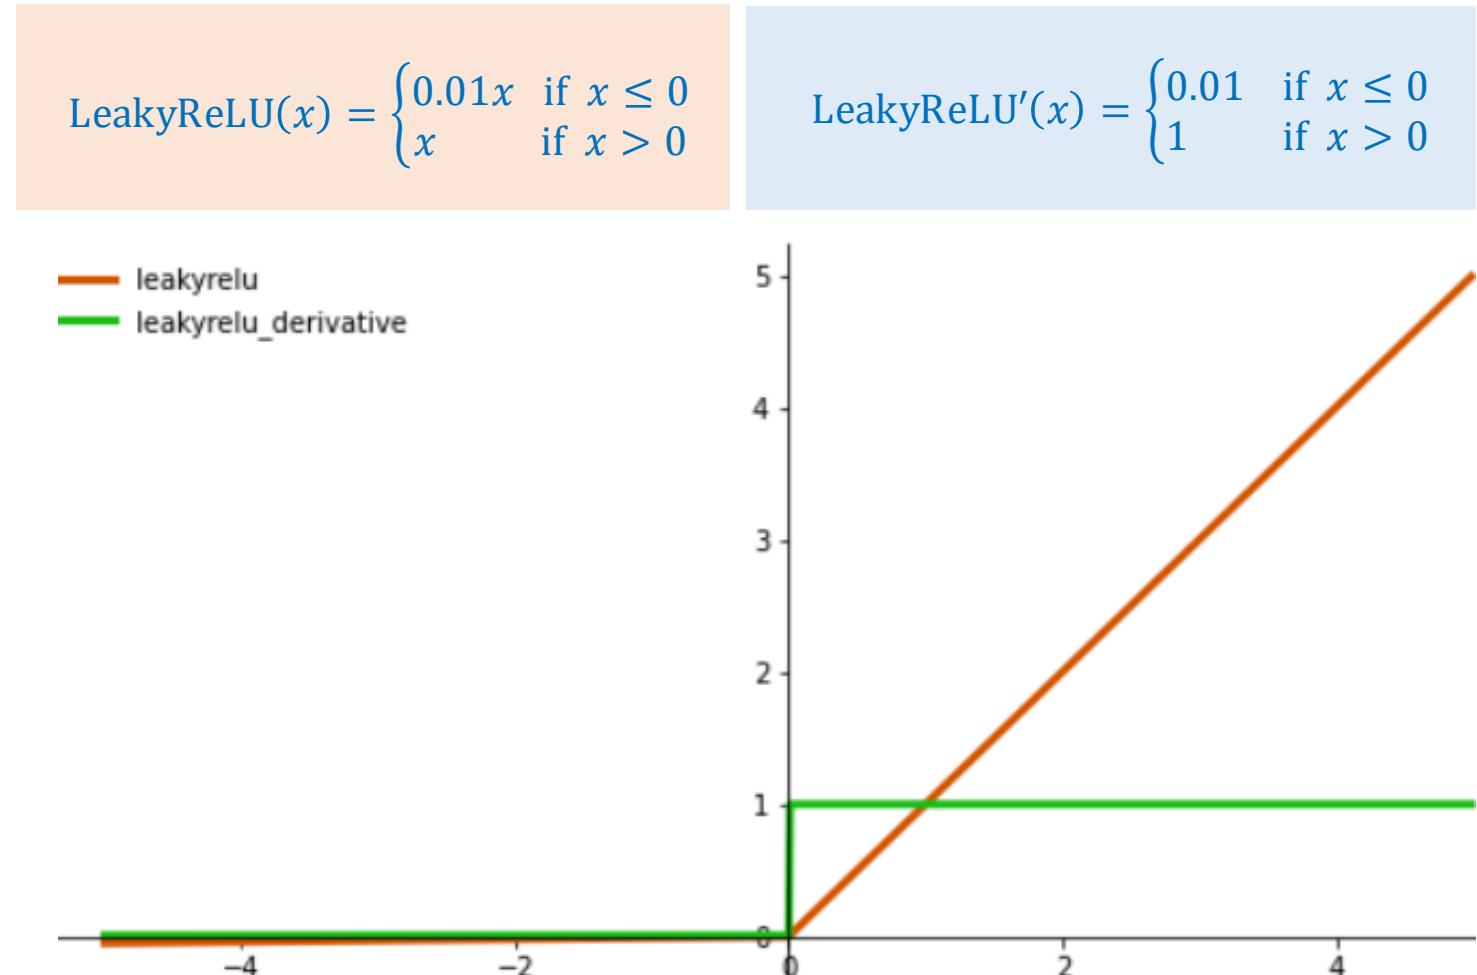


Review

❖ Activation function: LeakyReLU

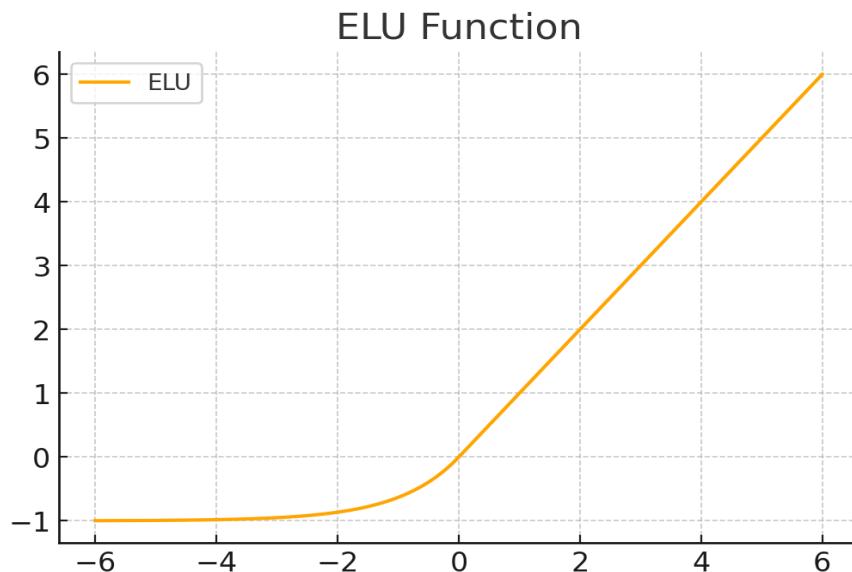


Domain: $(-\infty, +\infty)$,
Range: $(-\infty, +\infty)$



Review

❖ Activation function: ELU

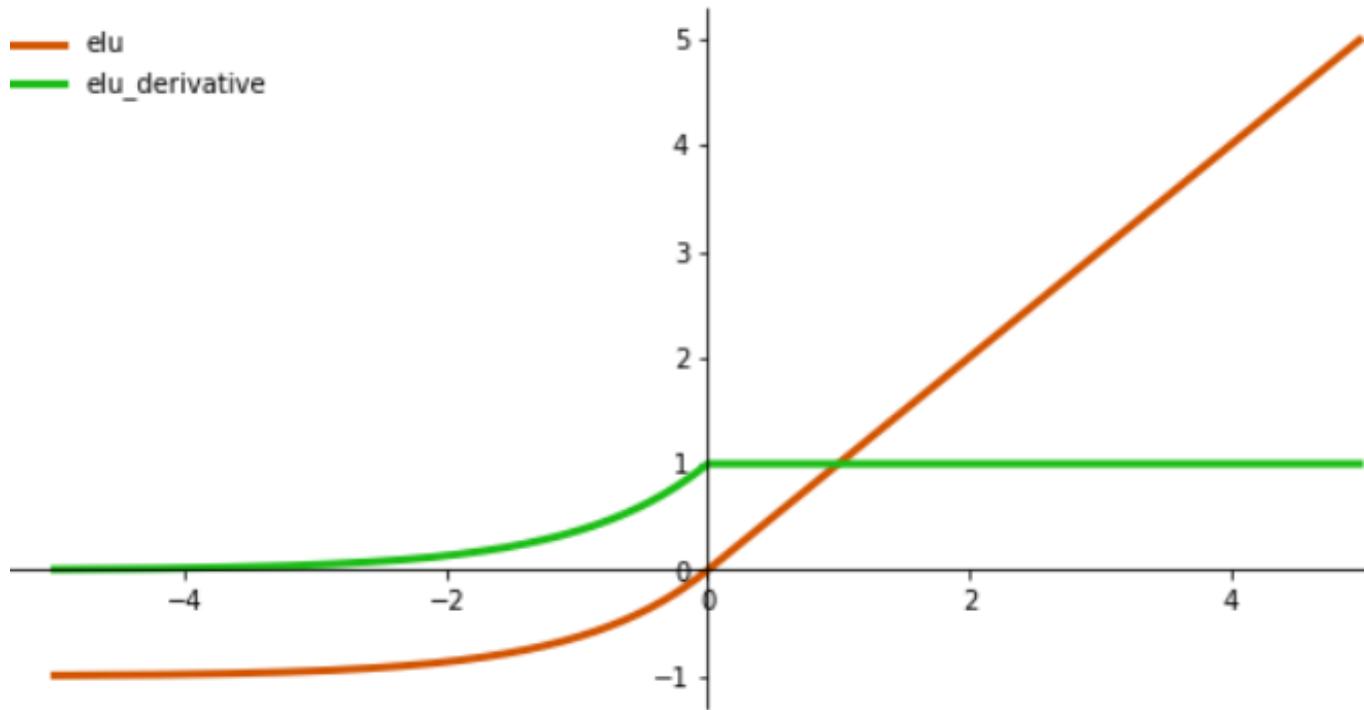


Domain: $(-\infty, +\infty)$,
Range: $(-\alpha, +\infty)$ with $\alpha > 0$

2015

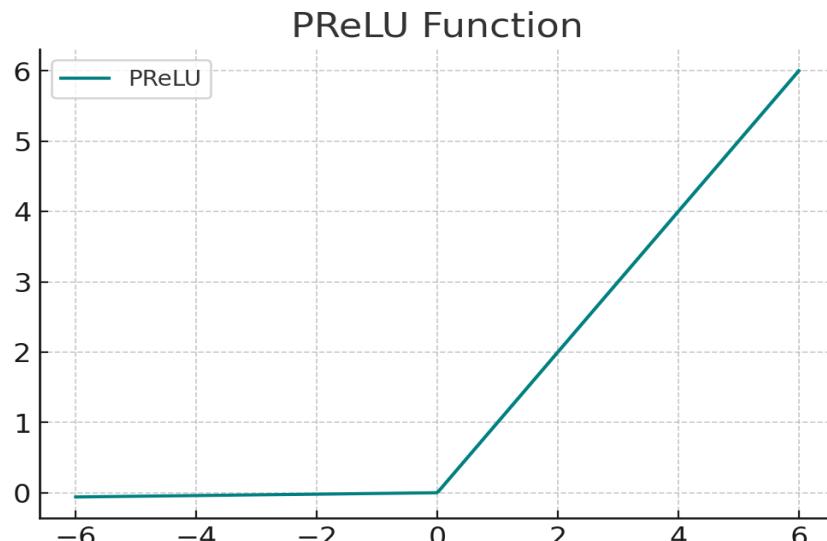
$$\text{ELU}(x) = \begin{cases} \alpha(e^x - 1) & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$$

elu
elu_derivative



Review

❖ Activation function: PReLU

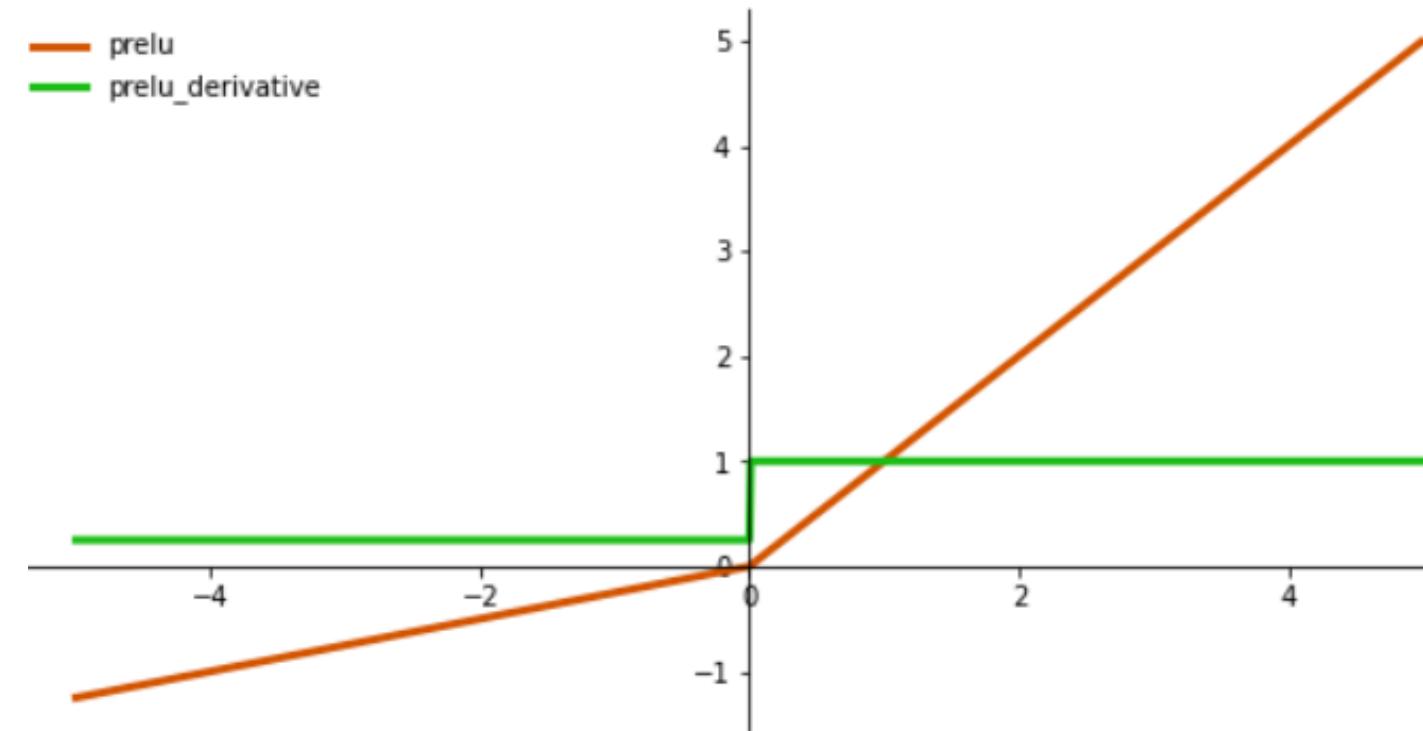


2015

$$\text{PReLU}(x) = \begin{cases} ax & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$$

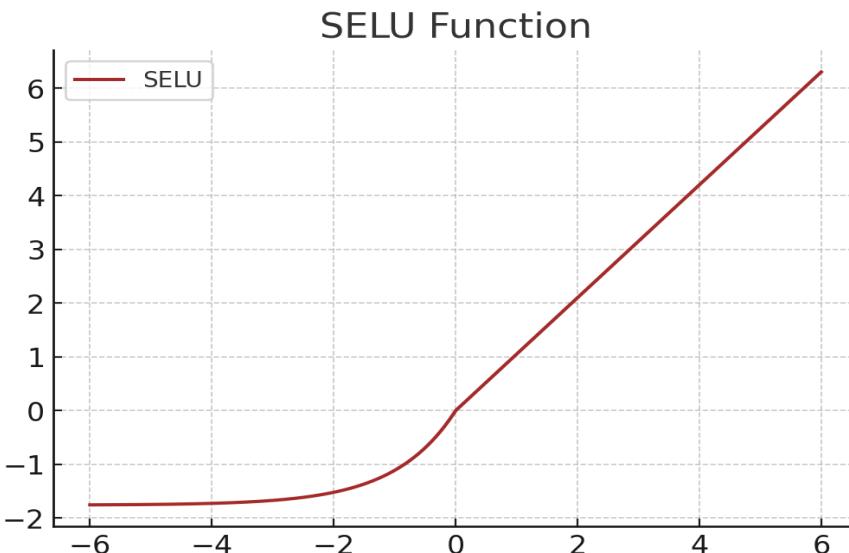
$$\text{PReLU}'(x) = \begin{cases} \alpha & \text{if } x \leq 0 \\ 1 & \text{if } x > 0 \end{cases}$$

— prelu
— prelu_derivative



Review

❖ Activation function: SELU

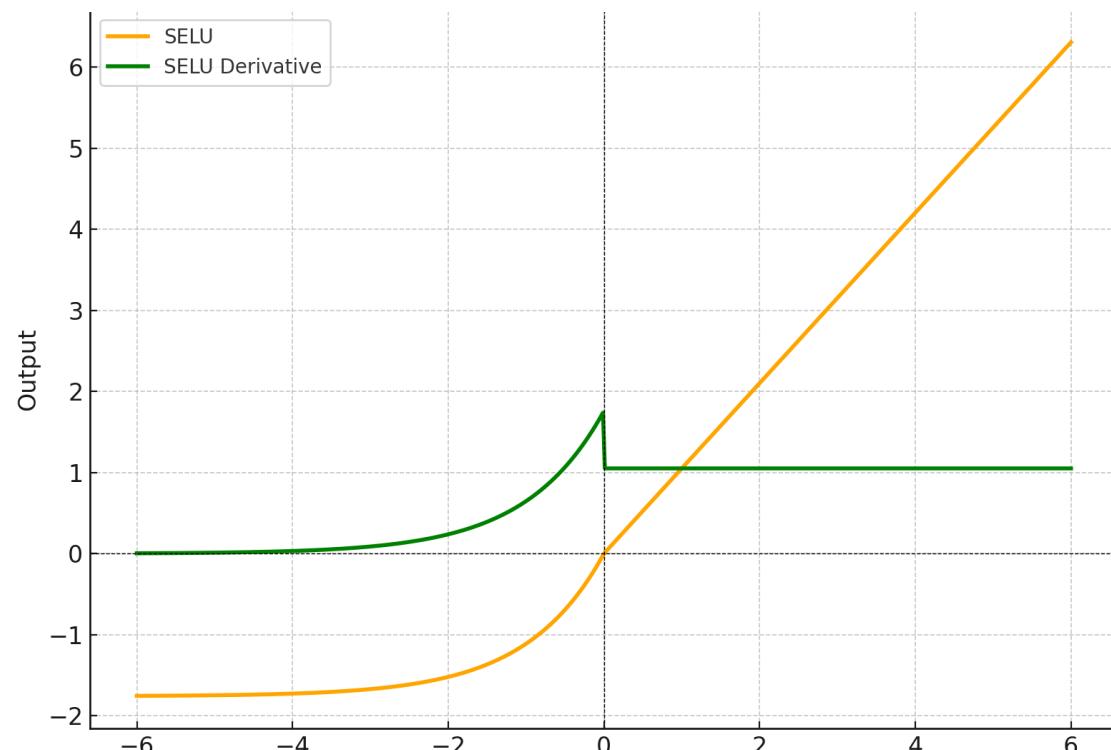


Domain: $(-\infty, +\infty)$,
Range: $(-\lambda\alpha, +\infty)$
with $\lambda \approx 1.0507$, $\alpha \approx 1.6733$

2017

$$\text{SELU}(x) = \begin{cases} \lambda x & \text{if } x \geq 0 \\ \lambda\alpha(e^x - 1) & \text{if } x < 0 \end{cases}$$
$$\lambda \approx 1.0507$$
$$\alpha \approx 1.6733$$

$$\text{SELU}'(x) = \begin{cases} \lambda & \text{if } x \leq 0 \\ \lambda\alpha e^x & \text{if } x > 0 \end{cases}$$

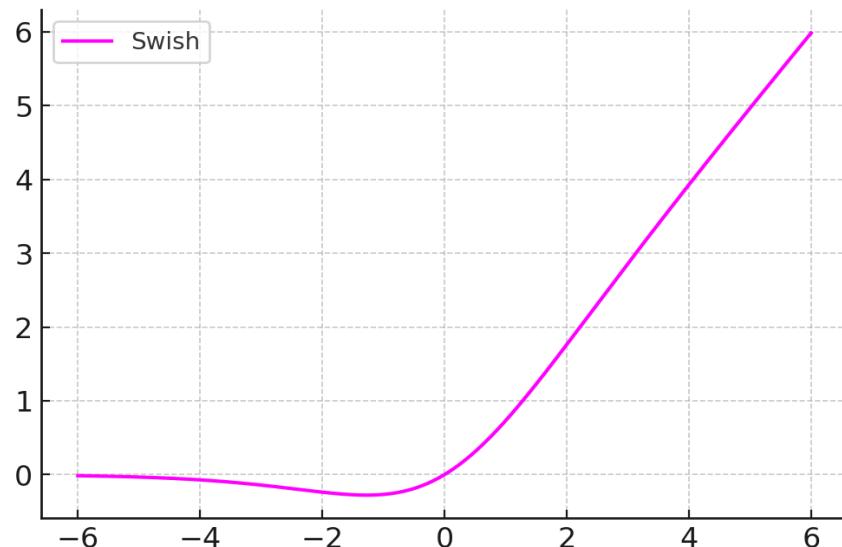


Review

❖ Activation function: Swish

2017

$$swish(x) = x * \frac{1}{1 + e^{-x}}$$

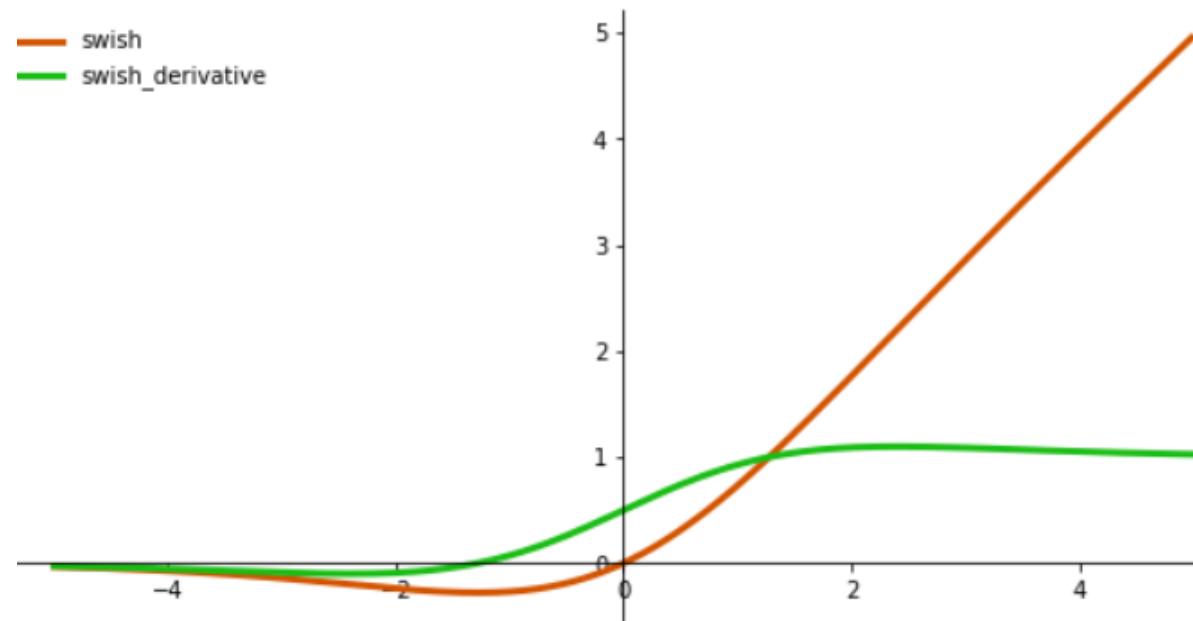


Domain: $(-\infty, +\infty)$,
Range: $(-\infty, +\infty)$

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$
$$swish(x) = \frac{x}{1 + e^{-x}} = x \sigma(x)$$

$$swish'(x) = swish(x) + \sigma(x)(1 - swish(x))$$

— swish
— swish_derivative



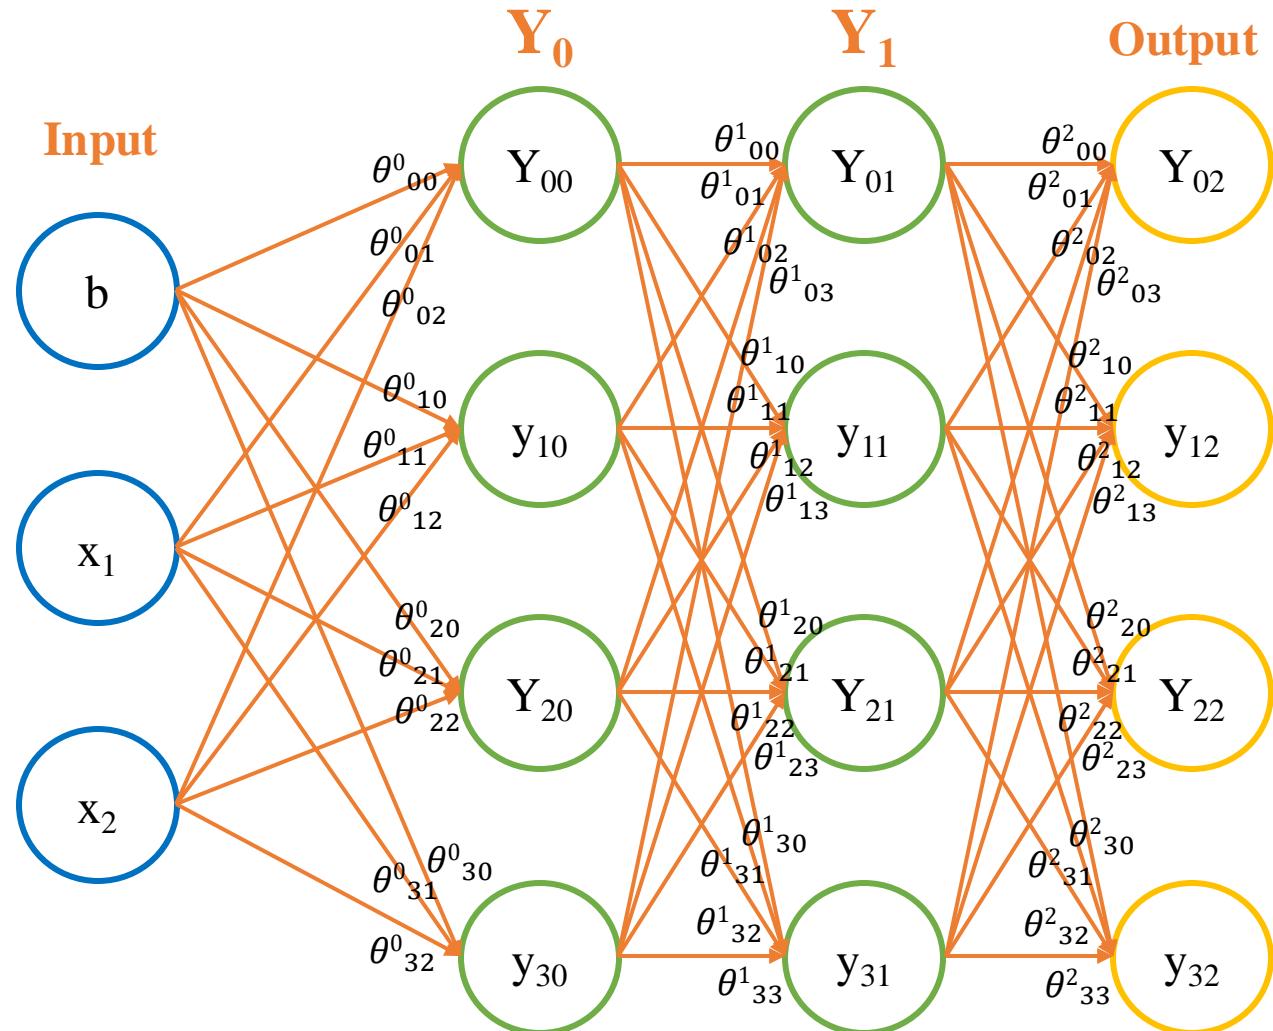


Review

❖ Activation function

Function	Pros	Cons	Applications
Sigmoid	Good for probabilities. Simple.	Vanishing gradient. Not centered.	Binary classification (e.g., logistic regression).
Tanh	Zero-centered. Good for RNNs.	Vanishing gradient.	RNN hidden layers (e.g., sentiment analysis, language modeling).
ReLU	Simple and efficient. No saturation.	Dying ReLU problem.	Default for CNNs, MLPs.
Softplus	Smooth ReLU approximation. Continuous and smooth.	Slower than ReLU.	Alternative to ReLU for smooth outputs in regression tasks
Leaky ReLU	Small slope for $x < 0$. Solves Dying ReLU. Simple.	Small bias in outputs.	Networks with deeper layers (e.g., GANs like DCGAN for generating images).
PReLU	Learnable slope for $x < 0$. Adaptive activation.	Higher computation cost.	Advanced deep networks (e.g., face recognition in VGGFace models).
ELU	Smooth. Zero-centered. Reduces vanishing gradient.	Slower than ReLU.	Robust networks (e.g., deep networks for speech recognition).
SELU	Self-normalizing. Stable training for deep layers.	Needs specific normalization.	Self-normalizing networks (SNNs) (e.g., tabular data in structured prediction).
Swish	Smooth. Non-monotonic. Improves deep model accuracy.	Computationally expensive.	State-of-the-art architectures (e.g., EfficientNet for image classification, Transformers for NLP tasks).

❖ Activation function

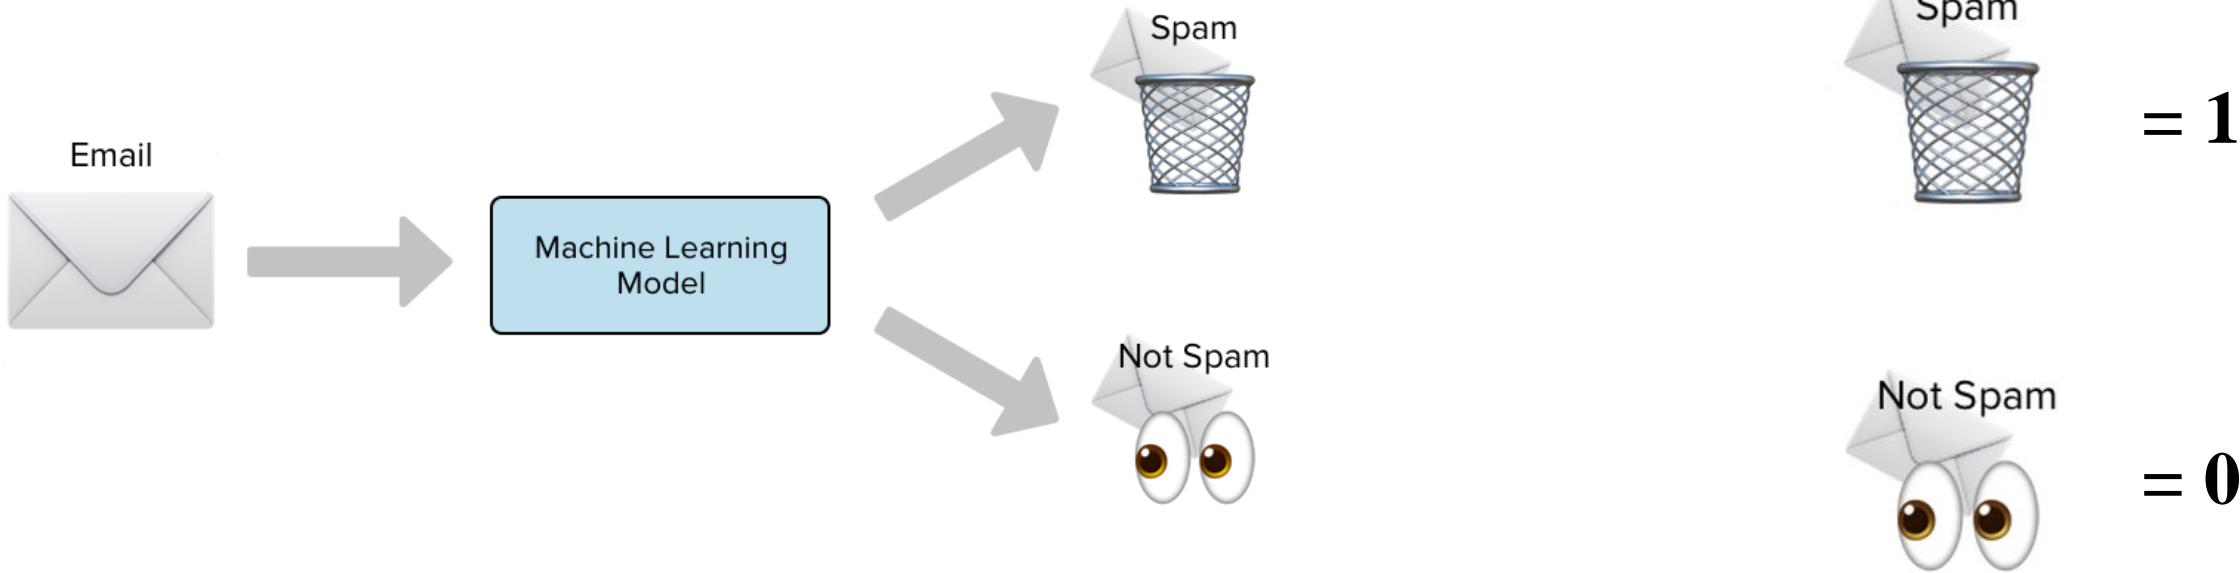


Question: Which activation function should we use?

Review

❖ Case study 1

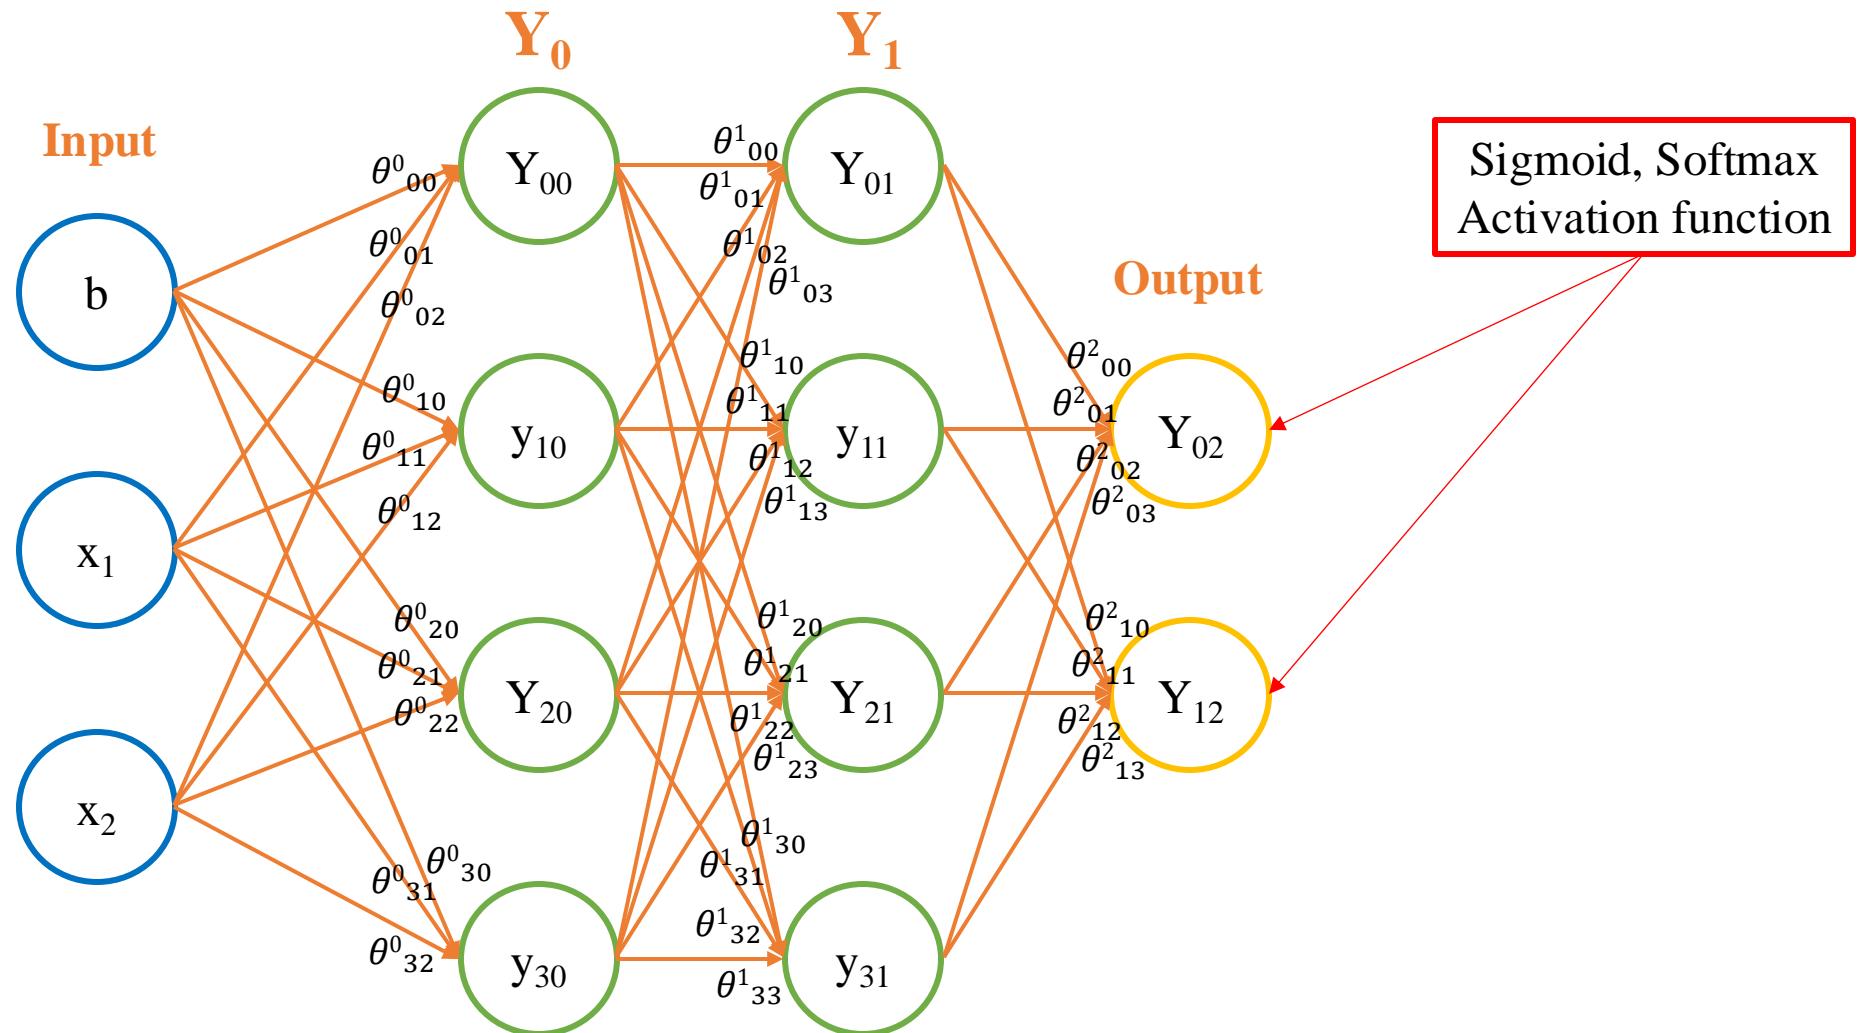
Problem output:
categorial value



We could use Sigmoid,
Tanh, Softmax... for
Output Layer.

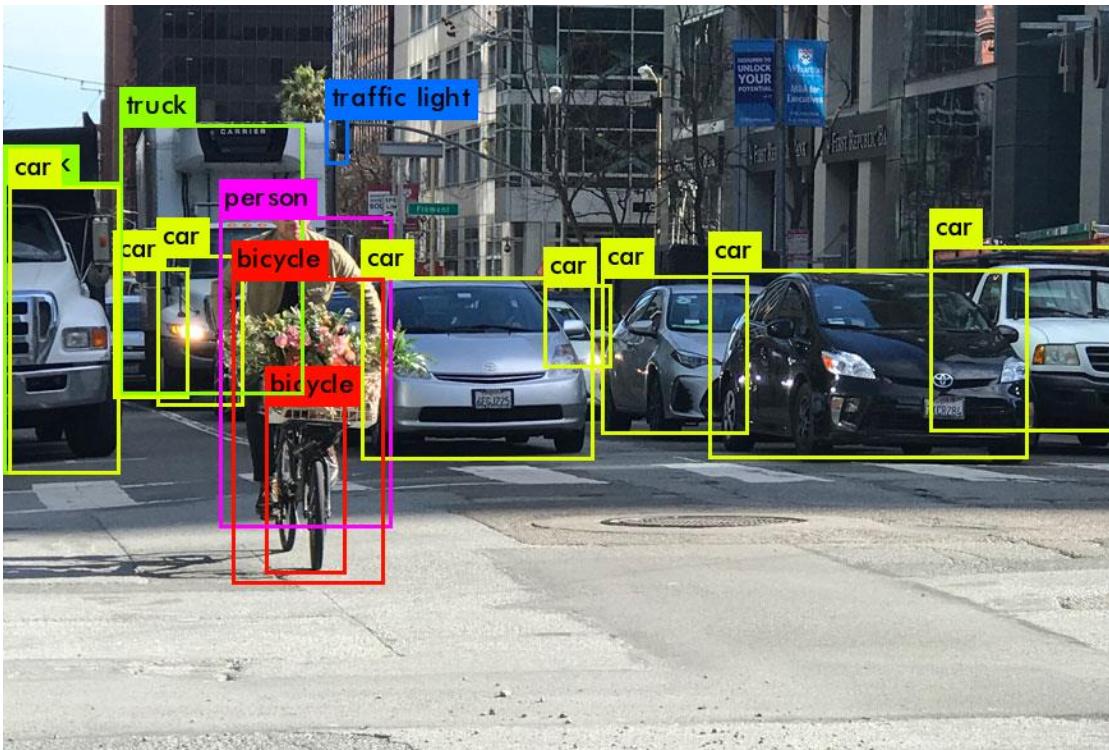
Review

❖ Case study 1



Review

❖ Case study 2



(x_left, y_left) width

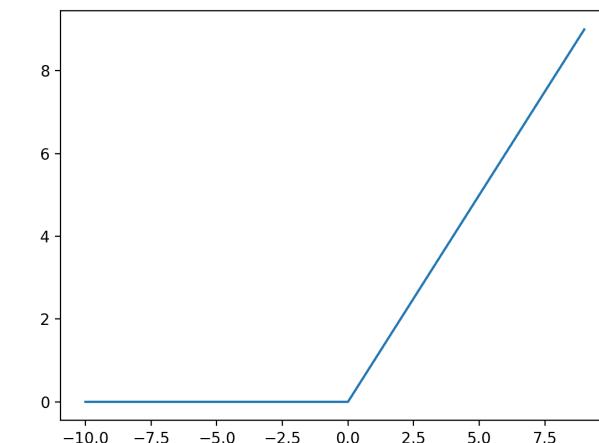
height



(x_right, y_right)

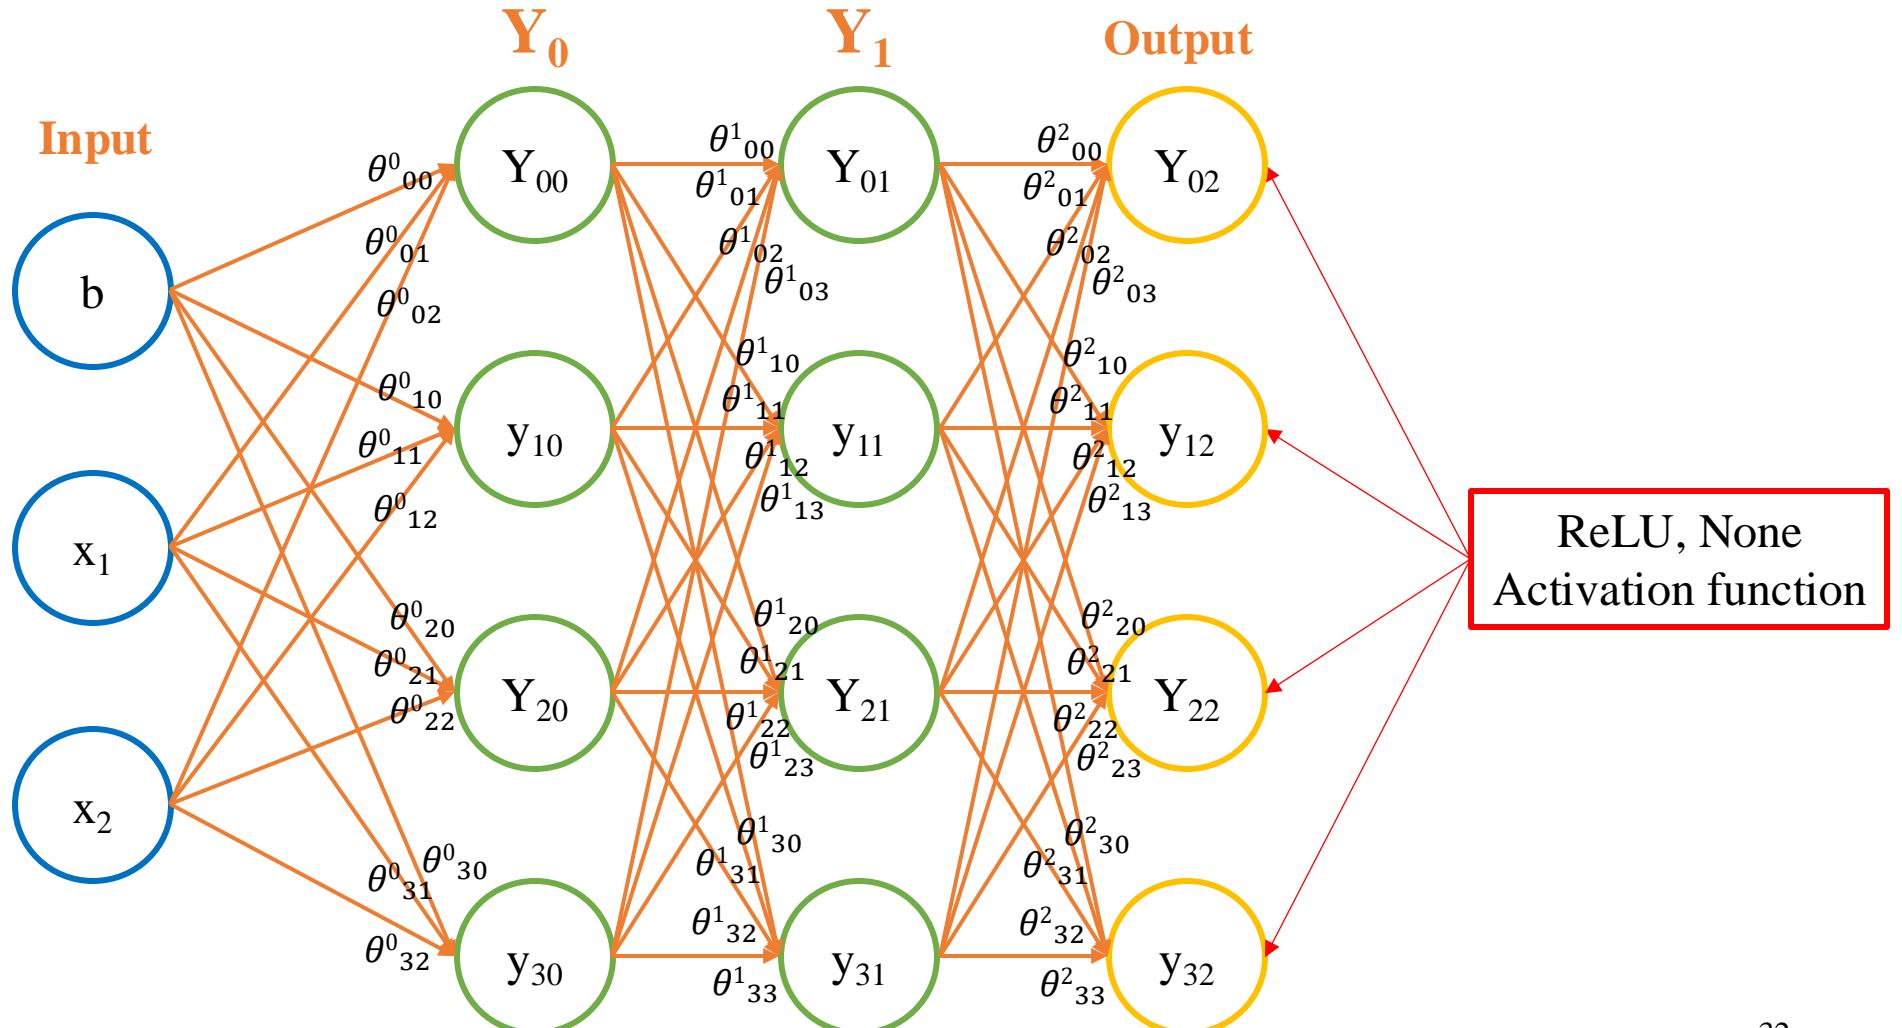
We could use
ReLU, None...
for Output Layer.

Object BB
Coordinate: positive
floating value.



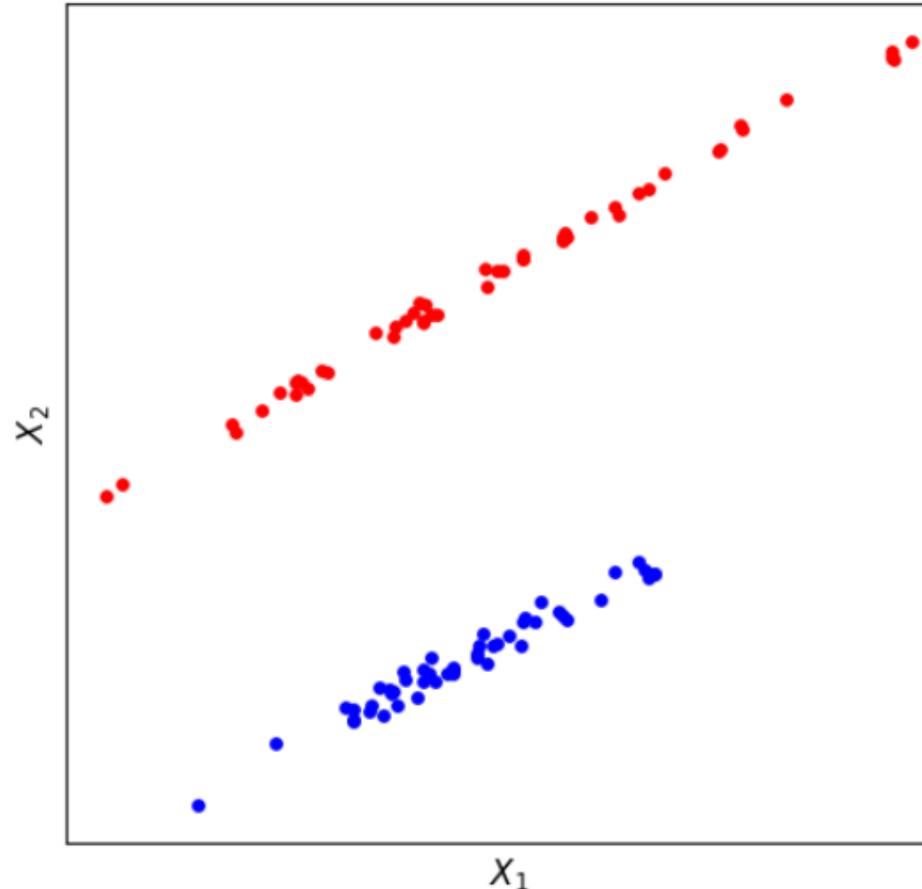
Review

❖ Case study 2



Review

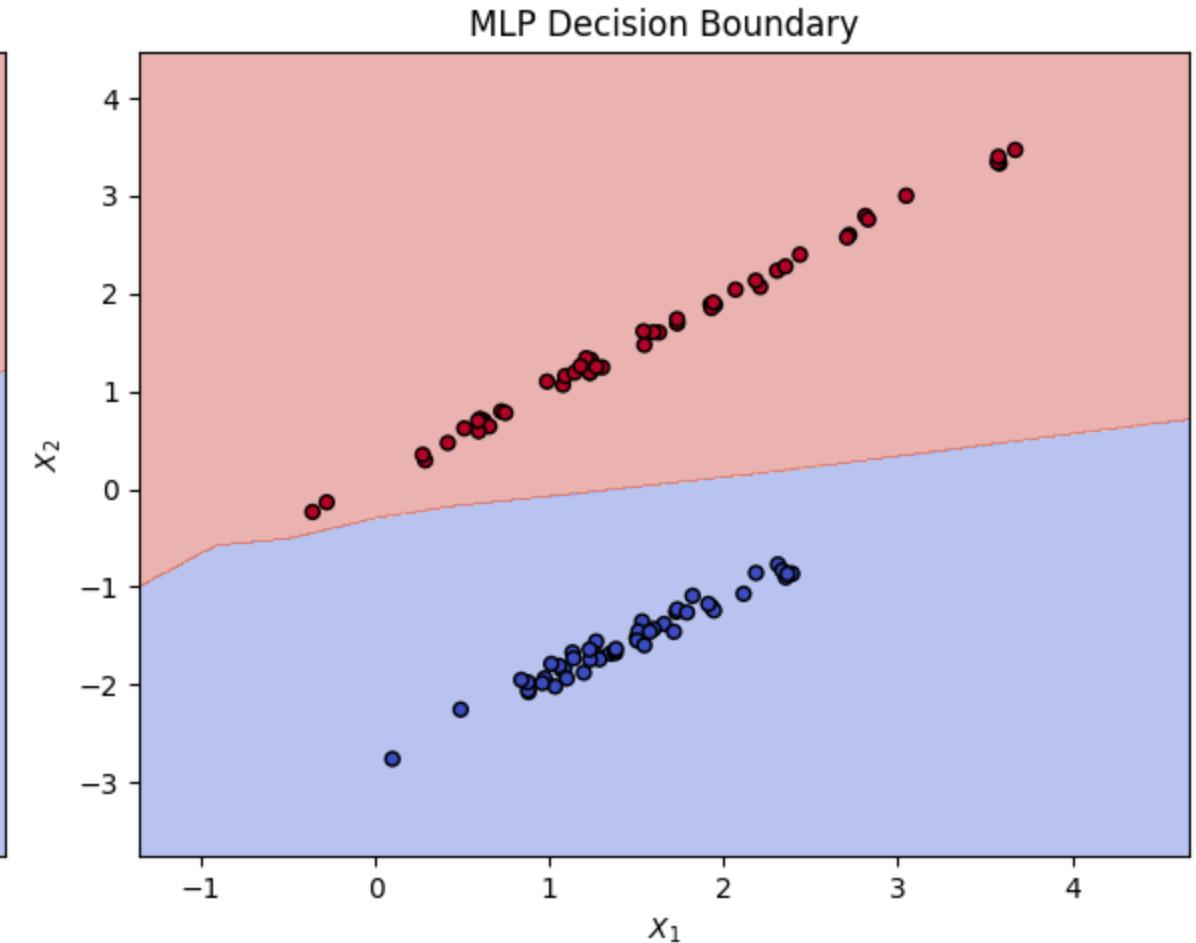
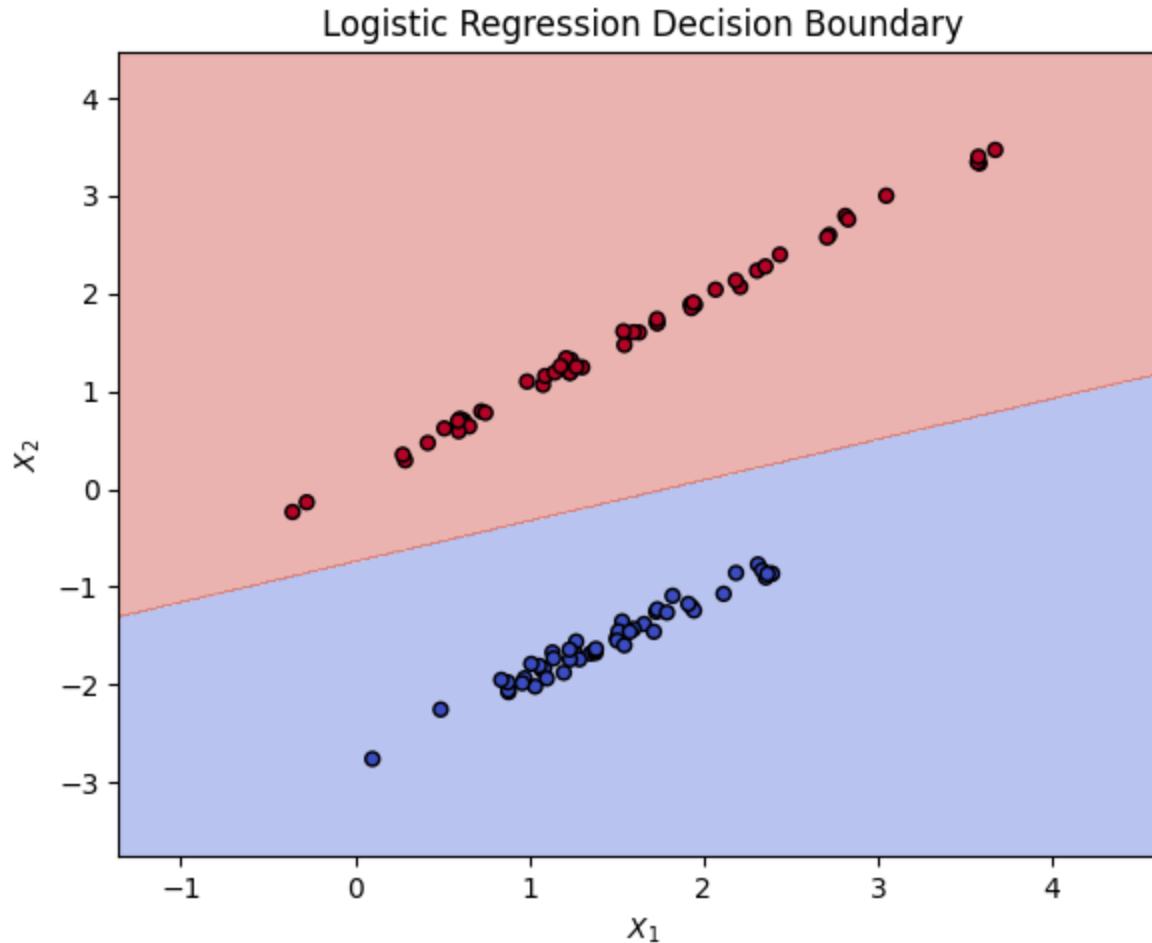
❖ Why we need MLP?



A simple (linear) classification dataset.

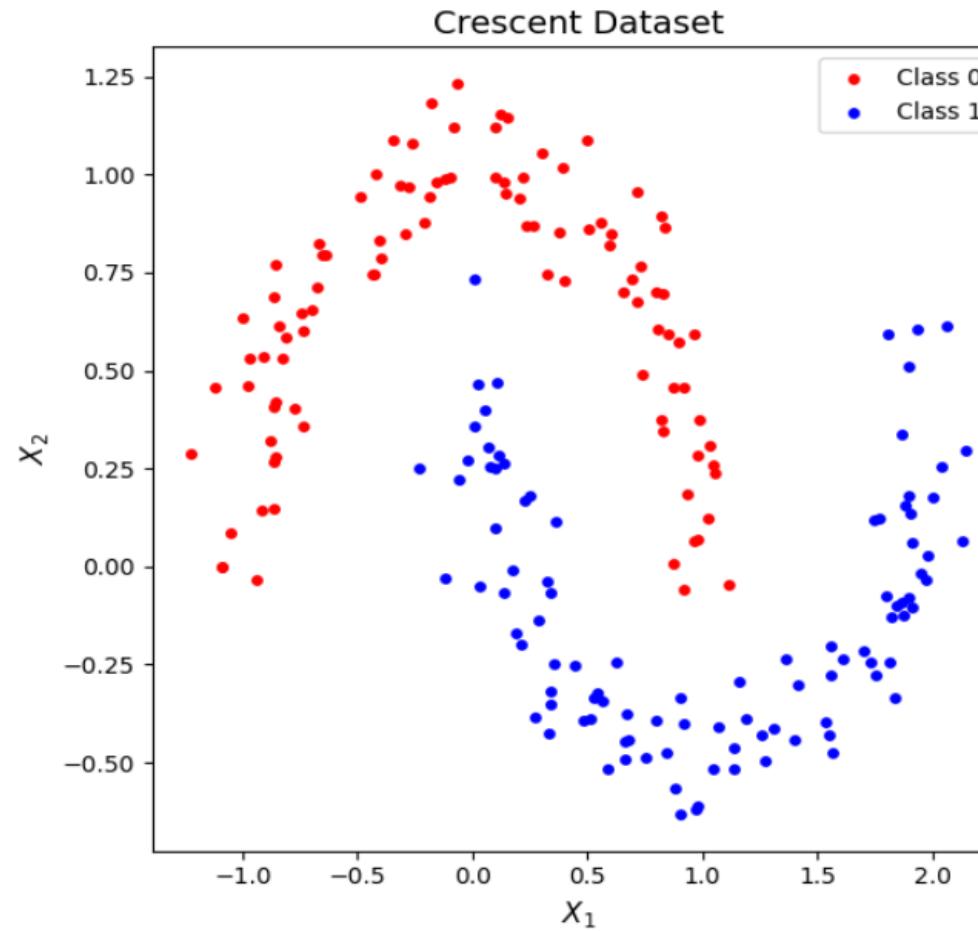
Review

❖ Why we need MLP?



Review

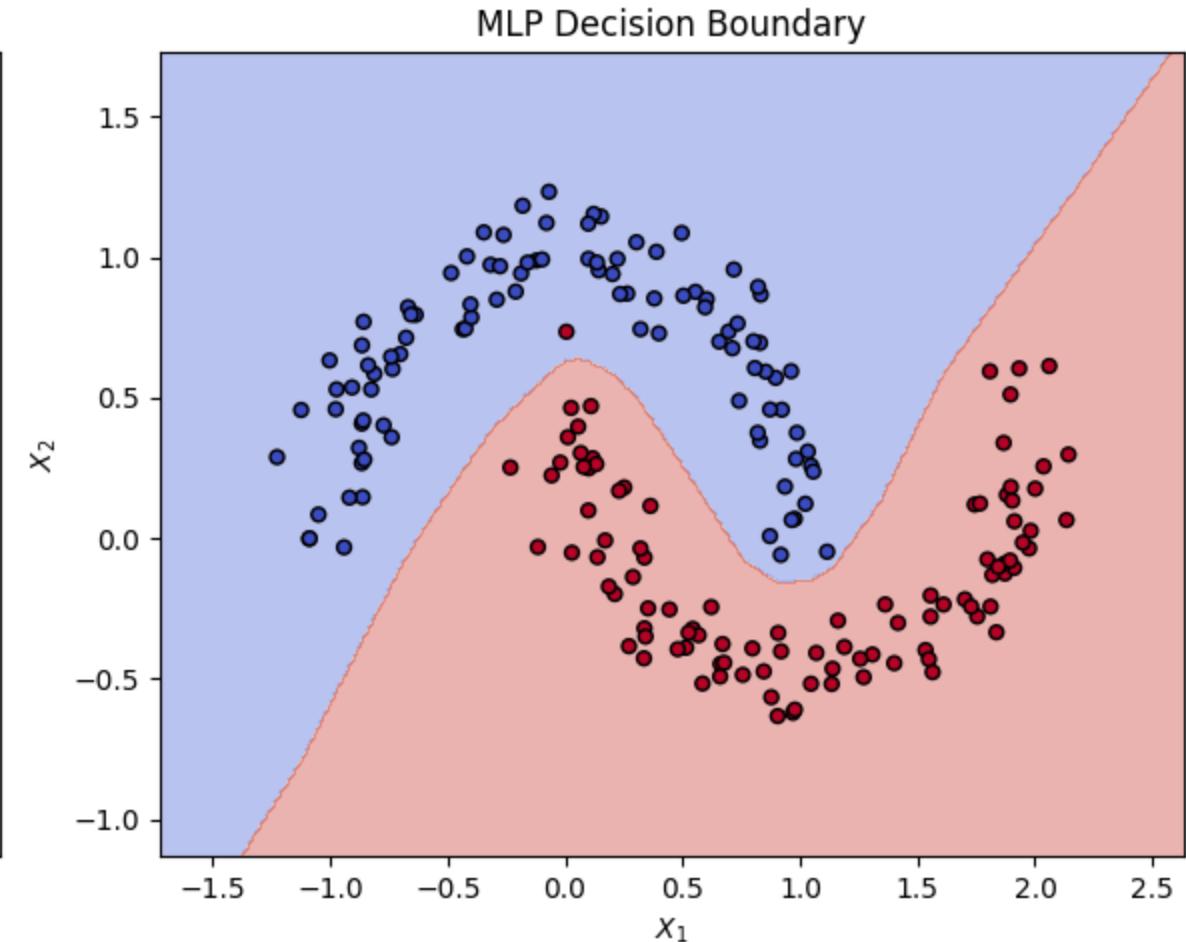
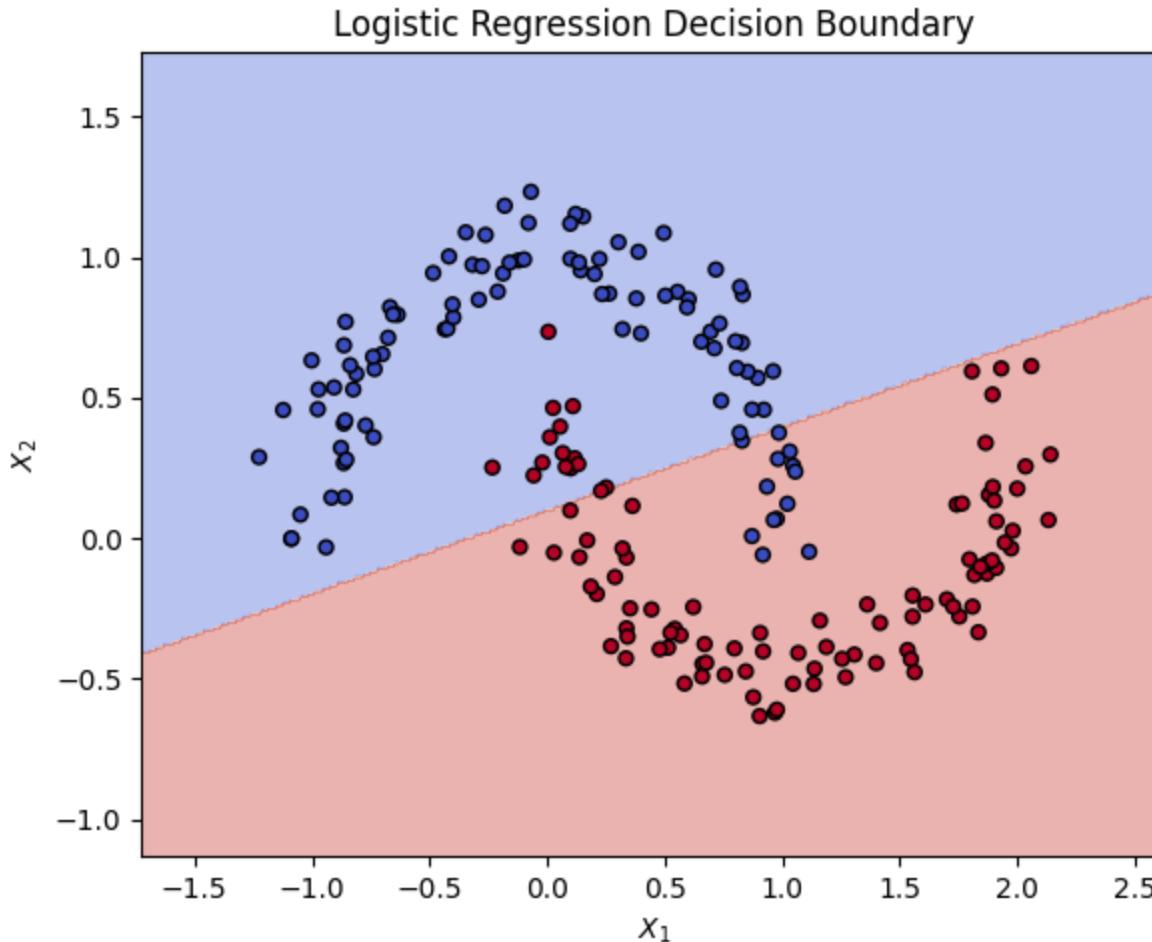
❖ Why we need MLP?



A more complex (non-linear) classification dataset.

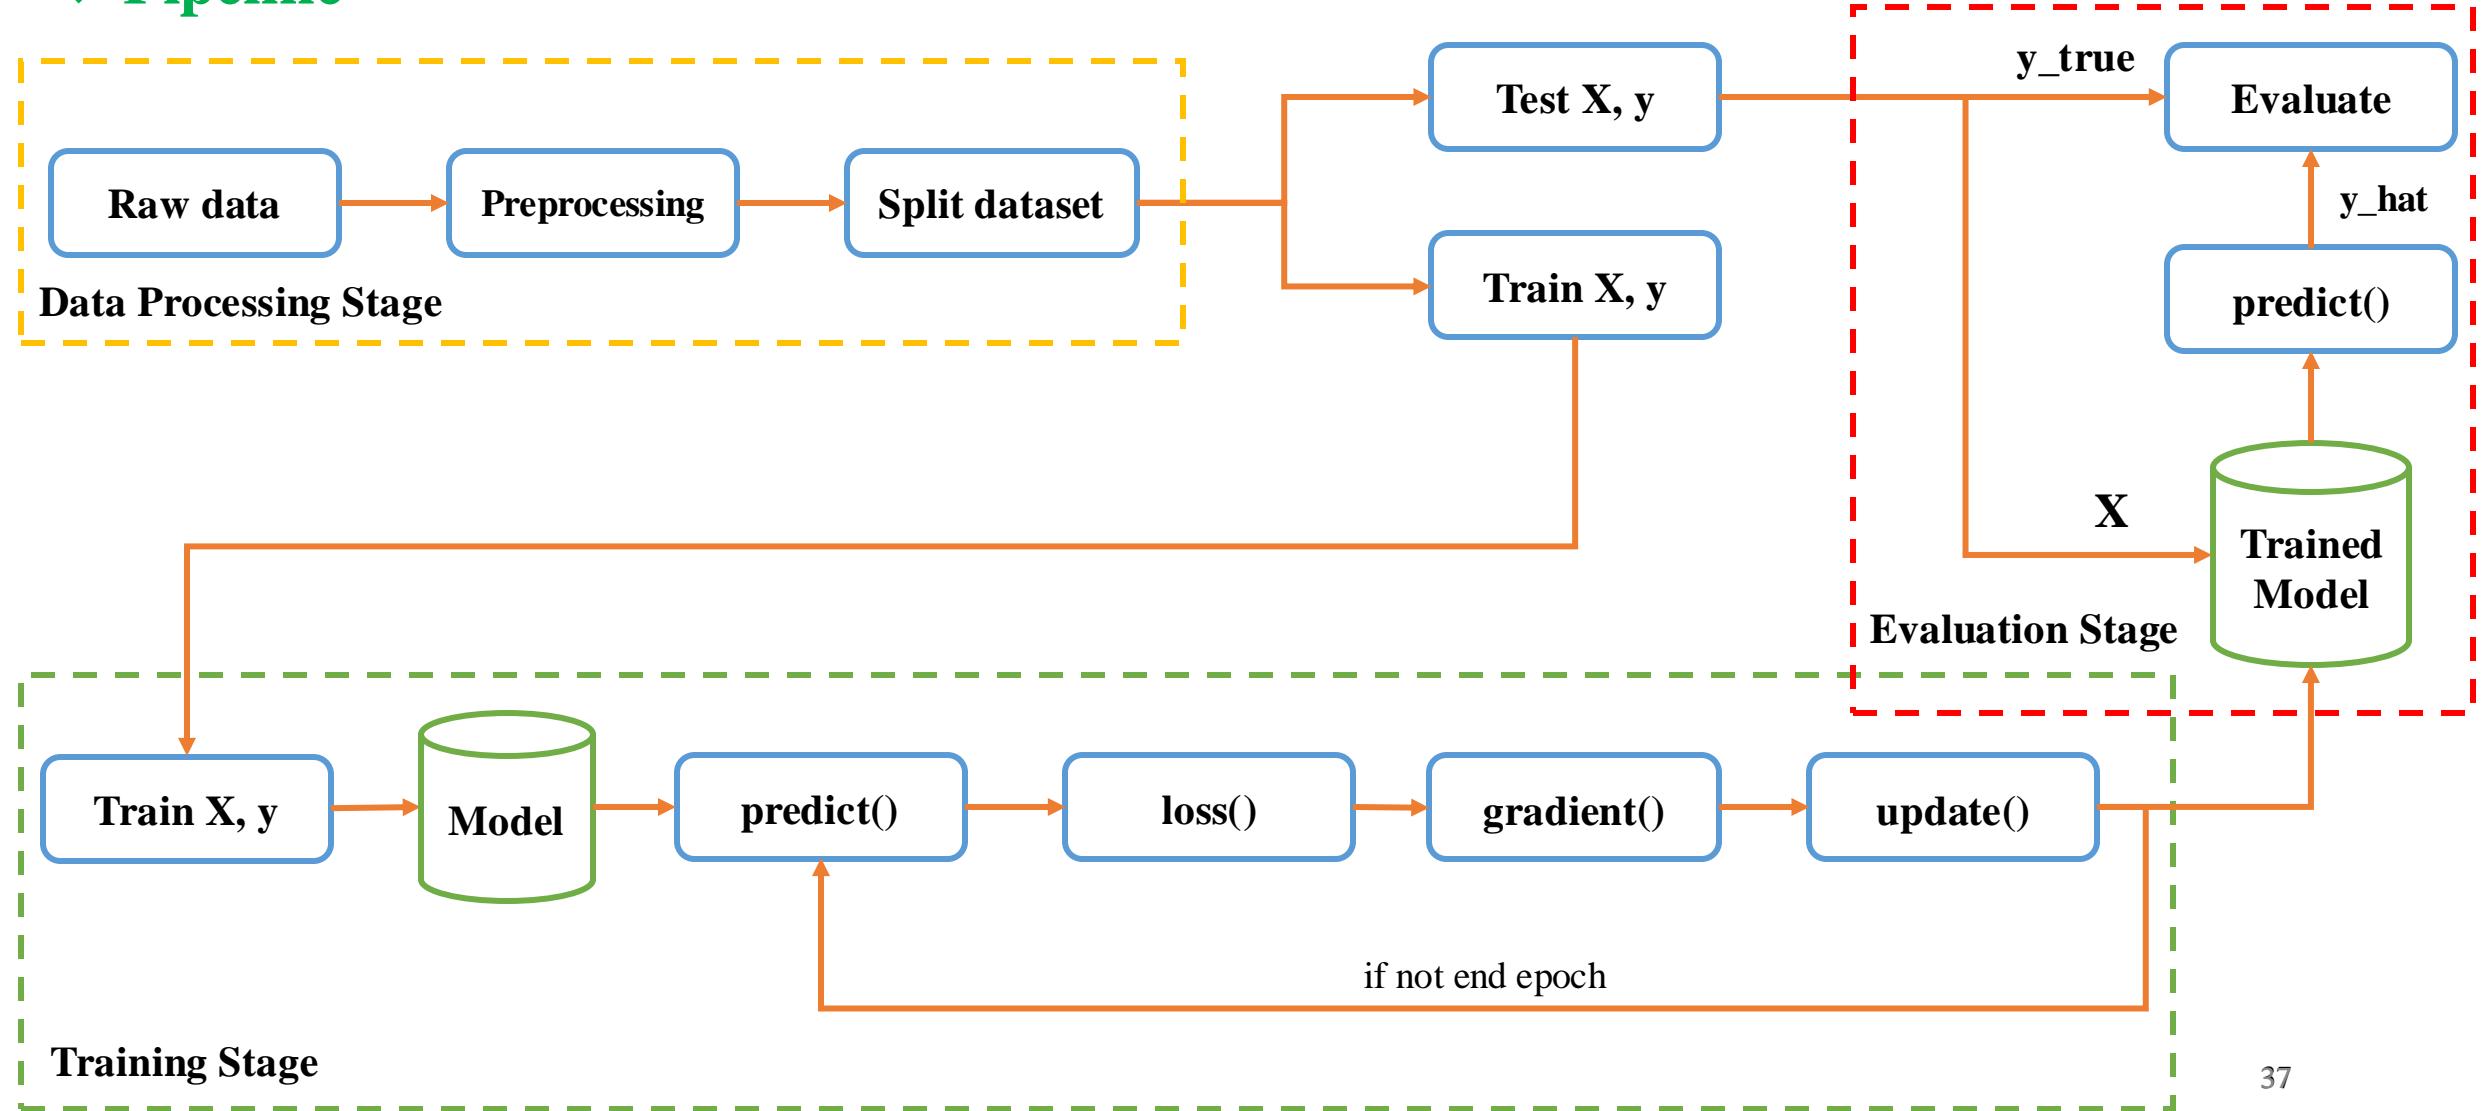
Review

❖ Why we need MLP?



Review

❖ Pipeline



Car Performance Prediction

Car Performance Prediction

❖ Introduction

Description: Given a dataset about [car performance](#), build a Multi-Layer Perceptron (MLP) model capable of accurately predicting the class labels based on the given features.



Car Performance Prediction

❖ Step 1: Import libraries and set random seed

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import torch
5 import torch.nn as nn
6 import torch.nn.functional as F
7 from torch.utils.data import Dataset, DataLoader
8
9 from sklearn.model_selection import train_test_split
10 from sklearn.preprocessing import StandardScaler
11
12 device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
13 random_state = 59
14 np.random.seed(random_state)
15 torch.manual_seed(random_state)
16 if torch.cuda.is_available():
17     torch.cuda.manual_seed(random_state)
```



Car Performance Prediction

❖ Step 2: Read dataset

```
1 dataset_path = '/content/Auto MPG_data.csv'  
2 dataset = pd.read_csv(dataset_path)  
3 dataset
```

	MPG	Cylinders	Displacement	Horsepower	Weight	Acceleration	Model Year	Europe	Japan	USA
0	18.0	8	307.0	130.0	3504.0	12.0	70	0	0	1
1	15.0	8	350.0	165.0	3693.0	11.5	70	0	0	1
2	18.0	8	318.0	150.0	3436.0	11.0	70	0	0	1
3	16.0	8	304.0	150.0	3433.0	12.0	70	0	0	1
4	17.0	8	302.0	140.0	3449.0	10.5	70	0	0	1
...
387	27.0	4	140.0	86.0	2790.0	15.6	82	0	0	1
388	44.0	4	97.0	52.0	2130.0	24.6	82	1	0	0
389	32.0	4	135.0	84.0	2295.0	11.6	82	0	0	1
390	28.0	4	120.0	79.0	2625.0	18.6	82	0	0	1
391	31.0	4	119.0	82.0	2720.0	19.4	82	0	0	1

392 rows × 10 columns

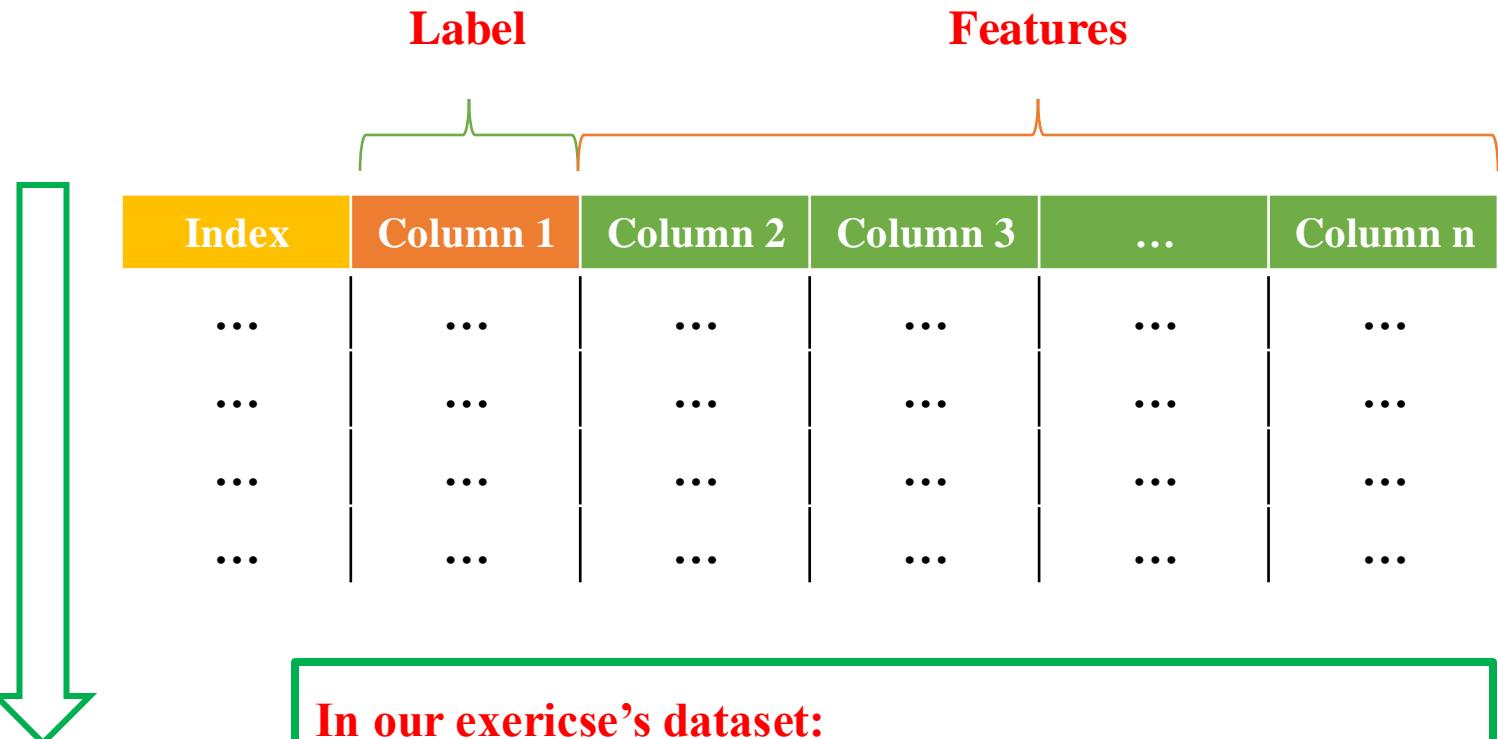
Car Performance Prediction

❖ Step 3: Split X, y



Independent Variable (Cause)

Dependent Variable (Result)



In our exercise's dataset:

- ❖ Independent variables (features): from column 2 to n .
- ❖ Dependent variables (labels): first column.

Car Performance Prediction

❖ Step 3: Split X, y

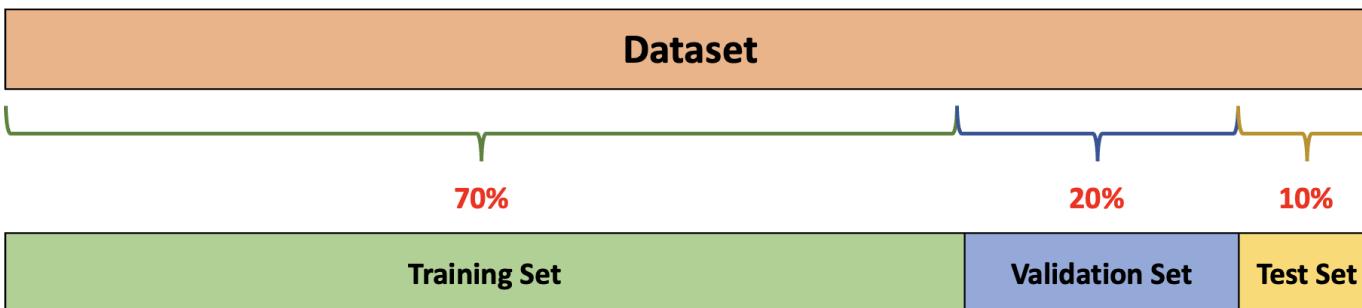
	MPG	Cylinders	Displacement	Horsepower	Weight	Acceleration	Model Year	Europe	Japan	USA
0	18.0	8	307.0	130.0	3504.0	12.0	70	0	0	1
1	15.0	8	350.0	165.0	3693.0	11.5	70	0	0	1
2	18.0	8	318.0	150.0	3436.0	11.0	70	0	0	1
3	16.0	8	304.0	150.0	3433.0	12.0	70	0	0	1
4	17.0	8	302.0	140.0	3449.0	10.5	70	0	0	1
...
387	27.0	4	140.0	86.0	2790.0	15.6	82	0	0	1
388	44.0	4	97.0	52.0	2130.0	24.6	82	1	0	0
389	32.0	4	135.0	84.0	2295.0	11.6	82	0	0	1
390	28.0	4	120.0	79.0	2625.0	18.6	82	0	0	1
391	31.0	4	119.0	82.0	2720.0	19.4	82	0	0	1

392 rows × 10 columns

```
1 X = dataset.drop(columns='MPG').values  
2 y = dataset['MPG'].values
```

Car Performance Prediction

❖ Step 4: Split train, val, test set



```
1 val_size = 0.2
2 test_size = 0.125
3 is_shuffle = True
4
5 X_train, X_val, y_train, y_val = train_test_split(
6     X, y,
7     test_size=val_size,
8     random_state=random_state,
9     shuffle=is_shuffle
10 )
11
12 X_train, X_test, y_train, y_test = train_test_split(
13     X_train, y_train,
14     test_size=test_size,
15     random_state=random_state,
16     shuffle=is_shuffle
17 )
```

```
1 print(f'Number of training samples: {X_train.shape[0]}')
2 print(f'Number of val samples: {X_val.shape[0]}')
3 print(f'Number of test samples: {X_test.shape[0]}')
```

Number of training samples: 273
Number of val samples: 79
Number of test samples: 40

Car Performance Prediction

❖ Step 5: Normalization

Using `sklearn.preprocessing.StandardScaler()` to scale all values in dataset.

$$z = \frac{x_i - \mu}{\sigma}$$

```
1 normalizer = StandardScaler()  
2 X_train = normalizer.fit_transform(X_train)  
3 X_val = normalizer.transform(X_val)  
4 X_test = normalizer.transform(X_test)
```

```
1 X_train
```

```
array([[-0.84144139, -1.03791712, -1.03812675, ..., -0.44426166,  
       1.90449958, -1.27475488],  
      [ 0.36757703,  0.36824297,  0.22296166, ..., -0.44426166,  
      -0.52507231,  0.78446454],  
      [-0.84144139, -0.916522, -0.89800582, ..., -0.44426166,  
      -0.52507231,  0.78446454],  
      ...,  
      [ 1.57659545,  1.30905511,  1.34392915, ..., -0.44426166,  
      -0.52507231,  0.78446454],  
      [ 0.36757703,  0.43905678, -0.33752208, ..., -0.44426166,  
      -0.52507231,  0.78446454],  
      [ 0.36757703,  0.62114946, -0.05728021, ..., -0.44426166,  
      -0.52507231,  0.78446454]])
```

Note: We only use the train set to fit the scaler.

Car Performance Prediction

❖ Step 5: Normalization

```
1 print("Shape of X:", X.shape)
2 print("Shape of y:", y.shape)
3
4 print("First 5 rows of X:")
5 print(X[:5])
6 print("First 5 labels:")
7 print(y[:5])
8
```

Shape of X: (300, 2)

Shape of y: (300,)

First 5 rows of X:

```
[[0.          0.        ]
 [0.00096008 0.01005528]
 [0.01045864 0.01728405]
 [0.00087922 0.03029027]
 [0.00991727 0.03916803]]
```

First 5 labels:

```
[0 0 0 0 0]
```

```
1 normalizer = StandardScaler()
2 X_train = normalizer.fit_transform(X_train)
3 X_val = normalizer.transform(X_val)
4 X_test = normalizer.transform(X_test)
5
6 X_train = torch.tensor(X_train, dtype=torch.float32)
7 X_val = torch.tensor(X_val, dtype=torch.float32)
8 X_test = torch.tensor(X_test, dtype=torch.float32)
9 y_train = torch.tensor(y_train, dtype=torch.long)
10 y_val = torch.tensor(y_val, dtype=torch.long)
11 y_test = torch.tensor(y_test, dtype=torch.long)
```

Since the original has Numpy arrays form, we need to convert the datasets (features) into tensors of float32 data type as PyTorch libraries require the input data to be tensors.

Car Performance Prediction

❖ Step 5: Normalization

```
1 print("Shape of X:", X.shape)
2 print("Shape of y:", y.shape)
3
4 print("First 5 rows of X:")
5 print(X[:5])
6 print("First 5 labels:")
7 print(y[:5])
8
```

```
Shape of X: (300, 2)
Shape of y: (300,)
First 5 rows of X:
[[0.          0.        ]
 [0.00096008 0.01005528]
 [0.01045864 0.01728405]
 [0.00087922 0.03029027]
 [0.00991727 0.03916803]]
First 5 labels:
[0 0 0 0 0]
```

```
1 normalizer = StandardScaler()
2 X_train = normalizer.fit_transform(X_train)
3 X_val = normalizer.transform(X_val)
4 X_test = normalizer.transform(X_test)
5
6 X_train = torch.tensor(X_train, dtype=torch.float32)
7 X_val = torch.tensor(X_val, dtype=torch.float32)
8 X_test = torch.tensor(X_test, dtype=torch.float32)
9 y_train = torch.tensor(y_train, dtype=torch.long)
10 y_val = torch.tensor(y_val, dtype=torch.long)
11 y_test = torch.tensor(y_test, dtype=torch.long)
```

- Then convert the label datasets into tensors of long type, commonly used for classification tasks because labels are typically integers (e.g., 0, 1, 2).
- As PyTorch also requires labels in classification tasks to be tensors of type long to be compatible with loss functions (CrossEntropyLoss)

Car Performance Prediction

❖ Step 6: Build Dataloader

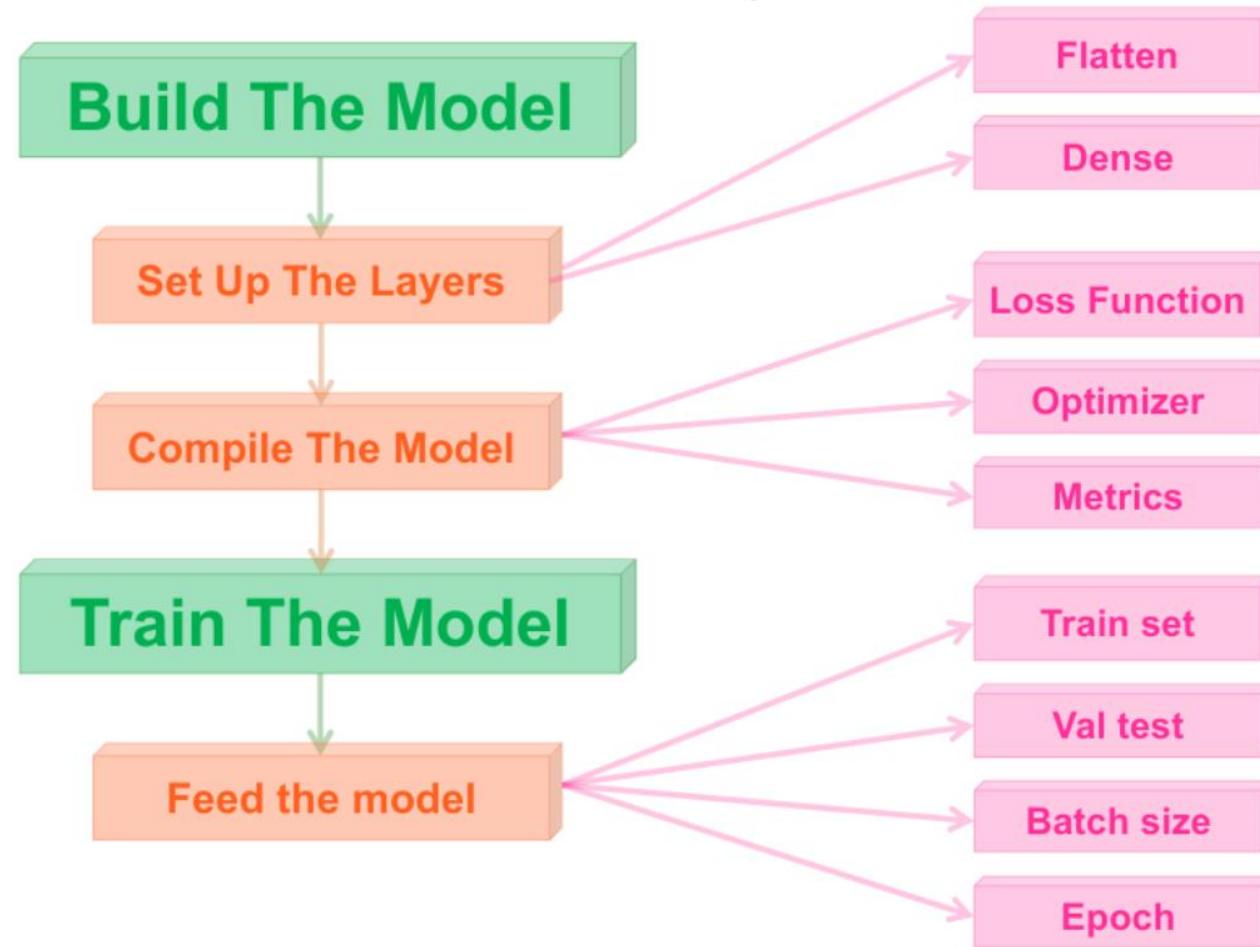
```
1 class CustomDataset(Dataset):
2     def __init__(self, X, y):
3         self.X = X
4         self.y = y
5
6     def __len__(self):
7         return len(self.y)
8
9     def __getitem__(self, idx):
10        return self.X[idx], self.y[idx]
```

```
1 batch_size = 32
2 train_dataset = CustomDataset(X_train, y_train)
3 val_dataset = CustomDataset(X_val, y_val)
4 train_loader = DataLoader(train_dataset,
5                           batch_size=batch_size,
6                           shuffle=True)
7 val_loader = DataLoader(val_dataset,
8                           batch_size=batch_size,
9                           shuffle=False)
```

To efficiently handle data for training and evaluation in PyTorch, we use **DataLoader**, which requires a custom **Dataset** to manage inputs and labels.

Car Performance Prediction

❖ Step 7: Build MLP model

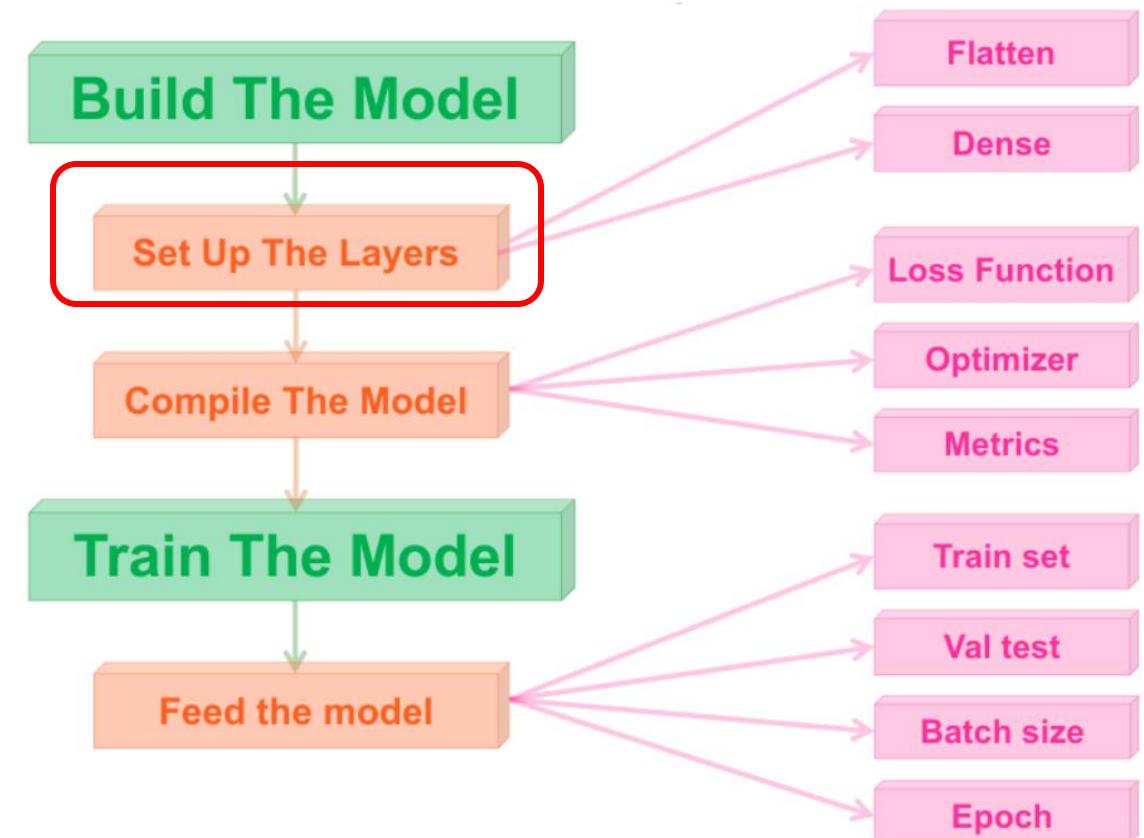


Pipeline to build and train a model in PyTorch.

Car Performance Prediction

❖ Step 7: Build MLP model

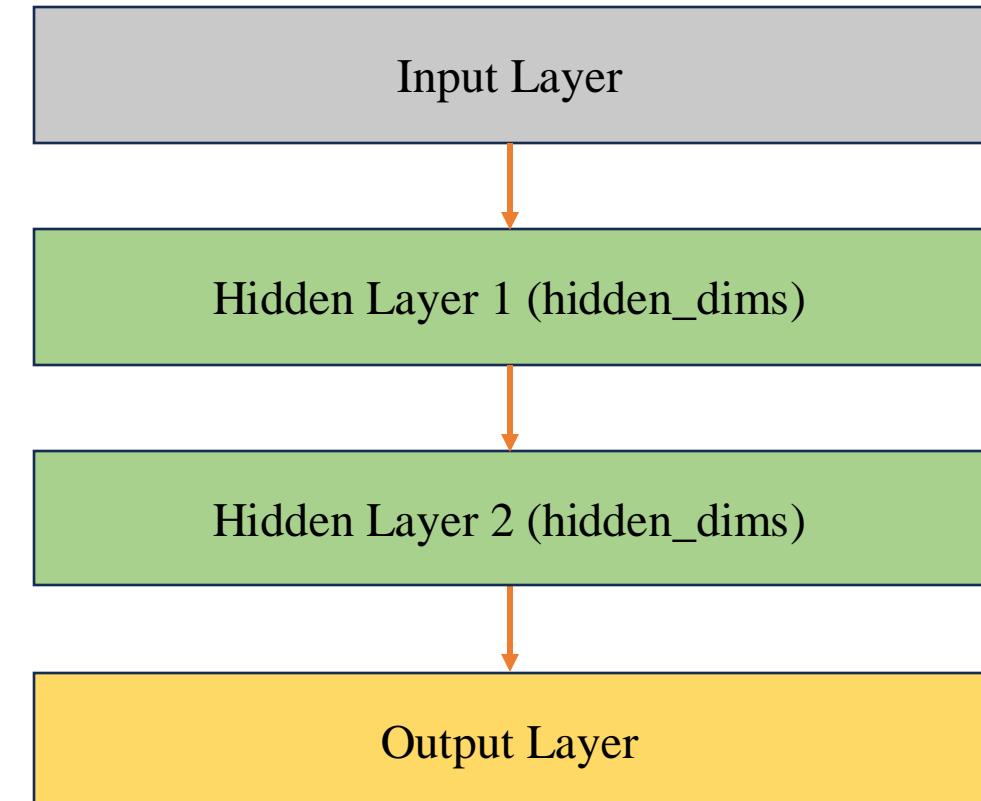
```
1 class MLP(nn.Module):
2     def __init__(self, input_dims, hidden_dims, output_dims):
3         super(MLP, self). init_()
4         self.linear1 = nn.Linear(input_dims, hidden_dims)
5         self.linear2 = nn.Linear(hidden_dims, hidden_dims)
6         self.output = nn.Linear(hidden_dims, output_dims)
7
8     def forward(self, x):
9         x = self.linear1(x)
10        x = F.relu(x)
11        x = self.linear2(x)
12        x = F.relu(x)
13        out = self.output(x)
14        return out.squeeze(1)
```



Car Performance Prediction

❖ Step 7: Build MLP model

```
1 class MLP(nn.Module):
2     def __init__(self, input_dims, hidden_dims, output_dims):
3         super(MLP, self).__init__()
4         self.linear1 = nn.Linear(input_dims, hidden_dims)
5         self.linear2 = nn.Linear(hidden_dims, hidden_dims)
6         self.output = nn.Linear(hidden_dims, output_dims)
7
8     def forward(self, x):
9         x = self.linear1(x)
10        x = F.relu(x)
11        x = self.linear2(x)
12        x = F.relu(x)
13        out = self.output(x)
14        return out.squeeze(1)
```

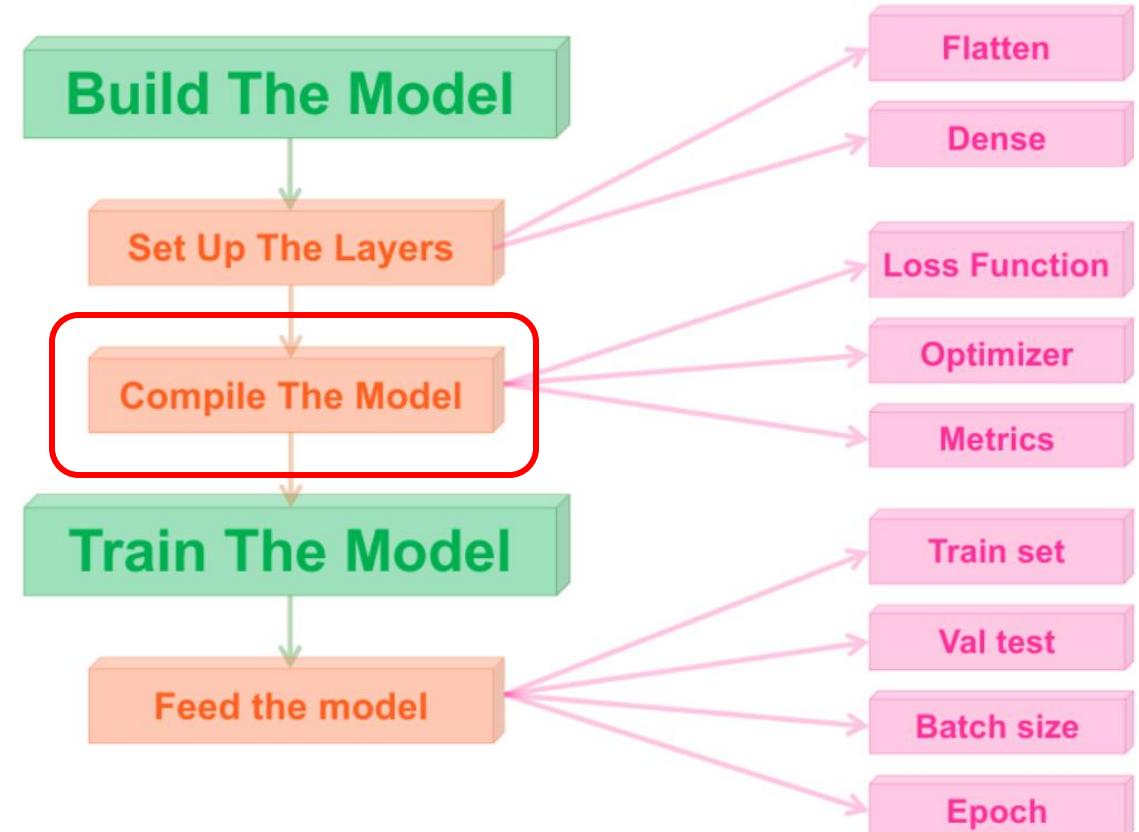


Car Performance Prediction

❖ Step 7: Declare model, optimizer and criterion

```
1 class MLP(nn.Module):
2     def __init__(self, input_dims, hidden_dims, output_dims):
3         super(MLP, self).__init__()
4         self.linear1 = nn.Linear(input_dims, hidden_dims)
5         self.linear2 = nn.Linear(hidden_dims, hidden_dims)
6         self.output = nn.Linear(hidden_dims, output_dims)
7
8     def forward(self, x):
9         x = self.linear1(x)
10        x = F.relu(x)
11        x = self.linear2(x)
12        x = F.relu(x)
13        out = self.output(x)
14        return out.squeeze(1)
```

```
1 input_dims = X_train.shape[1]
2 output_dims = 1
3 hidden_dims = 64
4 lr = 1e-2
5
6 model = MLP(input_dims=input_dims,
7             hidden_dims=hidden_dims,
8             output_dims=output_dims).to(device)
9 criterion = nn.MSELoss()
10 optimizer = torch.optim.SGD(model.parameters(), lr=lr)
```

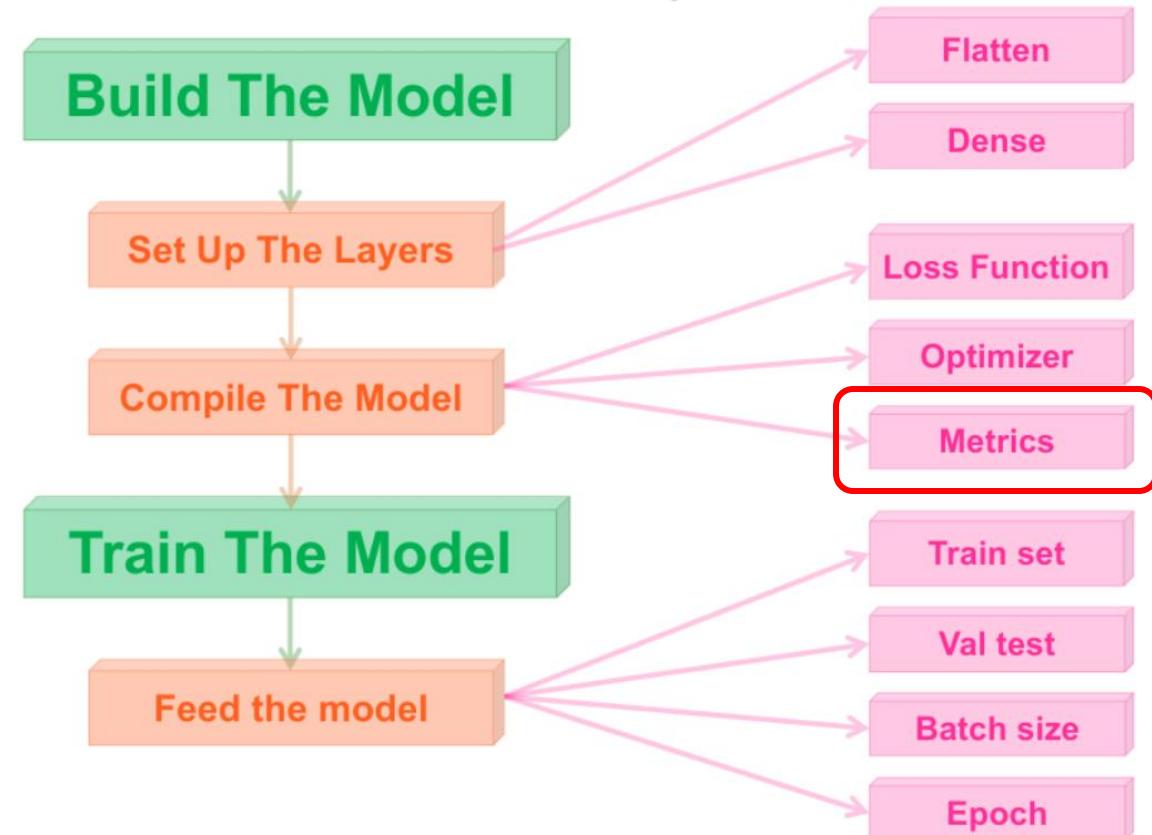


Car Performance Prediction

❖ Step 8: Create R2 Score function

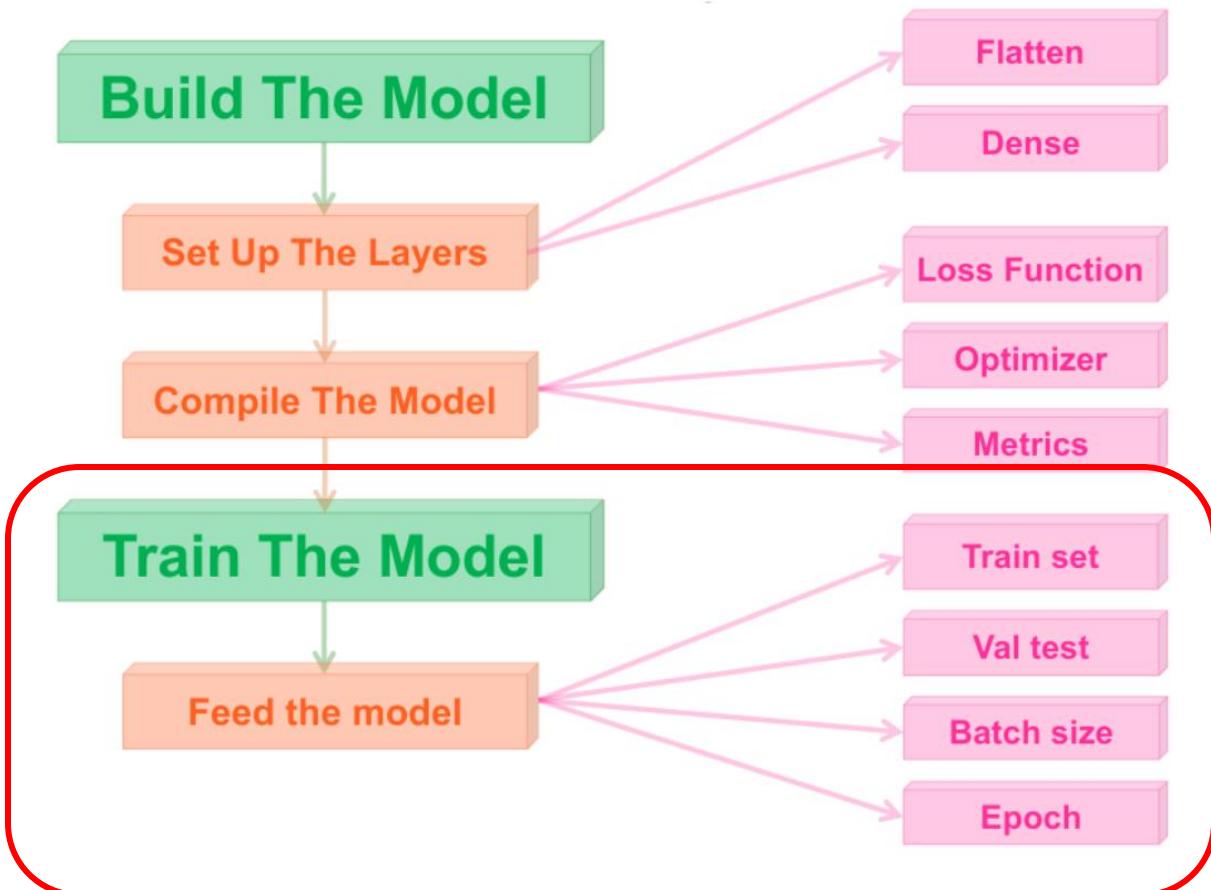
$$R^2 = 1 - \frac{SS_{res}}{SS_{tot}} = 1 - \frac{\sum_i (y_i - \hat{y}_i)^2}{\sum_i (y_i - \bar{y}_i)^2}$$

```
1 def r_squared(y_true, y_pred):
2     y_true = torch.Tensor(y_true).to(device)
3     y_pred = torch.Tensor(y_pred).to(device)
4     mean_true = torch.mean(y_true)
5     ss_tot = torch.sum((y_true - mean_true) ** 2)
6     ss_res = torch.sum((y_true - y_pred) ** 2)
7     r2 = 1 - (ss_res / ss_tot)
8     return r2
```



Car Performance Prediction

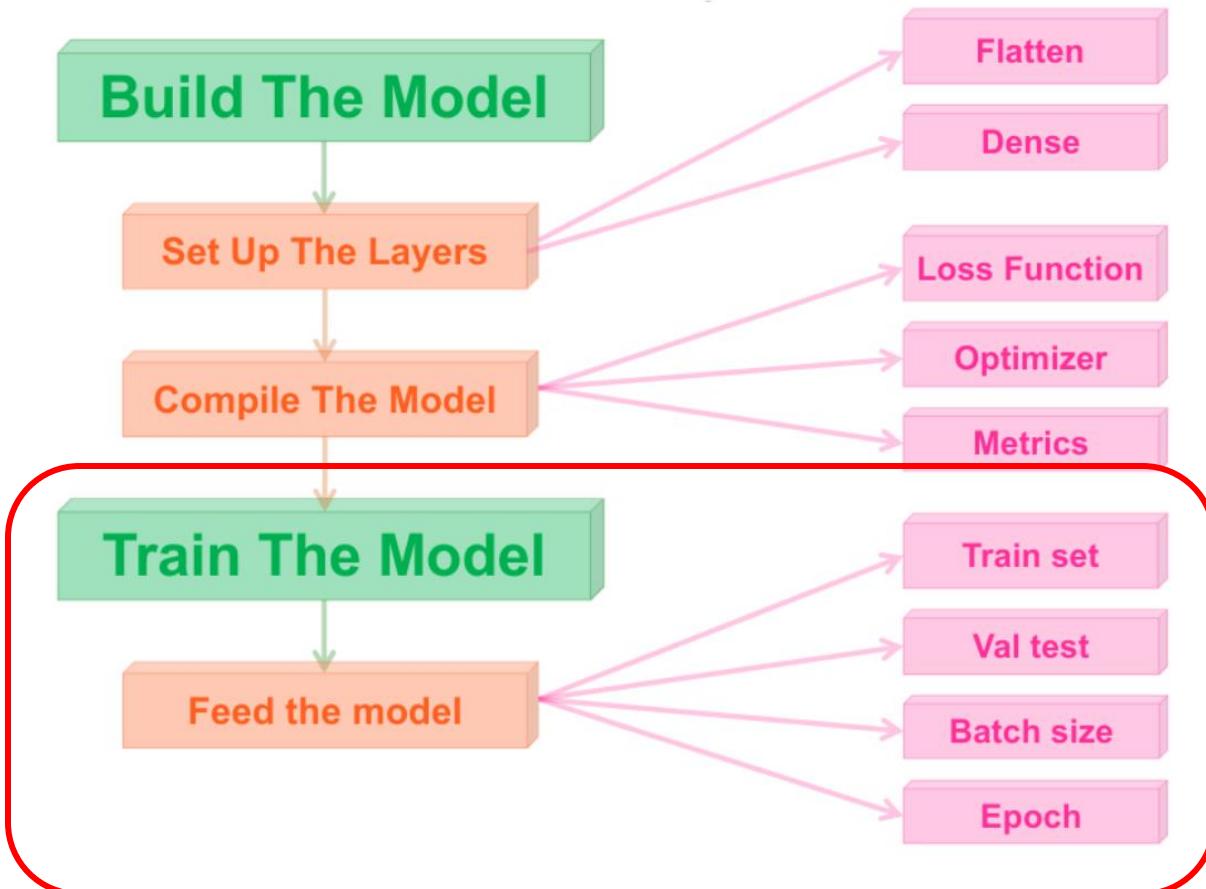
❖ Step 9: Training



```
1 epochs = 100
2 train_losses = []
3 val_losses = []
4 train_r2 = []
5 val_r2 = []
6
7 for epoch in range(epochs):
8     train_loss = 0.0
9     train_target = []
.0    val_target = []
.1    train_predict = []
.2    val_predict = []
.3    model.train()
.4
.5    for X_samples, y_samples in train_loader:
.6        X_samples = X_samples.to(device)
.7        y_samples = y_samples.to(device)
.8        optimizer.zero_grad()
.9        outputs = model(X_samples)
.00        train_predict += outputs.tolist()
.01        train_target += y_samples.tolist()
.02        loss = criterion(outputs, y_samples)
.03        loss.backward()
.04        optimizer.step()
.05        train_loss += loss.item()
.06    train_loss /= len(train_loader)
.07    train_losses.append(train_loss)
.08    train_r2.append(r_squared(train_target, train_predict))
```

Car Performance Prediction

❖ Step 9: Training (Evaluation)

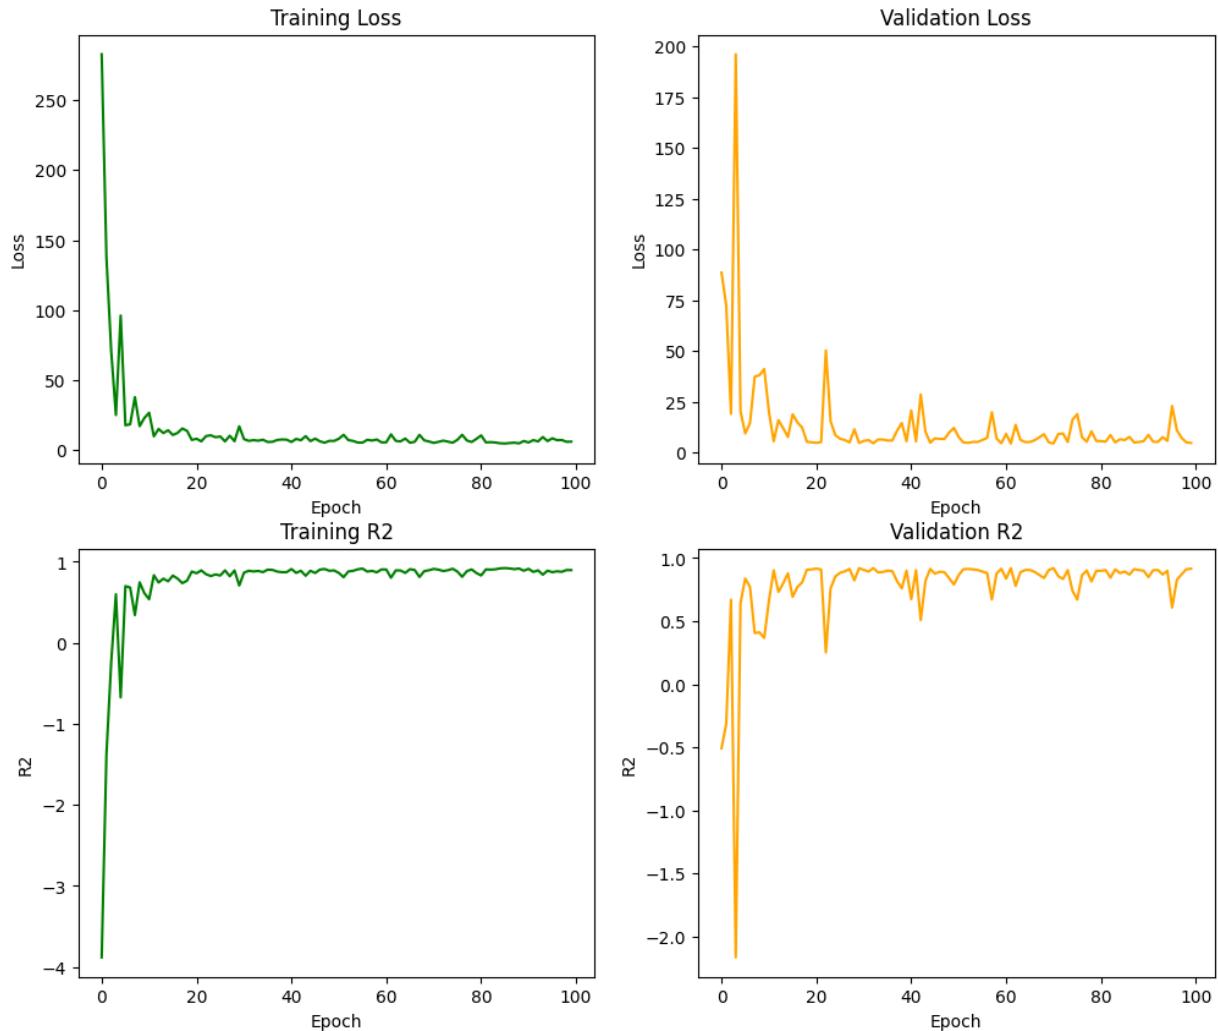


```
29 model.eval()
30 val_loss = 0.0
31 with torch.no_grad():
32     for X_samples, y_samples in val_loader:
33         X_samples = X_samples.to(device)
34         y_samples = y_samples.to(device)
35         outputs = model(X_samples)
36         val_predict += outputs.tolist()
37         val_target += y_samples.tolist()
38         loss = criterion(outputs, y_samples)
39         val_loss += loss.item()
40     val_loss /= len(val_loader)
41     val_losses.append(val_loss)
42     val_r2.append(r_squared(val_target, val_predict))
```

Car Performance Prediction

❖ Step 10: Evaluation

```
1 fig, ax = plt.subplots(2, 2, figsize=(12, 10))
2 ax[0, 0].plot(train_losses, color='green')
3 ax[0, 0].set(xlabel='Epoch', ylabel='Loss')
4 ax[0, 0].set_title('Training Loss')
5
6 ax[0, 1].plot(val_losses, color='orange')
7 ax[0, 1].set(xlabel='Epoch', ylabel='Loss')
8 ax[0, 1].set_title('Validation Loss')
9
10 ax[1, 0].plot(train_r2, color='green')
11 ax[1, 0].set(xlabel='Epoch', ylabel='R2')
12 ax[1, 0].set_title('Training R2')
13
14 ax[1, 1].plot(val_r2, color='orange')
15 ax[1, 1].set(xlabel='Epoch', ylabel='R2')
16 ax[1, 1].set_title('Validation R2')
17
18 plt.show()
```



Car Performance Prediction

❖ Step 10: Evaluation

$$R^2 = 1 - \frac{SS_{res}}{SS_{tot}} = 1 - \frac{\sum_i (y_i - \hat{y}_i)^2}{\sum_i (y_i - \bar{y}_i)^2}$$

```
1 def r_squared(y_true, y_pred):
2     y_true = torch.Tensor(y_true).to(device)
3     y_pred = torch.Tensor(y_pred).to(device)
4     mean_true = torch.mean(y_true)
5     ss_tot = torch.sum((y_true - mean_true) ** 2)
6     ss_res = torch.sum((y_true - y_pred) ** 2)
7     r2 = 1 - (ss_res / ss_tot)
8     return r2
```

```
1 model.eval()
2 with torch.no_grad():
3     y_hat = model(X_val)
4     val_set_r2 = r_squared(y_hat, y_val)
5     print('Evaluation on validation set:')
6     print(f'R2: {val_set_r2}')
```

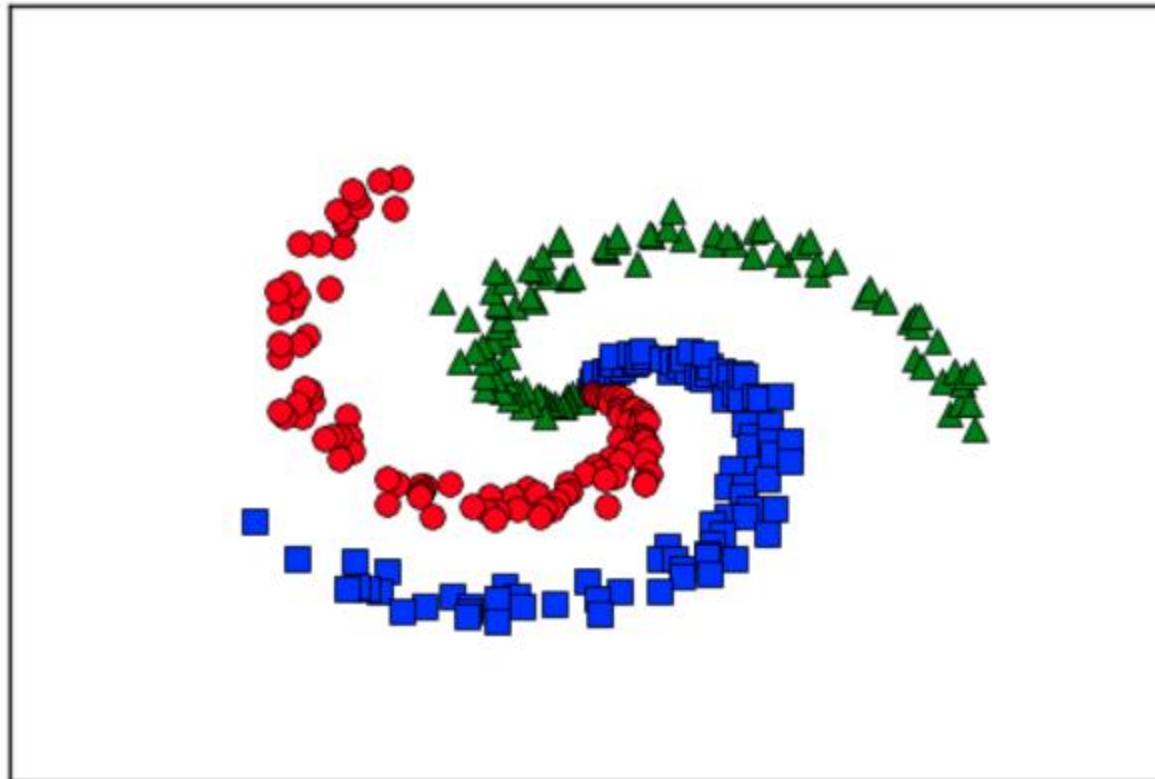
Evaluation on validation set:
R2: 0.893627941608429

Non-linear Data Classification

Non-linear Data Classification

❖ Introduction

Description: Classify non-linear data using the dataset [NonLinear_data.npy](#), build a Multi-Layer Perceptron (MLP) model capable of accurately predicting the class labels based on the given features.



Non-linear Data Classification

❖ Step 1: Import libraries and set random seed

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import torch
5 import torch.nn as nn
6 import torch.nn.functional as F
7 from torch.utils.data import Dataset, DataLoader
8
9 from sklearn.model_selection import train_test_split
10 from sklearn.preprocessing import StandardScaler
11
12 device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
13 random_state = 59
14 np.random.seed(random_state)
15 torch.manual_seed(random_state)
16 if torch.cuda.is_available():
17     torch.cuda.manual_seed(random_state)
```



❖ Step 2: Read Dataset and Visualize

```
1 data_path = '/content/NonLinear_data.npy'  
2 data = np.load(data_path, allow_pickle=True).item()  
3 X, y = data['X'], data['labels']
```

numpy.load

`numpy.load(file, mmap_mode=None, allow_pickle=False, fix_imports=True, encoding='ASCII', *, max_header_size=10000)` [\[source\]](#)

Load arrays or pickled objects from `.npy`, `.npz` or pickled files.

⚠ Warning

Loading files that contain object arrays uses the `pickle` module, which is not secure against erroneous or maliciously constructed data. Consider passing `allow_pickle=False` to load data that is known not to contain object arrays for the safer handling of untrusted sources.

```
1 print("Shape of X:", X.shape)  
2 print("Shape of y:", y.shape)  
3  
4 print("First 5 rows of X:")  
5 print(X[:5])  
6 print("First 5 labels:")  
7 print(y[:5])  
8
```

Shape of X: (300, 2)

Shape of y: (300,)

First 5 rows of X:

```
[[0.          0.          ]  
 [0.00096008  0.01005528]  
 [0.01045864  0.01728405]  
 [0.00087922  0.03029027]  
 [0.00991727  0.03916803]]
```

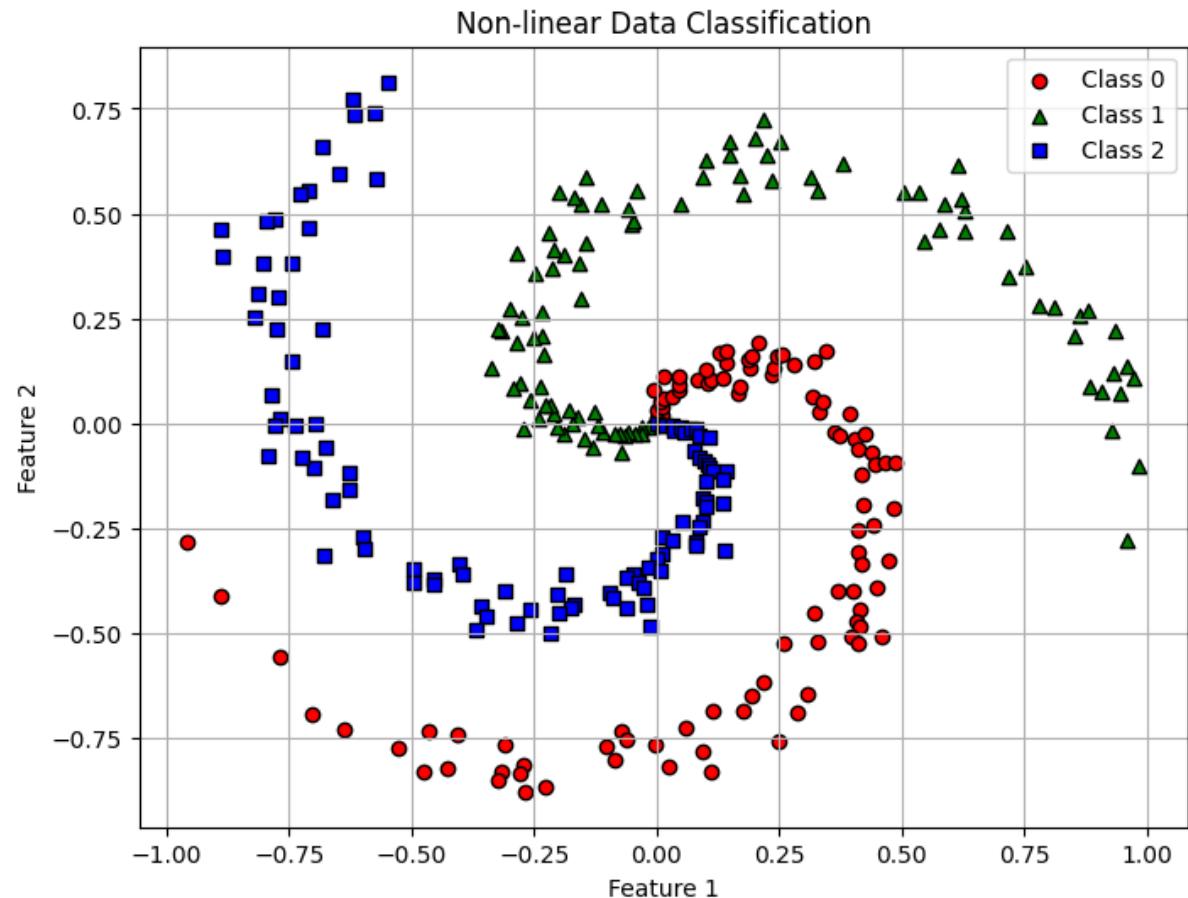
First 5 labels:

```
[0 0 0 0 0]
```

Non-linear Data Classification

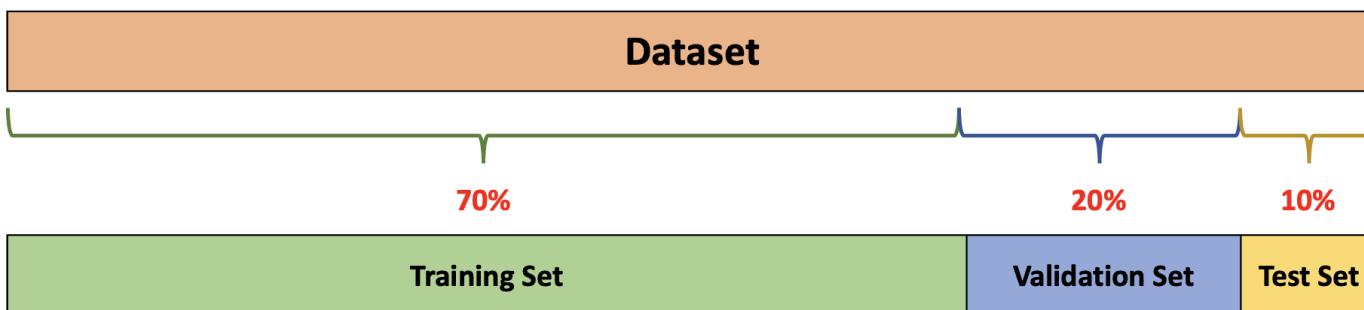
❖ Step 2: Read Dataset and Visualize

```
1 colors = ['red', 'green', 'blue']
2 markers = ['o', '^', 's']
3
4 plt.figure(figsize=(8, 6))
5 for i, class_label in enumerate(np.unique(y)):
6     plt.scatter(
7         X[y == class_label, 0],
8         X[y == class_label, 1],
9         c=colors[i],
10        marker=markers[i],
11        label=f'Class {class_label}',
12        edgecolor='k'
13    )
15 plt.title("Non-linear Data Classification")
16 plt.xlabel("Feature 1")
17 plt.ylabel("Feature 2")
18 plt.legend()
19 plt.grid(True)
20 plt.show()
```



Non-linear Data Classification

❖ Step 3: Split train, val, test set



```
1 val_size = 0.2
2 test_size = 0.125
3 is_shuffle = True
4
5 X_train, X_val, y_train, y_val = train_test_split(
6     X, y,
7     test_size=val_size,
8     random_state=random_state,
9     shuffle=is_shuffle
10 )
11
12 X_train, X_test, y_train, y_test = train_test_split(
13     X_train, y_train,
14     test_size=test_size,
15     random_state=random_state,
16     shuffle=is_shuffle
17 )
```

```
1 print(f'Number of training samples: {X_train.shape[0]}')
2 print(f'Number of val samples: {X_val.shape[0]}')
3 print(f'Number of test samples: {X_test.shape[0]}')
```

Number of training samples: 210
Number of val samples: 60
Number of test samples: 30

Non-linear Data Classification

❖ Step 4: Normalization

Using `sklearn.preprocessing.StandardScaler()` to scale all values in dataset. Only use the train set to fit the scaler.

$$z = \frac{x_i - \mu}{\sigma}$$

```
1 normalizer = StandardScaler()
2 X_train = normalizer.fit_transform(X_train)
3 X_val = normalizer.transform(X_val)
4 X_test = normalizer.transform(X_test)
5
6 X_train = torch.tensor(X_train, dtype=torch.float32)
7 X_val = torch.tensor(X_val, dtype=torch.float32)
8 X_test = torch.tensor(X_test, dtype=torch.float32)
9 y_train = torch.tensor(y_train, dtype=torch.long)
10 y_val = torch.tensor(y_val, dtype=torch.long)
11 y_test = torch.tensor(y_test, dtype=torch.long)
```

Non-linear Data Classification

❖ Step 5: Build Dataloader

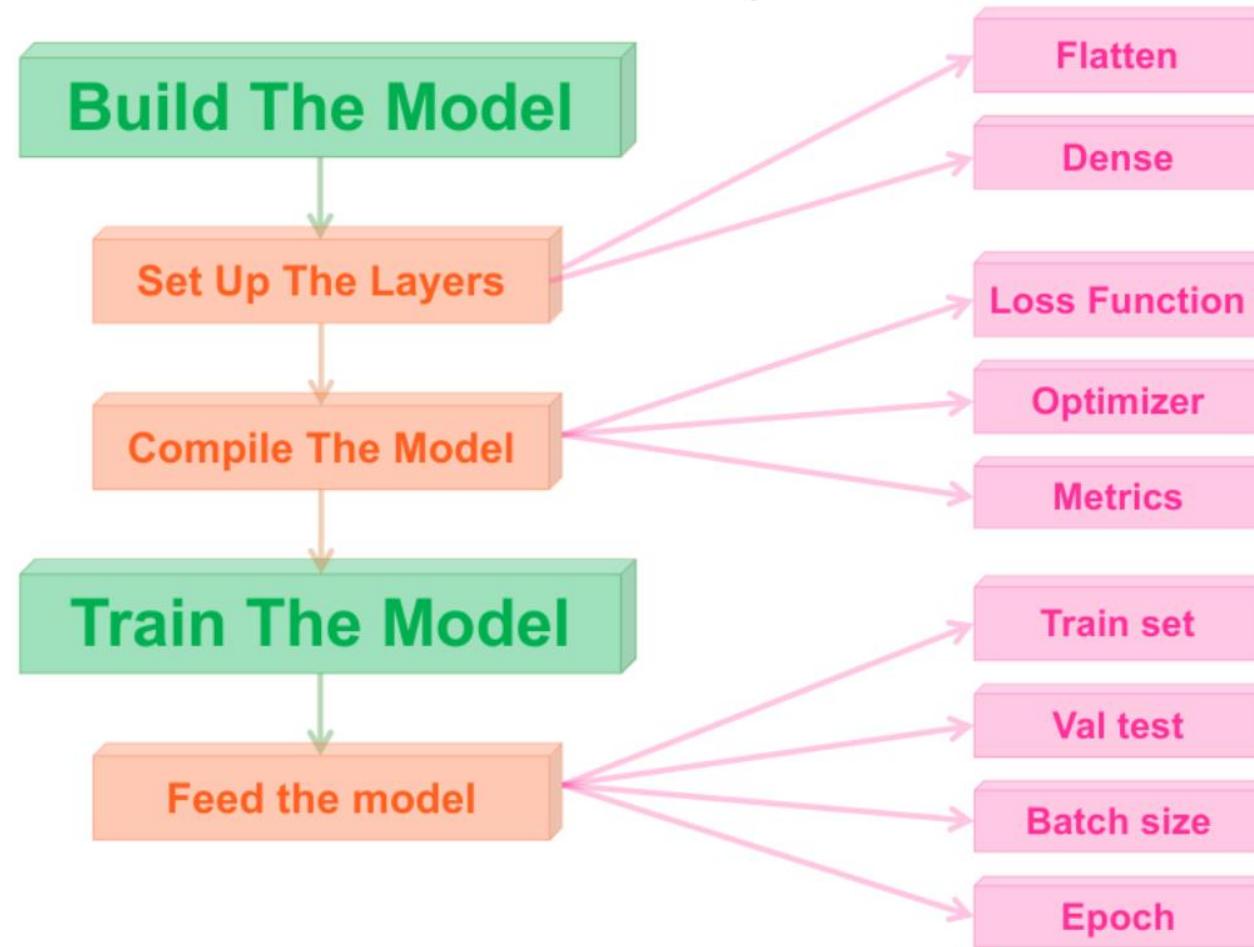
```
1 class CustomDataset(Dataset):  
2     def __init__(self, x, y):  
3         self.x = x  
4         self.y = y  
5  
6     def __len__(self):  
7         return len(self.y)  
8  
9     def __getitem__(self, idx):  
10        return self.x[idx], self.y[idx]
```

```
1 batch_size = 32  
2 train_dataset = CustomDataset(x_train, y_train)  
3 val_dataset = CustomDataset(x_val, y_val)  
4 test_dataset = CustomDataset(x_test, y_test)  
5 train_loader = DataLoader(train_dataset,  
6                           batch_size=batch_size,  
7                           shuffle=True)  
8 val_loader = DataLoader(val_dataset,  
9                         batch_size=batch_size,  
10                        shuffle=False)  
11 test_loader = DataLoader(test_dataset,  
12                          batch_size=batch_size,  
13                          shuffle=False)
```

To efficiently handle data for training and evaluation in PyTorch, we use **DataLoader**, which requires a custom **Dataset** to manage inputs and labels.

Non-linear Data Classification

❖ Step 6: Build MLP model

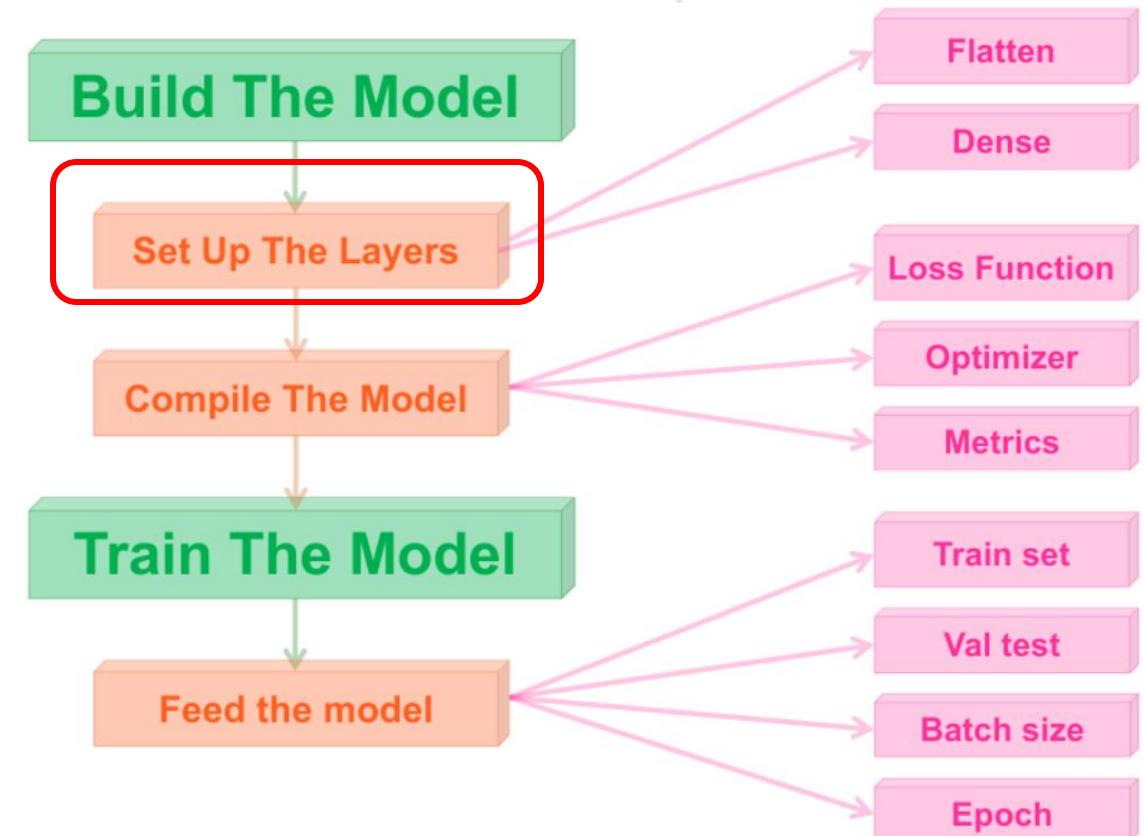


Pipeline to build and train a model in PyTorch.

Non-linear Data Classification

❖ Step 6: Build MLP model

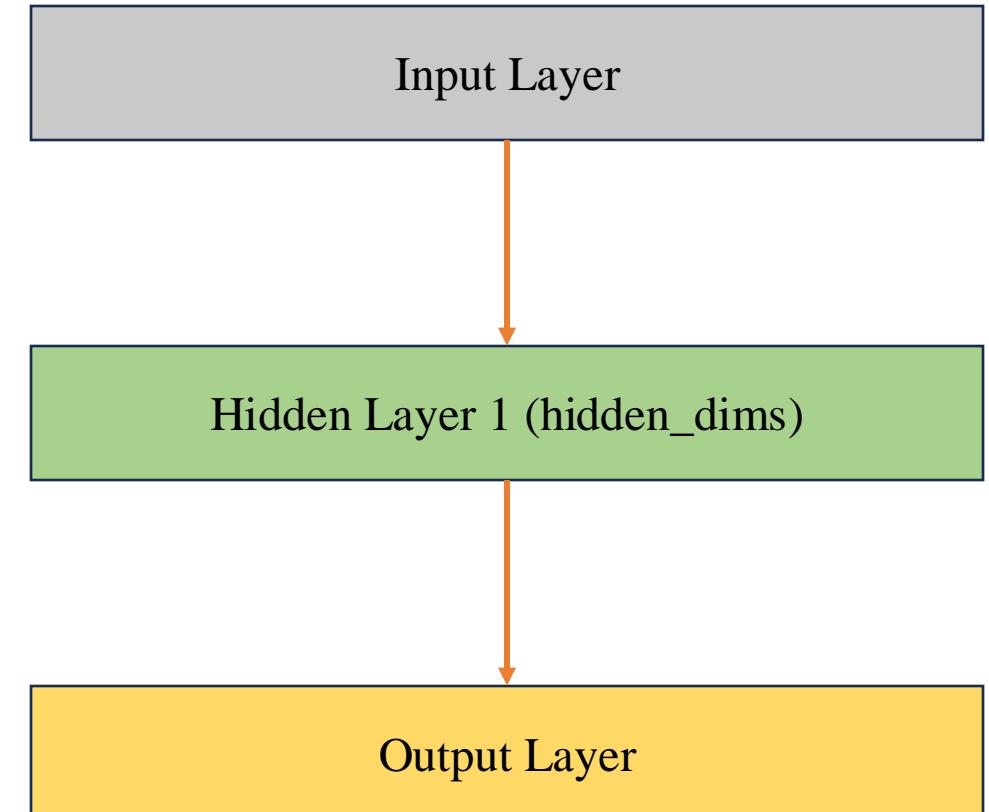
```
1 class MLP(nn.Module):
2     def __init__(self, input_dims, hidden_dims, output_dims):
3         super(MLP, self).__init__()
4         self.linear1 = nn.Linear(input_dims, hidden_dims)
5         self.output = nn.Linear(hidden_dims, output_dims)
6         self.relu = nn.ReLU()
7
8     def forward(self, x):
9         x = self.linear1(x)
10        x = self.relu(x)
11        out = self.output(x)
12        return out.squeeze(1)
```



Non-linear Data Classification

❖ Step 6: Build MLP model

```
1 class MLP(nn.Module):
2     def __init__(self, input_dims, hidden_dims, output_dims):
3         super(MLP, self).__init__()
4         self.linear1 = nn.Linear(input_dims, hidden_dims)
5         self.output = nn.Linear(hidden_dims, output_dims)
6         self.relu = nn.ReLU()
7
8     def forward(self, x):
9         x = self.linear1(x)
10        x = self.relu(x)
11        out = self.output(x)
12        return out.squeeze(1)
```

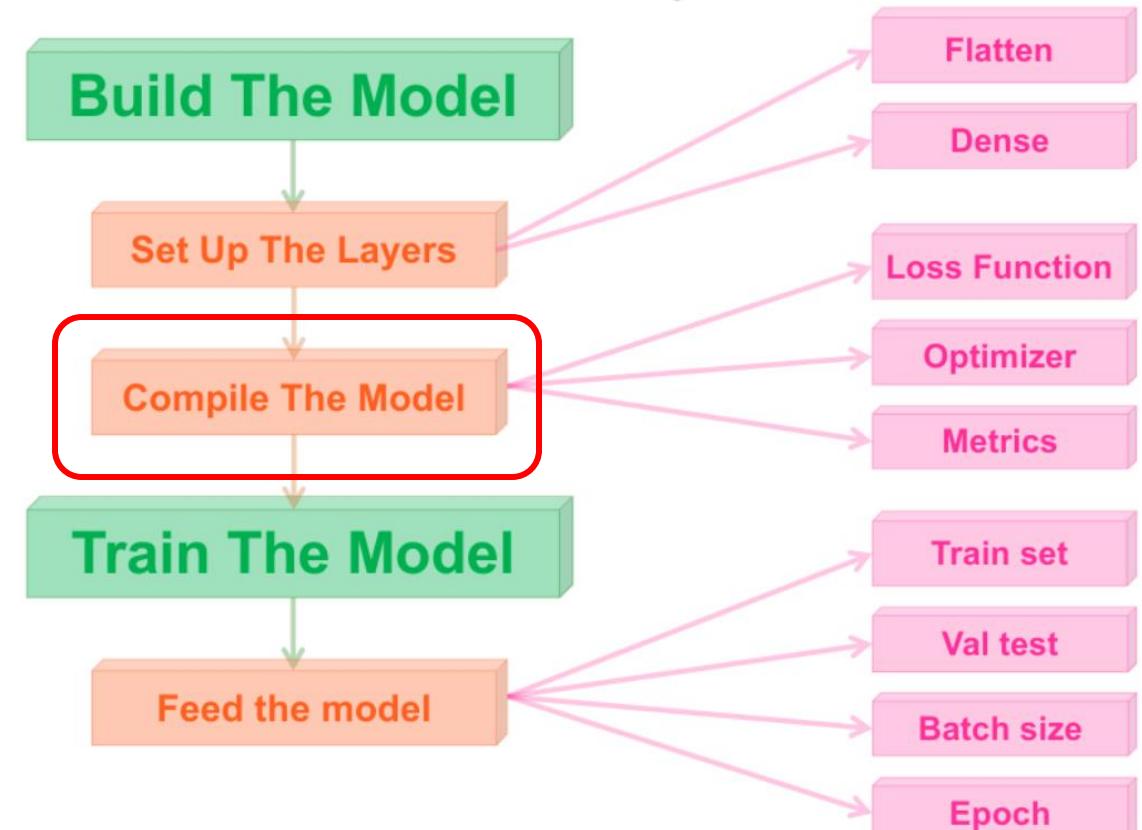


Non-linear Data Classification

❖ Step 6: Declare model, optimizer and criterion

```
1 class MLP(nn.Module):
2     def __init__(self, input_dims, hidden_dims, output_dims):
3         super(MLP, self).__init__()
4         self.linear1 = nn.Linear(input_dims, hidden_dims)
5         self.output = nn.Linear(hidden_dims, output_dims)
6
7     def forward(self, x):
8         x = self.linear1(x)
9         x = F.relu(x)
10        out = self.output(x)
11        return out.squeeze(1)
```

```
1 input_dims = X_train.shape[1]
2 output_dims = torch.unique(y_train).shape[0]
3 hidden_dims = 128
4 lr = 1e-1
5
6 model = MLP(input_dims=input_dims,
7             hidden_dims=hidden_dims,
8             output_dims=output_dims).to(device)
9 criterion = nn.CrossEntropyLoss()
10 optimizer = torch.optim.SGD(model.parameters(), lr=lr)
```

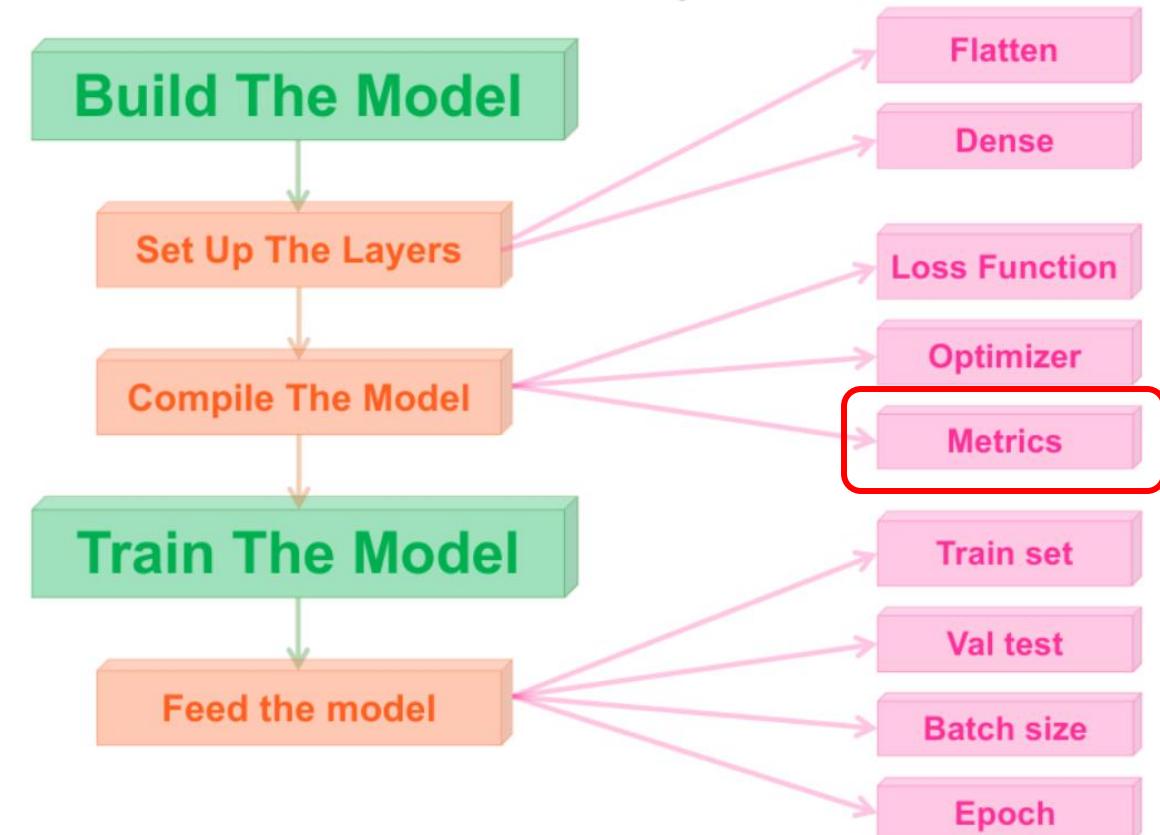


Non-linear Data Classification

❖ Step 7: Create accuracy function

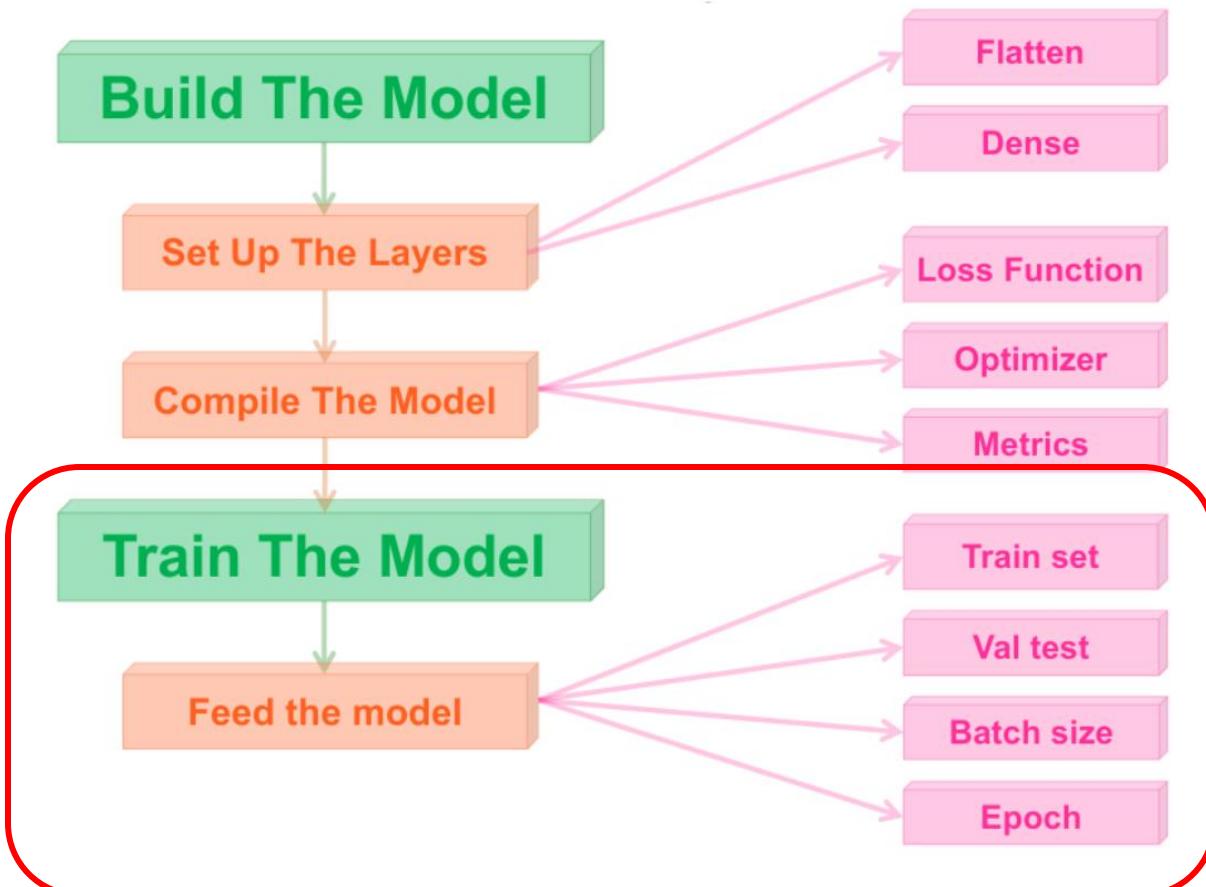
$$\text{accuracy} = \frac{\text{true_predictions}}{\text{n_samples}}$$

```
1 def compute_accuracy(y_hat, y_true):  
2     _, y_hat = torch.max(y_hat, dim=1)  
3     correct = (y_hat == y_true).sum().item()  
4     accuracy = correct / len(y_true)  
5     return accuracy
```



Non-linear Data Classification

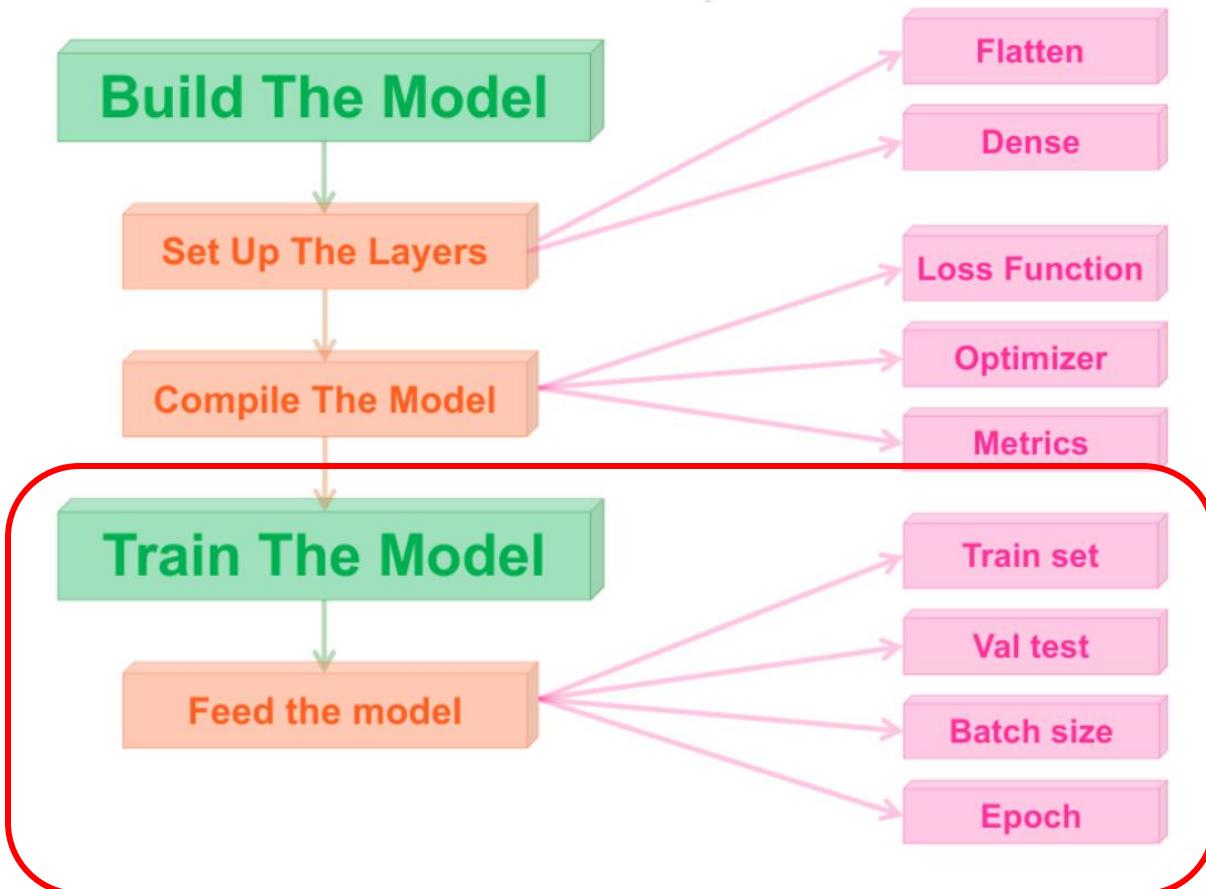
❖ Step 8: Training



```
1 epochs = 100
2 train_losses = []
3 val_losses = []
4 train_accs = []
5 val_accs = []
6
7 for epoch in range(epochs):
8     train_loss = 0.0
9     train_target = []
10    train_predict = []
11    model.train()
12    for X_samples, y_samples in train_loader:
13        X_samples = X_samples.to(device)
14        y_samples = y_samples.to(device)
15        optimizer.zero_grad()
16        outputs = model(X_samples)
17        loss = criterion(outputs, y_samples)
18        loss.backward()
19        optimizer.step()
20        train_loss += loss.item()
21
22        train_predict.append(outputs.detach().cpu())
23        train_target.append(y_samples.cpu())
24
25    train_loss /= len(train_loader)
26    train_losses.append(train_loss)
27
28    train_predict = torch.cat(train_predict)
29    train_target = torch.cat(train_target)
30    train_acc = compute_accuracy(train_predict, train_target)
31    train_accs.append(train_acc)
```

Non-linear Data Classification

❖ Step 8: Training (Evaluation)

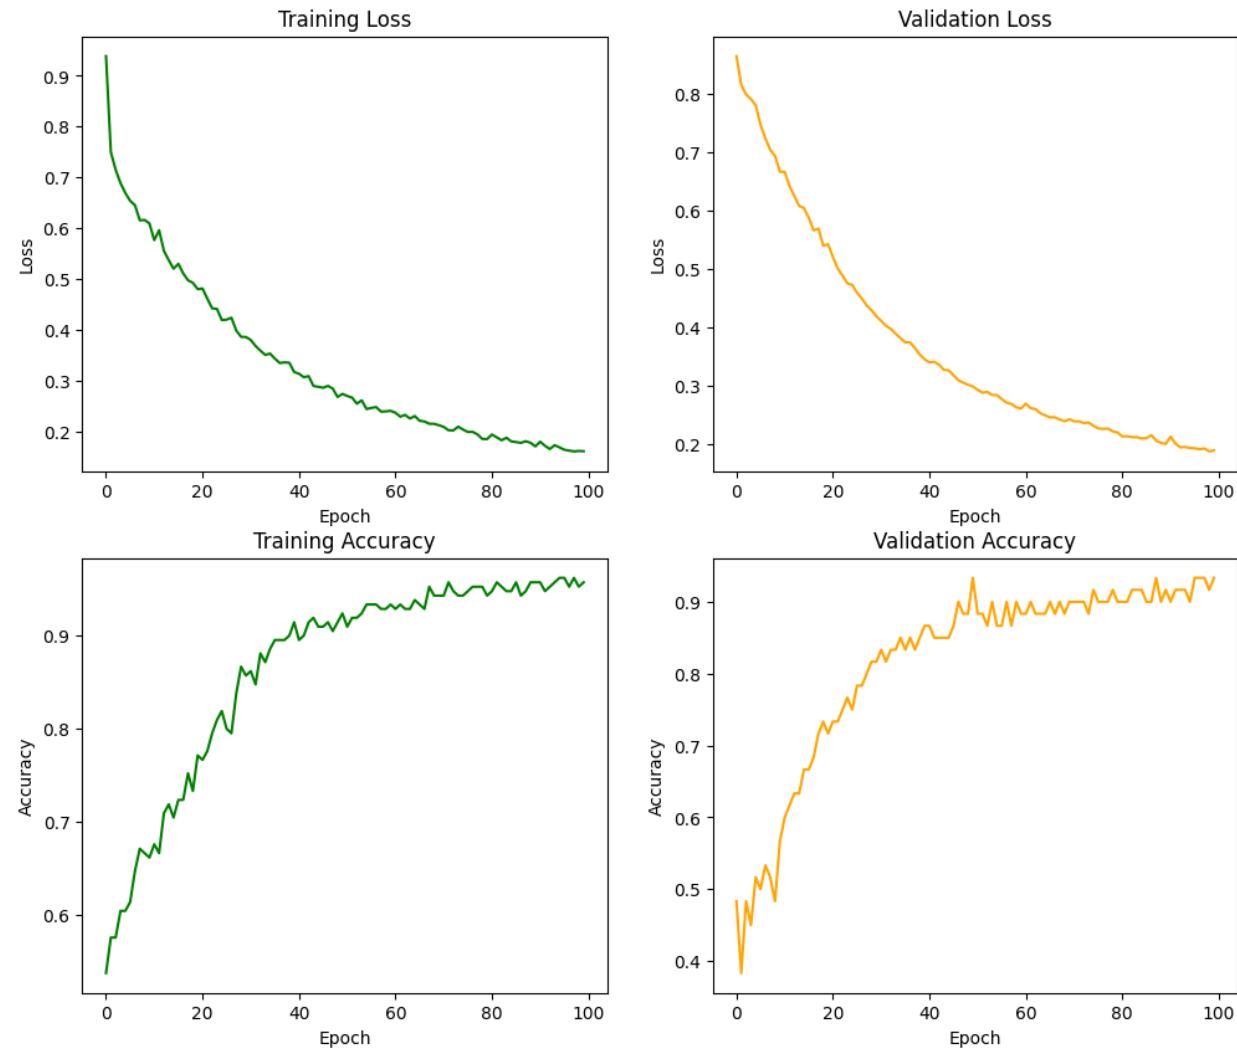


```
33 val_loss = 0.0
34 val_target = []
35 val_predict = []
36 model.eval()
37 with torch.no_grad():
38     for X_samples, y_samples in val_loader:
39         X_samples = X_samples.to(device)
40         y_samples = y_samples.to(device)
41         outputs = model(X_samples)
42         val_loss += criterion(outputs, y_samples).item()
43
44         val_predict.append(outputs.cpu())
45         val_target.append(y_samples.cpu())
46
47 val_loss /= len(val_loader)
48 val_losses.append(val_loss)
49
50 val_predict = torch.cat(val_predict)
51 val_target = torch.cat(val_target)
52 val_acc = compute_accuracy(val_predict, val_target)
53 val_acccs.append(val_acc)
```

Non-linear Data Classification

❖ Step 9: Evaluation

```
1 fig, ax = plt.subplots(2, 2, figsize=(12, 10))
2 ax[0, 0].plot(train_losses, color='green')
3 ax[0, 0].set(xlabel='Epoch', ylabel='Loss')
4 ax[0, 0].set_title('Training Loss')
5
6 ax[0, 1].plot(val_losses, color='orange')
7 ax[0, 1].set(xlabel='Epoch', ylabel='Loss')
8 ax[0, 1].set_title('Validation Loss')
9
10 ax[1, 0].plot(train_accs, color='green')
11 ax[1, 0].set(xlabel='Epoch', ylabel='Accuracy')
12 ax[1, 0].set_title('Training Accuracy')
13
14 ax[1, 1].plot(val_accs, color='orange')
15 ax[1, 1].set(xlabel='Epoch', ylabel='Accuracy')
16 ax[1, 1].set_title('Validation Accuracy')
17
18 plt.show()
```



Non-linear Data Classification

❖ Step 9: Evaluation

$$\text{accuracy} = \frac{\text{true_predictions}}{\text{n_samples}}$$

```
1 def compute_accuracy(y_hat, y_true):
2     _, y_hat = torch.max(y_hat, dim=1)
3     correct = (y_hat == y_true).sum().item()
4     accuracy = correct / len(y_true)
5     return accuracy
```

```
1 val_target = []
2 val_predict = []
3 model.eval()
4 with torch.no_grad():
5     for X_samples, y_samples in val_loader:
6         X_samples = X_samples.to(device)
7         y_samples = y_samples.to(device)
8         outputs = model(X_samples)
9
10        val_predict.append(outputs.cpu())
11        val_target.append(y_samples.cpu())
12
13    val_predict = torch.cat(val_predict)
14    val_target = torch.cat(val_target)
15    val_acc = compute_accuracy(val_predict, val_target)
16
17    print('Evaluation on val set:')
18    print(f'Accuracy: {val_acc}')
```

Evaluation on val set:
Accuracy: 0.9333333333333333

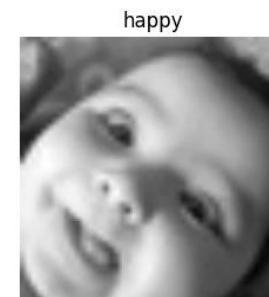
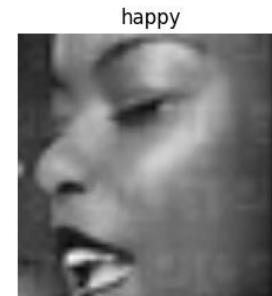
QUIZ

Sentiment Image Analysis

Sentiment Image Analysis

❖ Introduction

Description: Given an image dataset about [face emotion](#), build a Multi-Layer Perceptron (MLP) model capable of accurately predicting the class labels based on the given image.



Sentiment Image Analysis

❖ Step 1: Import libraries and set random seed

```
1 import cv2
2 import os
3 import numpy as np
4 import pandas as pd
5 import matplotlib.pyplot as plt
6 import torch
7 import torch.nn as nn
8 import torch.nn.functional as F
9 from torch.utils.data import Dataset, DataLoader
10 from torchvision.transforms import Resize
11 from torchvision.io import read_image
12
13 from sklearn.model_selection import train_test_split
14 from sklearn.preprocessing import StandardScaler
15
16 device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
17 random_state = 59
18 np.random.seed(random_state)
19 torch.manual_seed(random_state)
20 if torch.cuda.is_available():
21     torch.cuda.manual_seed(random_state)
```



Sentiment Image Analysis

❖ Step 2: Get number of classes in dataset

Because of the dataset structure, we can get all classnames by reading the name of each subfolder in a set.

```
1 train_dir = '/content/train'  
2 test_dir = '/content/test'  
3  
4 classes = os.listdir(train_dir)  
5 classes  
  
['happy', 'angry', 'sad', 'neutral', 'surprise', 'disgust', 'fear']
```

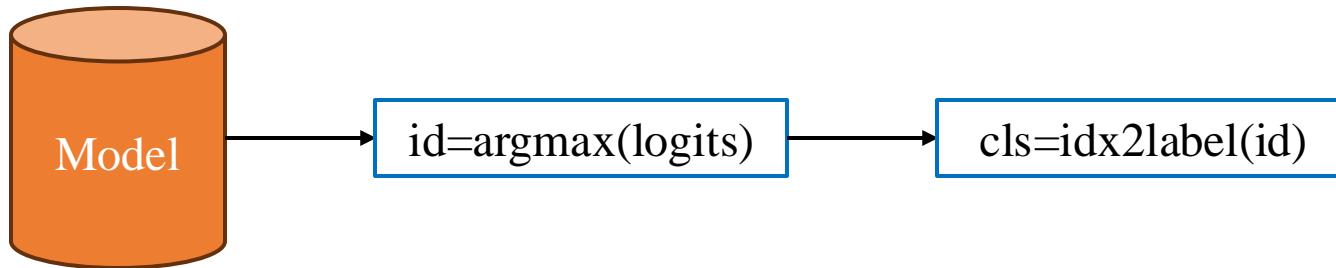


Sentiment Image Analysis

❖ Step 3: Create label-idx mapping

In practice, given a list of classnames, we will implement LabelEncoder in PyTorch as follows:

```
1 label2idx = {cls:idx for idx, cls in enumerate(classes)}  
2 idx2label = {idx:cls for cls, idx in label2idx.items()}
```



1 label2idx

```
{'happy': 0,  
'angry': 1,  
'sad': 2,  
'neutral': 3,  
'surprise': 4,  
'disgust': 5,  
'fear': 6}
```

1 idx2label

```
{0: 'happy',  
1: 'angry',  
2: 'sad',  
3: 'neutral',  
4: 'surprise',  
5: 'disgust',  
6: 'fear'}
```

Sentiment Image Analysis

❖ Step 4: Create PyTorch Dataset: `__init__()`

```
1 class ImageDataset(Dataset):
2     def __init__(self, img_dir, norm, label2idx,
3                  split='train', train_ratio=0.8):
4         self.resize = Resize(img_height, img_width)
5         self.norm = norm
6         self.split = split
7         self.train_ratio = train_ratio
8         self.img_dir = img_dir
9         self.label2idx = label2idx
10        self.img_paths, self.img_labels = self.read_img_files()
11
12        if split in ['train', 'val'] and 'train' in img_dir.lower():
13            train_data, val_data = train_test_split(
14                list(zip(self.img_paths, self.img_labels)),
15                train_size=train_ratio,
16                random_state=random_state,
17                stratify=self.img_labels
18            )
19
20            if split == 'train':
21                self.img_paths, self.img_labels = zip(*train_data)
22            elif split == 'val':
23                self.img_paths, self.img_labels = zip(*val_data)
```

- ▶  test
- ▶  train
-  FER-2013.zip

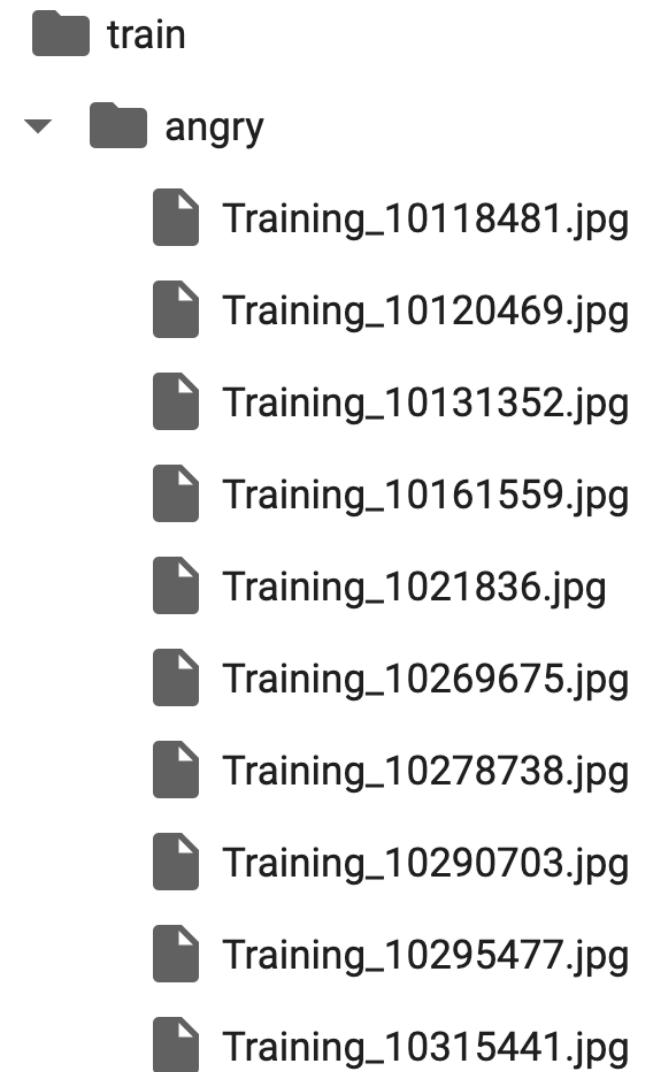
Because we currently have train and test set, we will split validation set from train set.

Sentiment Image Analysis

❖ Step 4: Create PyTorch Dataset: `__read_img_files__()`

We extract pairs of (img_path, img_label) from dataset folder:

```
25     def __read_img_files__(self):
26         img_paths = []
27         img_labels = []
28         for cls in self.label2idx.keys():
29             for img in os.listdir(os.path.join(self.img_dir, cls)):
30                 img_paths.append(os.path.join(self.img_dir, cls, img))
31                 img_labels.append(cls)
32
33         return img_paths, img_labels
```



Sentiment Image Analysis

❖ Step 4: Create PyTorch Dataset: `__len__()` and `__getitem__()`

```
35     def __len__(self):
36         return len(self.img_paths)
37
38     def __getitem__(self, idx):
39         img_path = self.img_paths[idx]
40         cls = self.img_labels[idx]
41         img = self.resize(read_image(img_path))
42         img = img.type(torch.float32)
43         label = self.label2idx[cls]
44         if self.norm:
45             img = (img/127.5) - 1
46         return img, label
```

read_image

```
torchvision.io.read_image(path: str, mode: ImageReadMode = ImageReadMode.UNCHANGED,
apply_exif_orientation: bool = False) → Tensor [SOURCE]
```

Reads a JPEG, PNG or GIF image into a 3 dimensional RGB or grayscale Tensor. Optionally converts the image to the desired format. The values of the output tensor are uint8 in [0, 255].

Resize

```
CLASS torchvision.transforms.Resize(size, interpolation=InterpolationMode.BILINEAR,
max_size=None, antialias=True) [SOURCE]
```

Resize the input image to the given size. If the image is torch Tensor, it is expected to have [..., H, W] shape, where ... means a maximum of two leading dimensions

Sentiment Image Analysis

❖ Step 4: Create PyTorch Dataset: `__len__()` and `__getitem__()`

```
35     def __len__(self):
36         return len(self.img_paths)
37
38     def __getitem__(self, idx):
39         img_path = self.img_paths[idx]
40         cls = self.img_labels[idx]
41         img = self.resize(read_image(img_path))
42         img = img.type(torch.float32)
43         label = self.label2idx[cls]
44         if self.norm:
45             img = (img/127.5) - 1
46         return img, label
```

read_image

```
torchvision.io.read_image(path: str, mode: ImageReadMode = ImageReadMode.UNCHANGED,
apply_exif_orientation: bool = False) → Tensor [SOURCE]
```

Reads a JPEG, PNG or GIF image into a 3 dimensional RGB or grayscale Tensor. Optionally converts the image to the desired format. The values of the output tensor are uint8 in [0, 255].

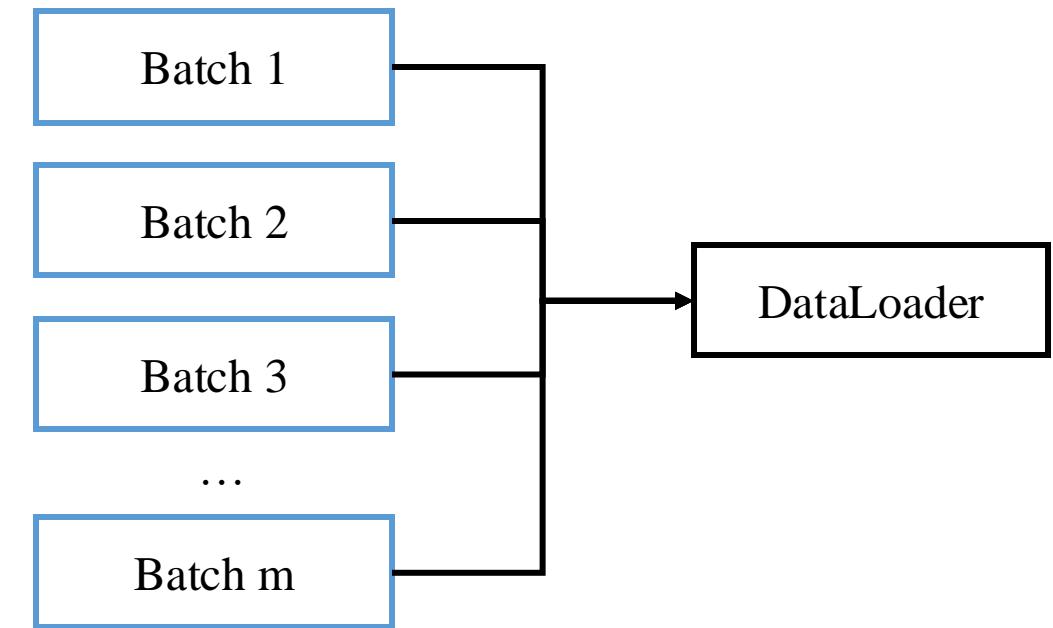
Resize

```
CLASS torchvision.transforms.Resize(size, interpolation=InterpolationMode.BILINEAR,
max_size=None, antialias=True) [SOURCE]
```

Resize the input image to the given size. If the image is torch Tensor, it is expected to have [..., H, W] shape, where ... means a maximum of two leading dimensions

Sentiment Image Analysis

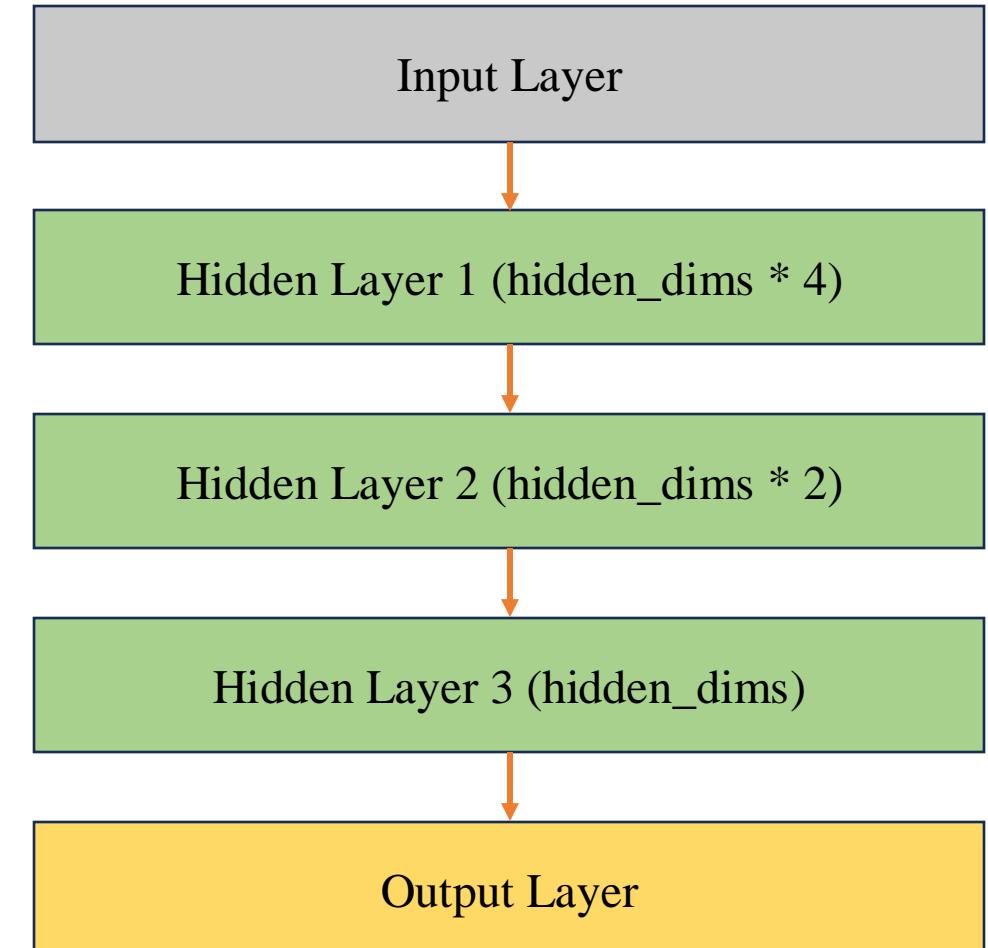
❖ Step 5: Create PyTorch Dataloader



Sentiment Image Analysis

❖ Step 6: Build MLP

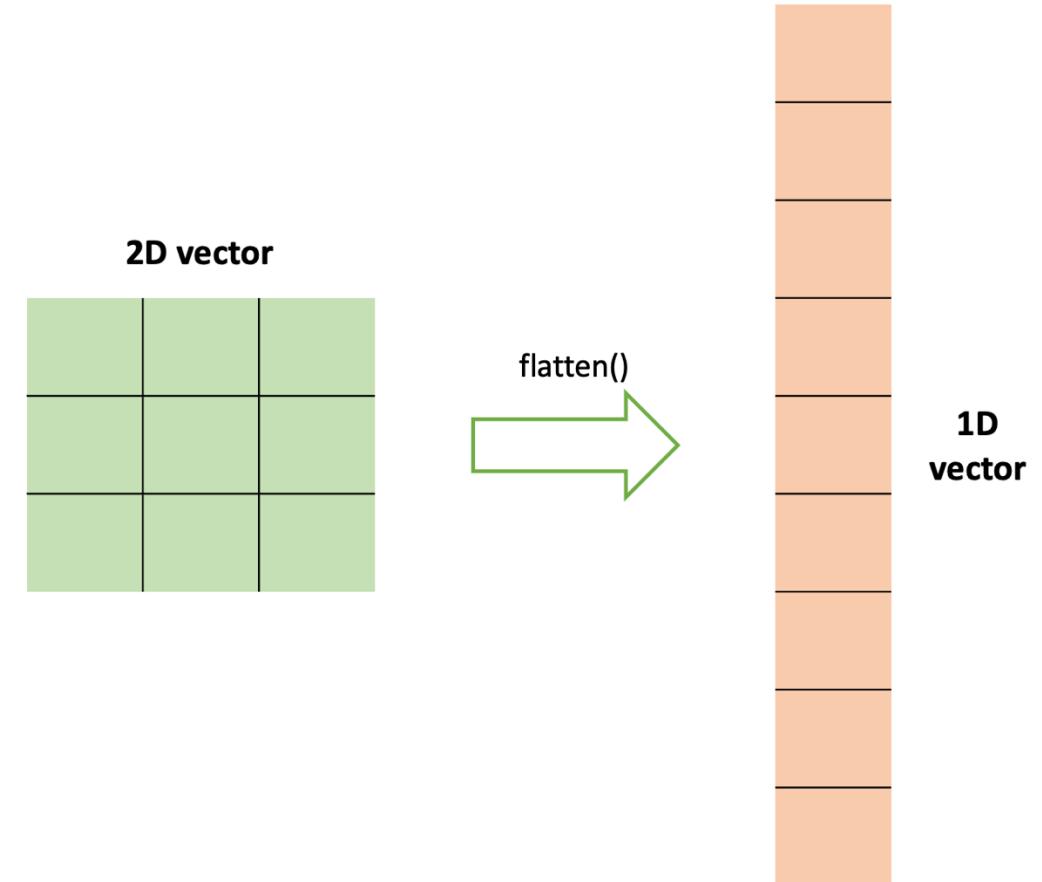
```
1 class MLP(nn.Module):
2     def __init__(self, input_dims, hidden_dims, output_dims):
3         super(MLP, self).__init__()
4         self.linear1 = nn.Linear(input_dims, hidden_dims*4)
5         self.linear2 = nn.Linear(hidden_dims*4, hidden_dims*2)
6         self.linear3 = nn.Linear(hidden_dims*2, hidden_dims)
7         self.output = nn.Linear(hidden_dims, output_dims)
8
9     def forward(self, x):
10        x = nn.Flatten()(x)
11        x = self.linear1(x)
12        x = F.relu(x)
13        x = self.linear2(x)
14        x = F.relu(x)
15        x = self.linear3(x)
16        x = F.relu(x)
17        out = self.output(x)
18        return out.squeeze(1)
```



Sentiment Image Analysis

❖ Step 6: Build MLP: Flatten()

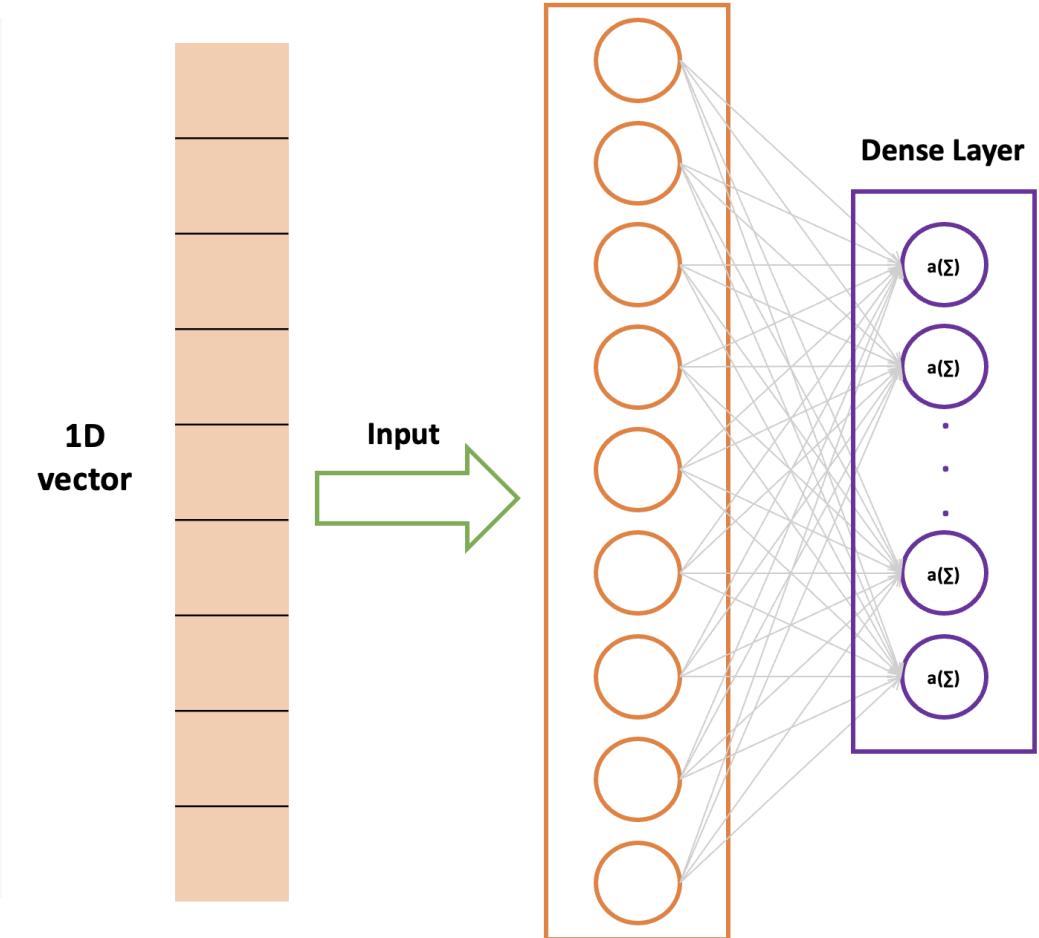
```
1 class MLP(nn.Module):
2     def __init__(self, input_dims, hidden_dims, output_dims):
3         super(MLP, self).__init__()
4         self.linear1 = nn.Linear(input_dims, hidden_dims*4)
5         self.linear2 = nn.Linear(hidden_dims*4, hidden_dims*2)
6         self.linear3 = nn.Linear(hidden_dims*2, hidden_dims)
7         self.output = nn.Linear(hidden_dims, output_dims)
8
9     def forward(self, x):
10        x = nn.Flatten()(x)
11        x = self.linear1(x)
12        x = F.relu(x)
13        x = self.linear2(x)
14        x = F.relu(x)
15        x = self.linear3(x)
16        x = F.relu(x)
17        out = self.output(x)
18        return out.squeeze(1)
```



Sentiment Image Analysis

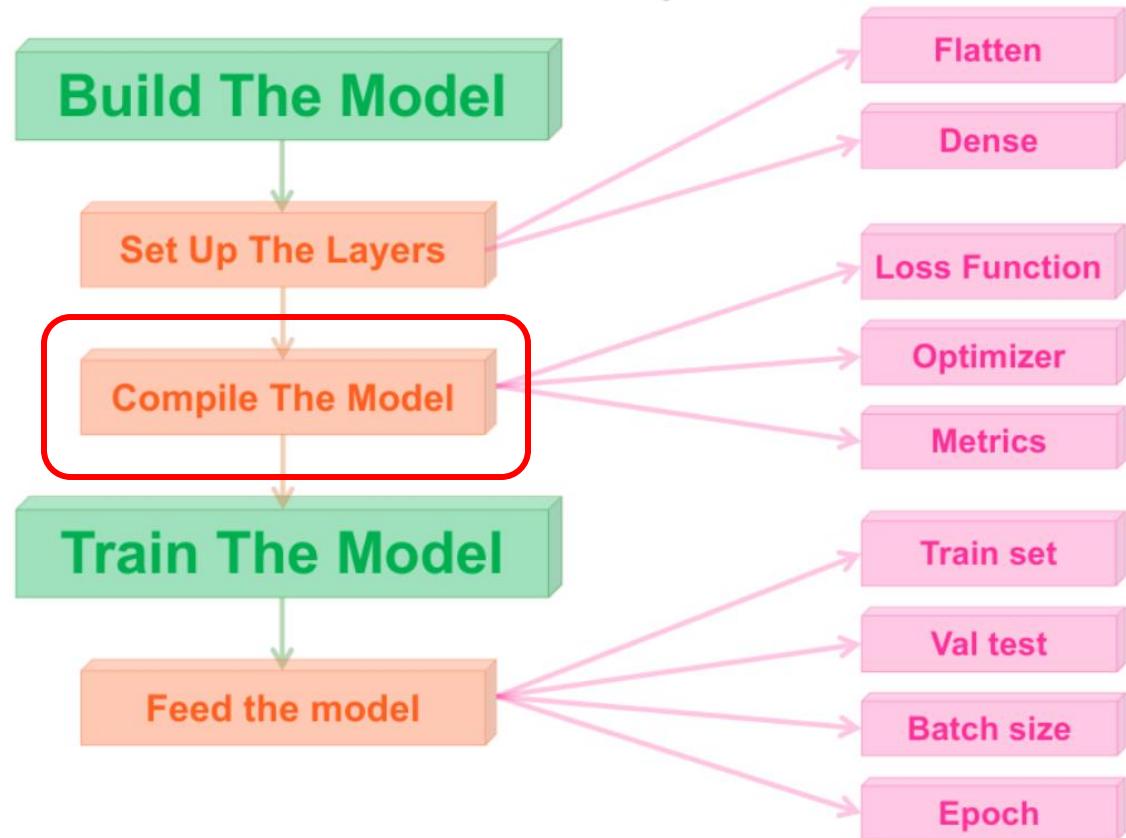
❖ Step 6: Build MLP: Feeding 1D vector image into the model

```
1 class MLP(nn.Module):
2     def __init__(self, input_dims, hidden_dims, output_dims):
3         super(MLP, self).__init__()
4         self.linear1 = nn.Linear(input_dims, hidden_dims*4)
5         self.linear2 = nn.Linear(hidden_dims*4, hidden_dims*2)
6         self.linear3 = nn.Linear(hidden_dims*2, hidden_dims)
7         self.output = nn.Linear(hidden_dims, output_dims)
8
9     def forward(self, x):
10        x = nn.Flatten()(x)
11        x = self.linear1(x)
12        x = F.relu(x)
13        x = self.linear2(x)
14        x = F.relu(x)
15        x = self.linear3(x)
16        x = F.relu(x)
17        out = self.output(x)
18        return out.squeeze(1)
```



Sentiment Image Analysis

❖ Step 7: Declare the model, optimizer and criterion



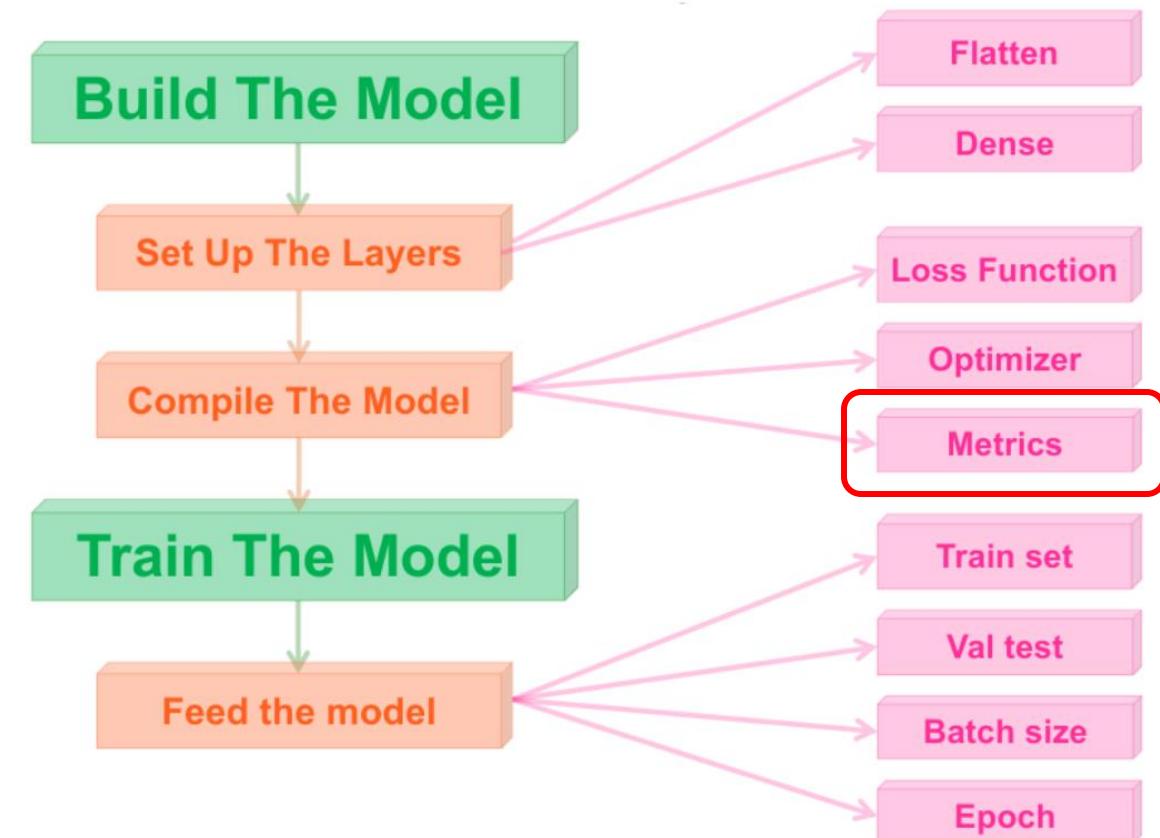
```
1 input_dims = img_height * img_width
2 output_dims = len(classes)
3 hidden_dims = 64
4 lr = 1e-2
5
6 model = MLP(input_dims=input_dims,
7               hidden_dims=hidden_dims,
8               output_dims=output_dims).to(device)
9 criterion = nn.CrossEntropyLoss()
10 optimizer = torch.optim.SGD(model.parameters(), lr=lr)
```

Sentiment Image Analysis

❖ Step 8: Create accuracy function

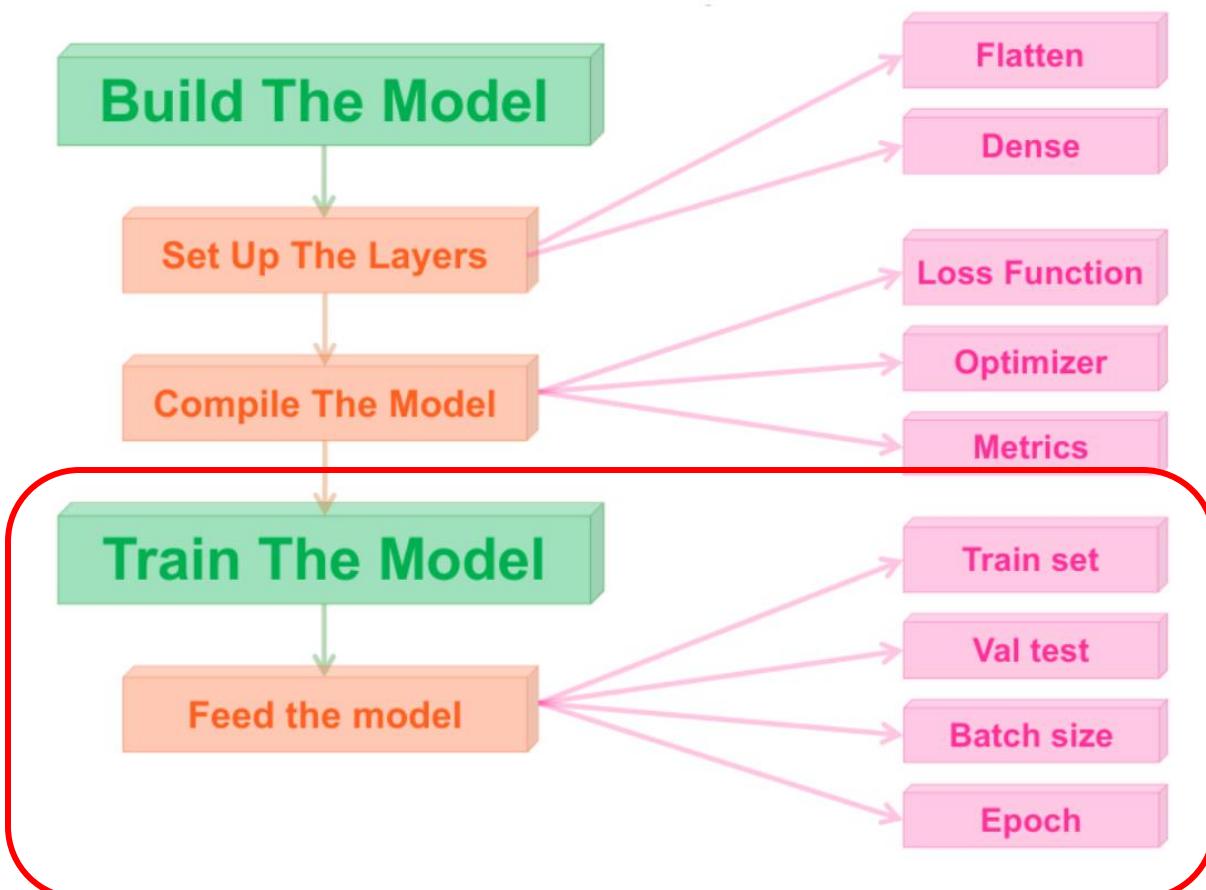
$$\text{accuracy} = \frac{\text{true_predictions}}{\text{n_samples}}$$

```
1 def compute_accuracy(y_hat, y_true):  
2     _, y_hat = torch.max(y_hat, dim=1)  
3     correct = (y_hat == y_true).sum().item()  
4     accuracy = correct / len(y_true)  
5     return accuracy
```



Sentiment Image Analysis

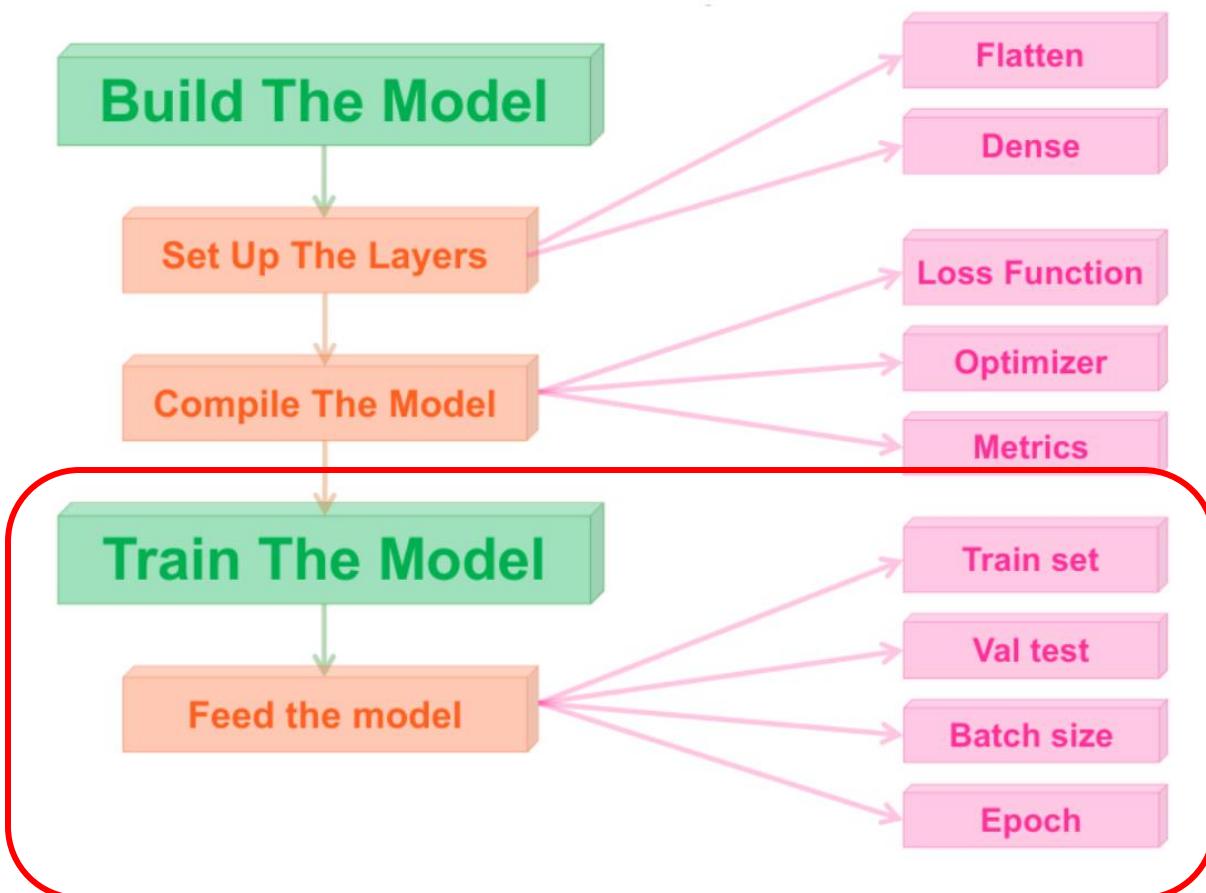
❖ Step 9: Training



```
1 epochs = 100
2 train_losses = []
3 val_losses = []
4 train_accs = []
5 val_accs = []
6
7 for epoch in range(epochs):
8     train_loss = 0.0
9     train_target = []
10    train_predict = []
11    model.train()
12    for X_samples, y_samples in train_loader:
13        X_samples = X_samples.to(device)
14        y_samples = y_samples.to(device)
15        optimizer.zero_grad()
16        outputs = model(X_samples)
17        loss = criterion(outputs, y_samples)
18        loss.backward()
19        optimizer.step()
20        train_loss += loss.item()
21
22        train_predict.append(outputs.detach().cpu())
23        train_target.append(y_samples.cpu())
24
25    train_loss /= len(train_loader)
26    train_losses.append(train_loss)
27
28    train_predict = torch.cat(train_predict)
29    train_target = torch.cat(train_target)
30    train_acc = compute_accuracy(train_predict, train_target)
31    train_accs.append(train_acc)
```

Sentiment Image Analysis

❖ Step 9: Training (Evaluation)

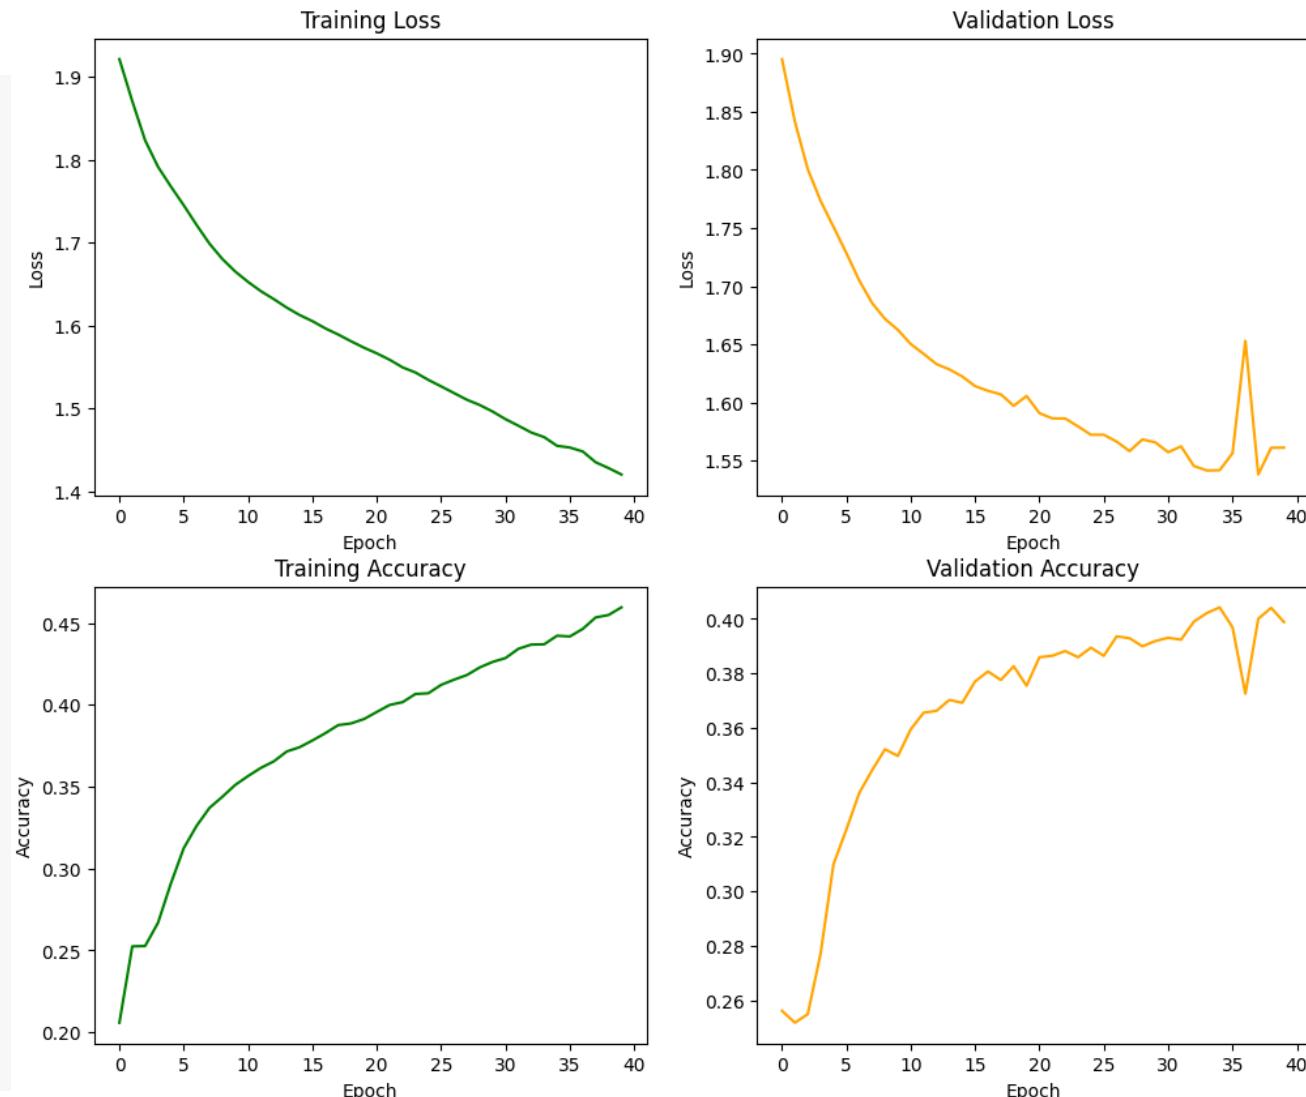


```
33 val_loss = 0.0
34 val_target = []
35 val_predict = []
36 model.eval()
37 with torch.no_grad():
38     for X_samples, y_samples in val_loader:
39         X_samples = X_samples.to(device)
40         y_samples = y_samples.to(device)
41         outputs = model(X_samples)
42         val_loss += criterion(outputs, y_samples).item()
43
44         val_predict.append(outputs.cpu())
45         val_target.append(y_samples.cpu())
46
47 val_loss /= len(val_loader)
48 val_losses.append(val_loss)
49
50 val_predict = torch.cat(val_predict)
51 val_target = torch.cat(val_target)
52 val_acc = compute_accuracy(val_predict, val_target)
53 val_acccs.append(val_acc)
```

Sentiment Image Analysis

❖ Step 10: Evaluation

```
1 fig, ax = plt.subplots(2, 2, figsize=(12, 10))
2 ax[0, 0].plot(train_losses, color='green')
3 ax[0, 0].set(xlabel='Epoch', ylabel='Loss')
4 ax[0, 0].set_title('Training Loss')
5
6 ax[0, 1].plot(val_losses, color='orange')
7 ax[0, 1].set(xlabel='Epoch', ylabel='Loss')
8 ax[0, 1].set_title('Validation Loss')
9
10 ax[1, 0].plot(train_accs, color='green')
11 ax[1, 0].set(xlabel='Epoch', ylabel='Accuracy')
12 ax[1, 0].set_title('Training Accuracy')
13
14 ax[1, 1].plot(val_accs, color='orange')
15 ax[1, 1].set(xlabel='Epoch', ylabel='Accuracy')
16 ax[1, 1].set_title('Validation Accuracy')
17
18 plt.show()
```



Sentiment Image Analysis

❖ Step 10: Evaluation

$$\text{accuracy} = \frac{\text{true_predictions}}{\text{n_samples}}$$

```
1 def compute_accuracy(y_hat, y_true):
2     _, y_hat = torch.max(y_hat, dim=1)
3     correct = (y_hat == y_true).sum().item()
4     accuracy = correct / len(y_true)
5     return accuracy
```

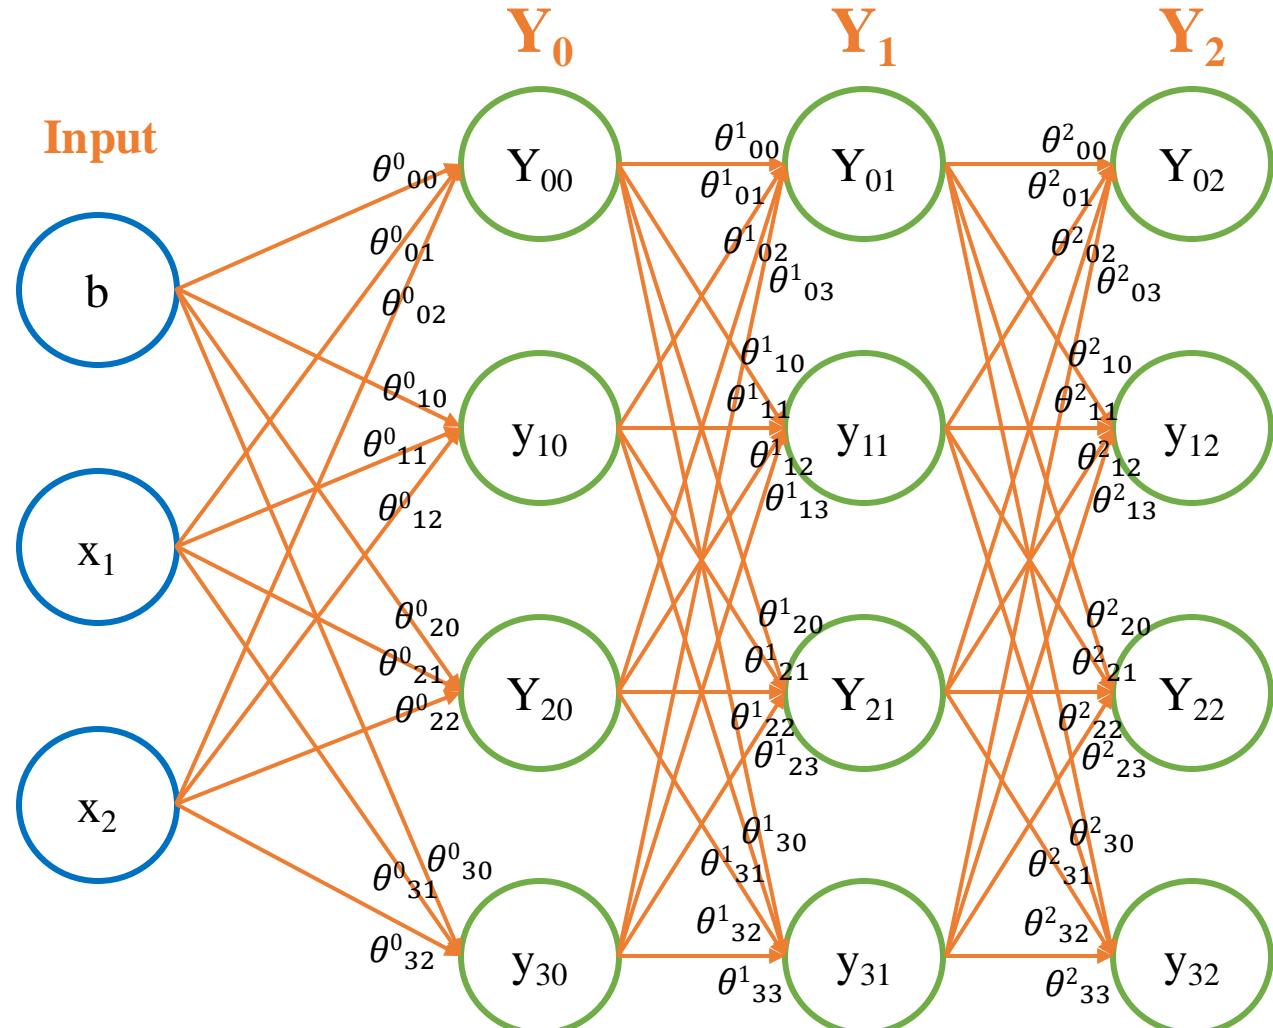
```
1 val_target = []
2 val_predict = []
3 model.eval()
4 with torch.no_grad():
5     for X_samples, y_samples in val_loader:
6         X_samples = X_samples.to(device)
7         y_samples = y_samples.to(device)
8         outputs = model(X_samples)
9
10        val_predict.append(outputs.cpu())
11        val_target.append(y_samples.cpu())
12
13    val_predict = torch.cat(val_predict)
14    val_target = torch.cat(val_target)
15    val_acc = compute_accuracy(val_predict, val_target)
16
17    print('Evaluation on val set:')
18    print(f'Accuracy: {val_acc}')
```

Evaluation on val set:
Accuracy: 0.39881574364332983

Summarization and Q&A

Summarization and Q&A

❖ Summarization



In this lecture, we have discussed about:

- Multilayer Perceptrons (MLPs).
- Activation functions.
- Practice how to implement and train MLP model on tabular data, non-linear data and image data using PyTorch.

Summarization and Q&A

❖ Q&A

